

SAM 3: SEGMENT ANYTHING WITH CONCEPTS

Anonymous authors

Paper under double-blind review

ABSTRACT

We present Segment Anything Model (SAM) 3, a unified model that detects, segments, and tracks objects in images and videos based on *concept prompts*, which we define as either short noun phrases (e.g., “yellow school bus”), image exemplars, or a combination of both. Promptable Concept Segmentation (PCS) takes such prompts and returns segmentation masks and unique identities for all matching object instances. To advance PCS, we build a scalable data engine that produces a high-quality dataset with 4M unique concept labels, including hard negatives, across images and videos. Our model consists of a vision backbone shared between an image-level detector and a memory-based video tracker. Recognition and localization are decoupled with a presence head, which significantly boosts detection accuracy. SAM 3 delivers a $2\times$ gain over existing systems in both image and video PCS, and improves previous SAM capabilities in interactive visual segmentation tasks. We open source SAM 3 along with our new Segment Anything with Concepts (SA-Co) benchmark.

1 INTRODUCTION

The ability to find and segment *anything* in a visual scene is foundational for multimodal AI, powering applications in robotics, content creation, augmented reality, data annotation, and scientific fields. The SAM series (Kirillov et al., 2023; Ravi et al., 2024) introduced the promptable segmentation task to segment objects in images and videos via interactive prompts, including visual inputs like points, boxes, and masks marking a specific object, or text inputs describing an object. However, SAM 1 and SAM 2 focus on visual prompts and segment a single object instance per prompt. While these methods achieved a breakthrough for this critical task, they did not address the broader task of finding and segmenting *all* instances of a concept appearing *anywhere* in the input (e.g., all “cats” in a video).

In this work, we present SAM 3, a model that achieves a step change in promptable segmentation in images and videos, improving Promptable Visual Segmentation (PVS) relative to SAM 2 and setting a new standard for *Promptable Concept Segmentation (PCS)*. We formalize the PCS task as taking text and/or image exemplars as input, and predicting instance and semantic masks for every single object matching the concept, while preserving object identities across video frames (§2). We focus on recognizing atomic visual concepts and thus constrain text to simple noun phrases (NPs), such as “red apple” or “striped cat”. Example outputs are shown in Fig. 1.

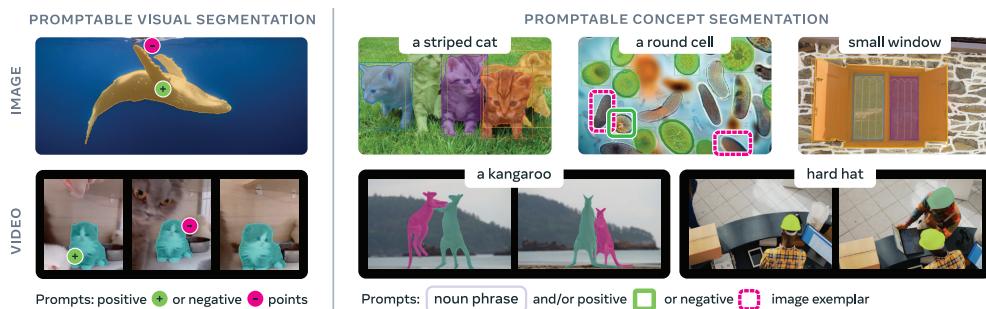


Figure 1: SAM 3 improves over SAM 2 on promptable *visual* segmentation with clicks (left) while advancing promptable *concept* segmentation (right) where users can segment all instances of a visual concept specified by a short noun phrase, image exemplars, or a combination of both.



Figure 2: Illustration of supported initial and optional interactive refinement prompts in the PCS task.

While SAM 3 is not designed for long referring language expressions or queries requiring reasoning, we show that it can be straightforwardly combined with a Multimodal Large Language Model (MLLM) to handle more complex language prompts.

The PCS task is inherently ambiguous due to its open-vocabulary nature, with many concepts having multiple interpretations; e.g., “small window” is subjective (small vs large) and has ambiguous boundaries (with shutters vs without). Our approach systematically accounts for these ambiguities at every stage, including data collection, metric design, and modeling. Consistent with previous SAM versions, SAM 3 is fully interactive, allowing users to resolve ambiguities by adding refinement prompts to guide the model towards their intended output.

Our *model* (§3) consists of a detector and a tracker that share a vision encoder (Bolya et al., 2025). The detector is a DETR-based (Carion et al., 2020) model conditioned on text, geometry, and image exemplars. To address the challenge of open-vocabulary concept detection, we introduce a separate *presence head* to decouple recognition and localization, which is especially effective when training with challenging negative phrases. The tracker inherits the SAM 2 transformer encoder-decoder architecture, supporting video segmentation and interactive refinement. The decoupled design for detection and tracking avoids task conflict, as the detector needs to be identity agnostic, while the tracker’s main objective is to separate identities in the video.

To unlock major performance gains, we build a scalable human- and model-in-the-loop *data engine* (§4) that annotates a large and diverse training dataset. We innovate upon prior data engine designs in three key ways: i) media curation: we curate more diverse media domains than past approaches that rely on homogeneous web sources, ii) label curation: we significantly increase label diversity and difficulty by leveraging an ontology and multimodal LLMs as “AI annotators” to generate noun phrases and hard negatives, iii) label verification: we double annotation throughput by fine-tuning MLLMs to be effective “AI verifiers” that achieve near-human performance. Starting from noisy media-phrase-mask pseudolabels, our data engine checks mask quality and exhaustivity using both human and AI verifiers, filtering out correctly labeled examples and identifying challenging error cases. Human annotators then focus on fixing these errors by manually correcting masks. This enables us to annotate high-quality training data with 4M unique phrases and 52M masks, and a synthetic dataset with 38M phrases and 1.4B masks. We additionally create the Segment Anything with Concepts (SA-Co) *benchmark* for PCS (§5) containing 214K unique concepts with exhaustive masks in 124K images and 1.7K videos, $> 50\times$ more concepts than existing benchmarks.

Our *experiments* (§6) show that SAM 3 sets a new state-of-the-art in promptable segmentation, e.g., reaching a zero-shot mask AP of 47.0 on LVIS vs the current best of 38.5, surpassing baselines on our new SA-Co benchmark by at least $2\times$, and improving upon SAM 2 on PVS benchmarks. Ablations (§B) verify that the choice of backbone, novel presence head, and adding hard negatives all boost results, and establish scaling laws on the PCS task for both our high-quality and synthetic datasets. We open-source the SA-Co benchmark and release the SAM 3 checkpoints and inference code. On an H200 GPU, SAM 3 runs in 30 ms for a single image with 100+ detected objects. In video, the inference latency scales with the number of objects, sustaining near real-time performance for ~ 5 concurrent objects. Next we dive into the task, and review related work in §A.

2 PROMPTABLE CONCEPT SEGMENTATION (PCS)

We define the Promptable Concept Segmentation task as follows: given an image or short video (≤ 30 secs), detect, segment and track all instances of a visual concept specified by a short text phrase, image exemplars, or a combination of both. We restrict concepts to those defined by simple noun phrases (NPs) consisting of a noun and optional modifiers. Noun-phrase prompts (when provided) are *global* to all frames of the image/video, while image exemplars can be provided on *individual* frames as positive or negative bounding boxes to iteratively *refine* the target masks (see Fig. 2).

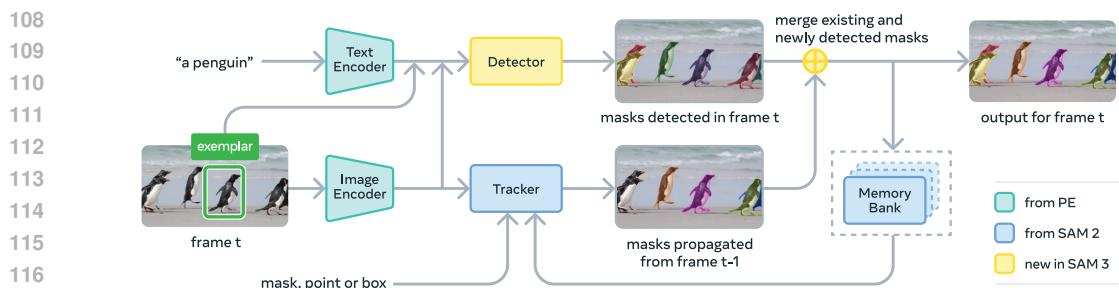


Figure 3: SAM 3 architecture overview. See Fig. 8 for a more detailed diagram.

All prompts must be consistent in their category definition, or the model’s behavior is undefined; e.g., “fish” cannot be refined with subsequent exemplar prompts of just the tail; instead the text prompt should be updated. Exemplar prompts are particularly useful when the model initially misses some instances, or when the concept is rare.

Our vocabulary includes any simple noun phrase groundable in a visual scene, which makes the task intrinsically ambiguous. There can be multiple interpretations of phrases arising from polysemy (“mouse” device vs. animal), subjective descriptors (“cozy”, “large”), vague or context-dependent phrases that may not even be groundable (“brand identity”), boundary ambiguity (whether “mirror” includes the frame), and factors like occlusion and blur that obscure object extent. While similar issues appear in large closed-vocabulary corpora (e.g., LVIS (Gupta et al., 2019)), they are alleviated by carefully curating the vocabulary and setting a clear definition of all the classes of interest. We address the ambiguity problem by collecting test annotations from three experts, adapting the evaluation protocol to allow multiple valid interpretations (§F.3), designing the data pipeline and guidelines to minimize ambiguity during annotation, and introducing an ambiguity module in the model (§D.2).

3 MODEL

SAM 3 is a generalization of SAM 2, supporting the new PCS task (§2) as well as the PVS task. It takes *concept* prompts (simple noun phrases, image exemplars) or *visual* prompts (points, boxes, masks) to define the *objects* to be (individually) segmented spatio-temporally. Image exemplars and visual prompts can be *iteratively* added on individual frames to *refine* the target masks—false positive and false negative objects can be *removed* or *added* respectively using image exemplars and an *individual* mask(let) can be refined using PVS in the style of SAM 2. Our architecture is broadly based on the SAM and (M)DETR (Carion et al., 2020; Kamath et al., 2021) series. Fig. 3 shows the SAM 3 architecture, consisting of a dual encoder-decoder transformer—a *detector* for image-level capabilities—which is used in combination with a *tracker* and memory for video. The detector and tracker ingest vision-language inputs from an aligned Perception Encoder (PE) backbone (Bolya et al., 2025). We present an overview below, see §D for details.

Detector Architecture. The architecture of the detector follows the general DETR paradigm. The image and text prompt are first encoded by PE and image exemplars, if present, are encoded by an exemplar encoder. We refer to the image exemplar tokens and text tokens jointly as “prompt tokens”. The fusion encoder then accepts the unconditioned embeddings from the image encoder and conditions them by cross-attending to the prompt tokens. The fusion is followed by a DETR-like decoder, where learned object queries cross-attend to the conditioned image embeddings from the fusion encoder. Each layer, for each object query, predicts a classification logit (in our case, a binary label of whether the object corresponds to the prompt), and a delta from the bounding box predicted by the previous level, following Zhu et al. (2020). We use box-region-positional bias (Lin et al., 2023) to help focalize the attention on each object, but unlike recent DETR models, we stick to vanilla attention. During training, we adopt dual supervision from DAC-DETR (Hu et al., 2023), and the Align loss (Cai et al., 2024). The mask head is adapted from MaskFormer (Cheng et al., 2021). In addition, we also have a semantic segmentation head, which predicts a binary label for every pixel in the image, indicating whether or not it corresponds to the prompt. See §D for details.

Presence Token. It can be difficult for each of the proposal queries to both recognize (what) and localize (where) an object in the image/frame. For the recognition component, contextual cues from the entire image are important. However, forcing proposal queries to understand the global context can be counterproductive, as it conflicts with the inherently local nature of the localization objective.

We decouple the recognition and localization steps by introducing a learned global *presence token*. This token is solely responsible for predicting whether the target concept in the form of a noun phrase (NP) is present in the image/frame, i.e. $p(\text{NP is present in input})$. Each proposal query q_i only needs to solve the localization problem $p(q_i \text{ is a match} \mid \text{NP is present in input})$. The final score for each proposal query is the product of its own score and the presence score.

Image Exemplars and Interactivity. SAM 3 supports image exemplars, given as a pair—a bounding box and an associated binary label (positive or negative)—which can be used in isolation or to supplement the text prompt. The model then detects all the instances that match the prompt. For example, given a positive bounding box on a dog, the model will detect *all* dogs in the image. This is different from the PVS task in SAM 1 and 2, where a geometric/visual prompt yields only a single object instance. Each image exemplar is encoded separately by the exemplar encoder using an embedding for the position, an embedding for the label, and ROI-pooled visual features, then concatenated and processed by a small transformer. The resulting prompt is concatenated to the text prompt to comprise the prompt tokens. Image exemplars can be *interactively* provided based on errors in current detections to refine the output.

Tracker and Video Architecture. Given a video and a prompt P , we use the detector and a tracker (see Fig. 3) to detect and track objects corresponding to the prompt throughout the video. On each frame, the detector finds new objects \mathcal{O}_t and the tracker propagates masklets \mathcal{M}_{t-1} (spatial-temporal masks) from frames at the previous time $t - 1$ to their new locations $\hat{\mathcal{M}}_t$ on the current frame at time t . We use a matching function to associate propagated masklets $\hat{\mathcal{M}}_t$ with new object masks emerging in the current frame \mathcal{O}_t .

$$\hat{\mathcal{M}}_t = \text{propagate}(\mathcal{M}_{t-1}), \quad \mathcal{O}_t = \text{detect}(I_t, P), \quad \mathcal{M}_t = \text{match_and_update}(\hat{\mathcal{M}}_t, \mathcal{O}_t).$$

Tracking an Object with SAM 2 Style Propagation. A masklet is initialized for every object detected on the first frame. Then, on each subsequent frame, the tracker module predicts the new masklet locations $\hat{\mathcal{M}}_t$ of those already-tracked objects based on their previous locations \mathcal{M}_{t-1} through a single-frame propagation step similar to the video object segmentation task in SAM 2. The tracker shares the same image/frame encoder (PE backbone) as the detector. After training the detector, we freeze PE and train the tracker as in SAM 2, including a prompt encoder, mask decoder, memory encoder, and a memory bank that encodes the object’s appearance using features from the past frames and conditioning frames (frames where the object is first detected or user-prompted). The memory encoder is a transformer with self-attention across visual features on the current frame and cross-attention from the visual features to the spatial memory features in the memory bank.

During inference, we only retain frames where the object is confidently present in the memory bank. The mask decoder is a two-way transformer between the encoder hidden states and the output tokens. To handle ambiguity, we predict three output masks for every tracked object on each frame along with their confidence, and select the most confident output as the predicted mask on the current frame.

Matching and Updating Based on Detections. After obtaining the tracked masks $\hat{\mathcal{M}}_t$, we match them with the current frame detections \mathcal{O}_t through a simple IoU based *matching function* (§D.3) and add them to \mathcal{M}_t on the current frame. We further spawn new masklets for all newly detected objects that are not matched. The merging might suffer from ambiguities, especially in crowded scenes. We address this with two temporal disambiguation strategies outlined next.

First, we use temporal information in the form of a *masklet detection score* (§D.3) to measure how consistently a masklet is matched to a detection within a temporal window (based on the number of past frames where it was matched to a detection). If a masklet’s detection score falls below a threshold, we suppress it. Second, we use the detector outputs to resolve specific failure modes of the tracker due to occlusions or distractors. We periodically *re-prompt* the tracker with high-confidence *detection* masks \mathcal{O}_t , replacing the tracker’s own predictions $\hat{\mathcal{M}}_t$. This ensures that the memory bank has recent and reliable references (other than the tracker’s own predictions).

Instance Refinement with Visual Prompts. After obtaining the initial set of masks (or masklets), SAM 3 allows refining individual masks(lets) using positive and negative clicks. Specifically, given the user clicks, we apply the prompt encoder to encode them, and feed the encoded prompt into the mask decoder to predict an adjusted mask. In videos the mask is then propagated across the entire video to obtain a refined masklet.

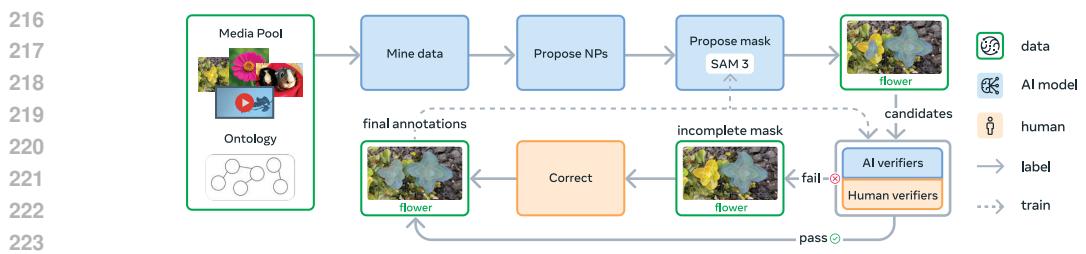


Figure 4: Overview of the final SAM 3 data engine. See §F.1 for details of collected data.

Training Stages. We train SAM 3 in four stages that progressively add data and capabilities: 1) Perception Encoder (PE) pre-training, 2) detector pre-training, 3) detector fine-tuning, and 4) tracker training with a frozen backbone. See §D.4.1 for details.

4 DATA ENGINE

Achieving a step change on PCS with SAM 3 requires training on a large and diverse set of concepts and visual domains, beyond any existing dataset (see Figure 10). Our solution is to build an efficient data engine that iteratively generates annotated data via a feedback loop with SAM 3, human annotators, and *AI annotators*, actively mining media-phrase pairs on which the current version of SAM 3 fails to produce high-quality training data to further improve the model. By delegating certain tasks to *AI annotators*—models that match or surpass human accuracy—we more than double the throughput compared to a human-only annotation pipeline. We develop the data engine in four phases, with each phase increasing the use of *AI models* to steer human effort to the most challenging failure cases, alongside expanding visual domain coverage. Phases 1-3 focus only on images, with Phase 4 expanding to videos. We describe the key steps here; details and metrics are in §E.

Data Engine Components (Fig. 4). Media inputs (image or video) are mined from a large pool with the help of a curated ontology. An AI model proposes noun phrases (NPs) describing visual concepts, followed by another model (e.g., SAM 3) that generates candidate instance masks for each proposed NP. The proposed masks are verified by a two-step process: first, in *Mask Verification (MV)* annotators accept or reject masks based on their quality and relevance to the NP. Second, in *Exhaustivity Verification (EV)* annotators check if all instances of the NP have been masked in the input. Any media-NP pairs that did not pass the exhaustivity check are sent to a manual correction stage, where humans add, remove or edit masks (using SAM 1 in a browser based tool), or use “group” masks for small, hard to separate objects. Annotators may reject ungroundable or ambiguous phrases.

Phase 1: Human Verification. At first, data mining is done by randomly sampling images and NP proposal is done with a simple captioner and parser. The initial mask proposal model is SAM 2 prompted with the output of an off-the-shelf open-vocabulary detector, and verifiers are all human. The collected 4.3M image-NP pairs form the initial SA-Co/HQ dataset. We train SAM 3 on this data and use it as the mask proposal model for the next phase.

Phase 2: Human + AI Verification. In this next phase, we use human accept/reject labels from the MV and EV tasks collected in Phase 1 to fine-tune Llama 3.2 (Dubey et al., 2024) to create AI verifiers to automatically perform the MV and EV tasks. These models receive image-phrase-mask triplets and produce multiple-choice judgements of mask quality or exhaustivity. This new auto-verification process allows human effort to be focused on the most challenging cases. We continue to re-train SAM 3 on newly collected data and update it 6 times. As SAM 3 and AI verifiers improve, a higher proportion of labels are auto-generated, further accelerating data collection. The introduction of AI verifiers for MV and EV doubles the data engine’s throughput. We further upgrade the NP proposal step to a Llama-based pipeline that also proposes hard negative NPs adversarial to SAM 3. This phase adds 122M image-NP pairs to SA-Co/HQ.

Phase 3: Scaling and Domain Expansion. In this phase, we use AI models to mine increasingly challenging cases and broaden domain coverage in SA-Co/HQ to 15 datasets (Fig. 11). In new domains, the MV AI verifier performs well zero-shot, while the EV AI verifier improves with modest domain-specific human supervision. We also expand concept coverage to long-tail, fine-grained concepts by extracting NPs from each image’s alt-text where available and by mining concepts from a 22.4M node *SA-Co ontology* based on Wikidata (17 top-level categories, 72 sub-categories). We iterate SAM 3 training 7 times and AI verifiers 3 times, and add 19.5M image-NP pairs to SA-Co/HQ.



Figure 5: Example video (top) and images (bottom) from SA-Co with annotated phrases and instance masks/IDs.

Phase 4: Video Annotation. This phase extends the data engine to video. We use a mature SAM 3 to collect targeted quality annotations that capture video-specific challenges. The data mining pipeline applies scene/motion filters, content balancing, ranking, and targeted searches. Video frames are sampled (randomly or by object density) and sent to the image annotation flow (from phase 3). *Masklets* are produced with SAM 3 (now extended to video) and post-processed via deduplication and removal of trivial masks. Because video annotation is more difficult, we concentrate humans on likely failures by favoring clips with many crowded objects and tracking failures. The collected video data SA-Co/VIDEO consists of 52.5K videos and 467K masklets. See §E.6 for details.

5 SEGMENT ANYTHING WITH CONCEPTS (SA-CO) DATASET

Training Data. We collect three *image datasets* for the PCS task: i) SA-Co/HQ, the high-quality image data collected from the data engine in phases 1-4, ii) SA-Co/SYN, a synthetic dataset of images labeled by a mature version of the data engine (phase 3) without any human involvement, and iii) SA-Co/EXT, fifteen external datasets which have instance mask annotations, enriched with hard negatives using our SA-Co ontology pipeline. Notably in the SA-Co/HQ dataset we annotate 5.2M images and 4M unique NPs, making it the largest high-quality open-vocab segmentation dataset. We also annotate a *video dataset*, SA-Co/VIDEO, containing 52.5K videos and 24.8K unique NPs, forming 134K video-NP pairs. The videos have an average length of 84.1 frames at 6 fps. See §F.1 for details including full statistics, comparison with existing datasets and the distribution of concepts.

SA-Co Benchmark. The SA-Co evaluation benchmark has all together 214K unique phrases, 126K images and videos, and over 3M media-phrase pairs with challenging hard negative labels to test open-vocabulary recognition. It consists of several splits: SA-Co/Gold has seven domains and each image-NP pair is annotated by three different annotators (used measure human performance); SA-Co/Silver has ten domains and only one human annotation per image-NP pair; SA-Co/Bronze and SA-Co/Bio are nine existing datasets either with existing mask annotations or masks generated by using boxes as prompts to SAM 2. The SA-Co/VEval benchmark has three domains and one annotator per video-NP pair. See Tab. 29 for dataset statistics and Fig. 5 for example annotations.

Metrics. We aim to measure the usefulness of the model in downstream applications. Detection metrics such as average precision (AP) do not account for calibration, which means that models can be difficult to use in practice. To remedy this, we only evaluate predictions with confidence above 0.5, effectively introducing a threshold that mimics downstream usages and enforces good calibration. The PCS task can be naturally split into two sub-tasks, *localization* and *classification*. We evaluate localization using *positive macro F1* (pmF_1) on positive media-phrase pairs with at least one ground-truth and one predicted mask. Classification is measured with *image-level Matthews Correlation Coefficient* (IL-MCC) which ranges in $[-1, 1]$ and evaluates binary prediction at the image level (“is the object present or not?”) without regard for the quality of the masks. Our main metric, *classification-gated F1* (CGF_1), combines these as follows: $\text{CGF}_1 = 100 * \text{pmF}_1 * \text{IL-MCC}$. Full definitions are in §F.3.

Model	Instance Segmentation					Box Detection					Semantic Segmentation						
	LVIS		SA-Co			LVIS		COCO		SA-Co			ADE-847		PC-59	Cityscapes	
	CGF ₁	AP	Gold	Silver	Bronze	Bio	CGF ₁	AP	AP	AP _o	Gold	Silver	Bronze	Bio	mIoU	mIoU	mIoU
Human _{min}	—	—	74.2	—	—	—	—	—	—	76.2	—	—	—	—	—	—	—
Human _{max}	—	—	81.4	—	—	—	—	—	—	83.6	—	—	—	—	—	—	—
OWLv2	35.5	—	22.9	12.1	8.5	1.1	35.2	35.2	38.2	42.4	23.3	11.7	8.5	1.6	—	—	—
OWLv2*	45.7	43.4	34.3	19.3	19.3	0.2	47.4	45.5	46.1	23.9	35.4	19.2	19.4	0.3	—	—	—
gDino-T	32.9	—	9.1	7.4	11.1	0.6	33.8	20.5	45.7	35.3	9.4	6.8	12.3	0.7	—	—	—
LLMDet-L	48.1	36.3	12.9	12.1	18.8	0.5	53.7	42.0	55.6	49.8	13.5	11.5	20.6	0.6	—	—	—
APE-D*	—	53.0†	27.3	15.0	19.7	0.0	—	59.6†	58.3†	—	29.4	15.9	21.8	0.0	9.2†	58.5†	44.2†
DINO-X	—	38.5†	27.7δ	—	—	—	—	52.4†	56.0†	—	29.4δ	—	—	—	—	—	—
Gemini 2.5	19.8	—	16.4	10.5	8.9	10.2	23.7	—	—	—	18.7	12.2	10.1	12.0	—	—	—
SAM 3	52.8	47.0	65.0	57.1	49.5	59.3	57.5	51.7	53.5	55.5	67.7	58.0	53.1	60.0	14.7	59.4	65.1

Table 1: Evaluation on image concept segmentation with text. AP_o corresponds to COCO-O accuracy, * partially trained on LVIS. † from original papers, δ from DINO-X API. Gray numbers indicate usage of respective closed set training data (LVIS/COCO). Upper and lower bound for human performance given, see §F.4 for details.

Handling Ambiguity. We collect 3 annotations per NP on SA-Co/Gold. We measure *oracle* accuracy comparing each prediction to all ground truths and selecting the best score. See §F.3.

6 EXPERIMENTS

We evaluate SAM 3 across image/video segmentation, few-shot adaptation to detection and counting benchmarks, and segmenting from complex language queries with SAM 3 + MLLM. We also show a subset of ablations, with more in §B. References, more results and details are in §G.

Image PCS with Text. We evaluate instance segmentation, box detection, and semantic segmentation on external and our benchmarks. SAM 3 is prompted with a single NP at a time, and predicts instance masks, bounding boxes, or semantic masks. As baselines, we evaluate OWLv2, GroundingDino, and LLMDet on box detection, and prompt SAM 1 with their boxes to evaluate segmentation. We also compare to APE, DINO-X, and Gemini 2.5 Flash, a generalist LLM. Tab. 1 shows that zero-shot, SAM 3 is competitive on closed-vocabulary COCO, COCO-O and on LVIS boxes, and is significantly better on LVIS masks. On open-vocabulary SA-Co/Gold SAM 3 achieves *double* the CGF₁ score of the strongest baseline OWLv2*, and 88% of the estimated lower bound on human performance. The improvements are even higher on the other SA-Co splits. Open vocabulary semantic segmentation results on ADE-847, PascalConcept-59, and Cityscapes show that SAM 3 outperforms APE, a strong specialist baseline. See §G.1 for details.

Model	ODinW13		RF-100VL		COCO				LVIS				ODinW13				
	AP ₀	AP ₁₀	AP ₀	AP ₁₀	AP	AP ⁺	AP ⁺	AP ⁺	AP	AP ⁺	AP ⁺	AP ⁺	AP	AP ⁺	AP ⁺	AP ⁺	
Gemini2.5-Pro	33.7	—	11.6	9.8	T-Rex2	52.2	—	58.5	—	45.8	—	65.8	—	50.3	—	61.8	—
gDino-T	49.7	—	15.7	33.7	SAM 3	53.5	56.8	75.7	76.0	51.7	53.4	75.5	77.0	59.9	62.5	81.9	79.6
gDino1.5-Pro	58.7	67.9	—	—													
SAM 3	59.9	71.6	14.3	35.7													

Table 2: Zero-shot and 10-shot transfer on in-the-wild datasets. Evaluation per prompt type: T (text-only), I (image-only), and T+I (combined text and image). AP⁺ is evaluated only on positives examples.

Few-Shot Adaptation. We evaluate zero- and few-shot transfer of SAM 3 on ODinW13 and Roboflow100-VL, with their original labels as prompts. We *do not* perform any prompt tuning. We fine-tune SAM 3 without mask loss, and report average bbox mAP in Tab. 2. SAM 3 achieves SoTA 10-shot performance, surpassing in-context prompting in Gemini and object detection experts (gDino); more details in §G.3. RF-100VL contains domains with specialized prompts that are out of SAM 3’s current scope, but SAM 3 adapts through fine-tuning more efficiently than baselines.

PCS with 1 Exemplar. We first evaluate image exemplars using a single input box sampled at random from the ground truth. This can be done only on “positive” data, where each prompted object appears in the image. We report the corresponding AP⁺ in Tab. 3 across three settings: text prompt (T), exemplar image (I), and both text and image (T+I); SAM 3 outperforms prior state-of-the-art T-Rex2 by a healthy margin on COCO (+17.2), LVIS (+9.7), and ODinW (+20.1). See §G.2 for more details and results on SA-Co/Gold.

PCS with K Exemplars. Next, we evaluate SAM 3 in an interactive setting, simulating collaboration with a human annotator. Starting with a text prompt, we iteratively add one exemplar prompt at a time: missed ground truths are candidate positive prompts, false positive detections are candidate negative prompts. Results (Fig. 6) are compared to a perfect PVS baseline, where we simulate the user manually fixing errors using ideal box-to-mask corrections. SAM 3’s PCS improves CGF₁ more quickly, as it generalizes from exemplars (e.g., detecting or suppressing similar objects), while PVS only corrects individual instances. After 3 clicks, interactive PCS outperforms text-only by +18.6 CGF₁ points and PVS refinement by +9.7. Performance plateaus after 4 clicks, as exemplars cannot fix poor-quality masks. Simulating a *hybrid* switch to PVS at this point yields further gains, showing that the two approaches are *complementary*.

Object Counting. We evaluate on object counting benchmarks CountBench and PixMo-Count to compare with several MLLMs using Accuracy (%) and Mean Absolute Error (MAE) from previous technical reports and our own evaluations. See Tab. 4 for results and §G.4 for more evaluation details. Compared to MLLMs, SAM 3 not only achieves good object counting accuracy, but also provides object segmentation that most MLLMs cannot provide.

Video PCS with Text. We evaluate video segmentation with text prompts on both our SA-Co/VEval benchmark and existing public benchmarks. For SA-Co/VEval, we report CGF₁ and pHOTA metrics (defined in §G.5) across its subsets (SA-V, YT-Temporal-1B, SmartGlasses). For public benchmarks, we use their official metrics. Baselines include GLEE, an open-vocabulary image and video segmentation model, “LLMDet + SAM 3 Tracker” (replacing our detector with LLMDet), and “SAM 3 Detector + T-by-D” (replacing our tracker with an association module based on the tracking-by-detection paradigm). In Tab. 5, SAM 3 largely outperforms these baselines, especially on benchmarks with a very large number of noun phrases. On SA-Co/VEval it reaches over 80% of human pHOTA. See §G.5 for more details.

Model	SA-Co/VEval benchmark test split						public benchmarks				
	SA-V (2.0K NPs)		YT-Temporal-1B (1.7K NPs)		SmartGlasses (2.4K NPs)		LVVIS (1.2K NPs)		BURST (482 NPs)	YTVIS21 (40 NPs)	OVIS (25 NPs)
	CGF ₁	pHOTA	CGF ₁	pHOTA	CGF ₁	pHOTA	test mAP	test HOTA	val mAP	val mAP	
Human	53.2	68.0	73.8	79.2	57.6	70.3	–	–	–	–	
GLEE [†] (all NPs at once)	0.2	8.8	2.1	18.9	0.1	5.5	20.8	28.4	62.2	38.7	
GLEE [†] (one NP at a time)	0.0	12.3	3.2	22.6	0.1	6.4	9.3	20.2	56.5	32.4	
LLMDet [†] + SAM 3 Tracker	3.6	31.2	9.2	41.5	0.0	7.4	15.5	35.1	32.4	47.5	
SAM 3 Detector + T-by-D	22.2	49.0	44.6	64.4	29.3	57.1	37.3	40.7	57.3	54.9	
SAM 3	27.8	53.9	51.7	69.2	38.2	62.9	38.2	45.9	56.9	59.9	

Table 5: Video PCS from a text prompt (open-vocabulary video instance segmentation) on SA-Co/VEval and public benchmarks (see Tab. 38 for more results and analyses). SAM 3 shows strong performance, especially on benchmarks with a large number of NPs. [†]: GLEE and LLMDet have not been trained on the SA-Co dataset, so their results should be seen as zero-shot on SA-Co/VEval.

PVS. We evaluate SAM 3 on a range of geometric tasks, including Video Object Segmentation (VOS) and interactive image segmentation. Tab. 6 compares SAM 3 to recent state-of-the-art methods on the VOS task. SAM 3 achieves significant improvements over SAM 2 on most benchmarks, particularly on the challenging MOSEv2 dataset, where SAM 3 outperforms prior work by 6 points. For the interactive image segmentation task, we evaluate SAM 3 on the 37 datasets benchmark introduced in SAM 2. As shown in Tab. 7, SAM 3 outperforms SAM 2 on average mIoU.

SAM 3 Agent. We experiment with an MLLM that uses SAM 3 as a tool, called *SAM 3 Agent*, to segment more complex text queries such as “people sitting down but not holding a gift box in

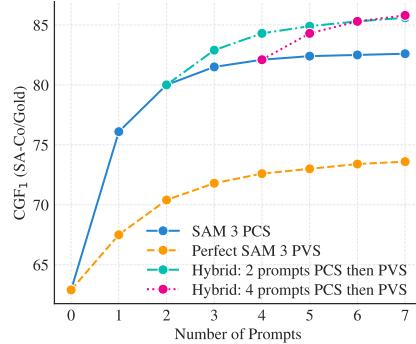


Figure 6: SAM 3’s interactive exemplar prompts vs the ideal PVS baseline on SA-Co. We report CGF₁ vs # of box prompts, averaged over all SA-Co/Gold phrases.

Model	CountBench		PixMo-Count	
	MAE ↓	Acc ↑	MAE ↓	Acc ↑
DINO-X	0.62	82.9	0.21	85.0
Qwen2-VL-72B	0.28	86.7	0.61	63.7
Molmo-72B	0.27	92.4	0.17	88.8
Gemini 2.5 Pro	0.24	92.4	0.38	78.2
SAM 3	0.11	95.6	0.22	87.3

Table 4: Accuracy on counting benchmarks. Gray indicates usage of training sets.

432	Model	$\mathcal{J} \& \mathcal{F}$					\mathcal{G} YTVOS19 val	$\mathcal{J} \& \dot{\mathcal{F}}$ MOSEv2 val	Model	Avg. mIoU	
		MOSEv1 val	DAVIS17 val	LVOSv2 val	SA-V val	SA-V test				1-click	5-clicks
433	SAMURAI	72.6	89.9	84.2	79.8	80.0	88.3	51.1	SAM 1 H	58.5	82.1
434	SAM2Long	75.2	91.4	85.9	81.1	81.2	88.7	51.5	SAM 2.1 L	64.8	84.4
435	SeC	75.3	91.3	86.5	82.7	81.7	88.6	53.8	SAM 3	65.4	85.0
436	SAM 2.1 L	77.9	90.7	79.6	77.9	78.4	89.3	47.9 [†]			
437	SAM 3	78.4	92.0	88.2	82.9	84.6	89.6	60.1			

Table 6: SAM 3 improves over SAM 2 in VOS. [†]: Zero-shot.

Table 7: Interactive image segmentation on the SA-37 benchmark.

439	Model	ReasonSeg (gIoU)				Omnilabel (AP)				
		val		test		val 2023		2023		
		All	<u>All</u>	Short	Long	descr	descr-S	descr-M	descr-L	
440	X-SAM	Phi-3-3.8B	56.6	57.8	47.7	56.0	12.0*	17.1*	11.4*	8.8*
441	SegZero	Qwen2.5-VL 7B	<u>62.6</u>	<u>57.5</u>	—	—	13.5*	20.7*	12.4*	9.1*
442	RSVP	GPT-4o	64.7	55.4	61.9	60.3	—	—	—	—
443	Overall SoTA Performance [†]		65	61.3	55.4	63.2	36.5	54.4	33.2	25.5
444	SAM 3 Agent	Qwen2.5-VL 7B	65.4	62.6	59.1	63.7	36.5	52.6	34.3	26.7
445	SAM 3 Agent	Llama4 Maverick	71.5	69.3	70.9	68.8	36.2	47.5	34.9	28.1
446	SAM 3 Agent	Qwen2.5-VL 72B	75.0	71.8	71.3	72.0	44.7	58.4	42.6	36.1
447	SAM 3 Agent	Gemini 2.5 Pro	76.0	73.8	74.0	73.7	46.7	54.6	46.2	38.7

Table 8: SAM 3 Agent results. Gray indicates fine-tuned results on ReasonSeg (train), * indicates reproduced results, underline indicates the main metric. [†]: LISA-13B-LLaVA1.5 for ReasonSeg; REAL for OmniLabel.

their hands”. The MLLM proposes noun phrase queries to prompt SAM 3 and analyzes the returned masks, iterating until the masks are satisfactory. Tab. 8 shows that SAM 3 Agent evaluated zero-shot on ReasonSeg (Lai et al., 2024) and OmniLabel (Schulter et al., 2023) surpasses prior work without training on any referring expression segmentation or reasoning segmentation data. SAM 3 Agent also outperforms previous zero-shot results on RefCOCO+ (Kazemzadeh et al., 2014) and RefCOCOg (Mao et al., 2016). SAM 3 can be combined with various MLLMs, with the same set of the system prompts for all those MLLMs, showing SAM 3’s robustness. See §H for more details.

CGF ₁	IL_MCC	pmF ₁	#/img	CGF ₁	IL_MCC	pmF ₁	EXT	SYN	HQ	CGF ₁	IL_MCC	pmF ₁	Model	CGF ₁	IL_MCC	pmF ₁	
×	57.6	0.77	74.7	0	31.8	0.44	70.2	✓	✗	✗	30.9	0.46	66.3	Human _{min}	74.2	0.88	84.7
✓	63.3	0.82	77.1	5	44.8	0.62	71.9	✓	✓	✗	39.7	0.57	70.6	SAM 3	65.0	0.82	79.7
			30	49.2	0.68	72.3	✓	✗	✓	51.8	0.71	73.2	+ EV AI	68.1	0.86	78.8	

(a) Presence token. (b) Hard Negatives. (c) Training data. (d) SAM 3 + AI verifiers.

Table 9: Selected model and data ablations on SA-Co/Gold. Numbers across tables are not directly comparable.

Selected Ablations. In Tab. 9 we report a subset of the more extensive ablations from §B. The presence head boosts CGF₁ by +5.7 (9a), mainly improving image-level recognition ability measured by IL_MCC. Tab. 9b shows that adding hard negatives significantly improves the model performance, most notably the image-level IL_MCC from 0.44 to 0.68. Tab. 9c shows that synthetic (SYN) training data improves over the external (EXT) by +8.8 CGF₁ and our high-quality (HQ) annotations add +14.6 CGF₁ on top of this baseline. We present detailed data scaling laws of the synthetic data and high-quality data in §B.2, showing their effectiveness on both in-domain and out-of-domain test sets. In Tab. 9d, we show how AI verifiers can improve pseudo-labels. Replacing the presence score from SAM 3 with a presence score from the exhaustivity verification (EV) AI verifier boosts IL_MCC by 4.6 points. Using the mask verification (MV) AI verifier to remove bad masks boosts pmF₁ by +2.5.

7 CONCLUSION

We present Segment Anything with *Concepts*, enabling open vocabulary text and exemplars as prompts in interactive segmentation. Our principal contributions are: (i) introducing the PCS task and SA-Co benchmark, (ii) a decoupled recognition-localization architecture that extends SAM 2 for PCS while retaining PVS capabilities, (iii) a high-quality efficient human and AI annotator in the loop data engine. SAM 3 achieves state-of-the-art results, doubling performance over prior systems for PCS on SA-Co in images and videos. We believe SAM 3 and the SA-Co benchmark will be important milestones and pave the way for future research and applications in computer vision.

486 REFERENCES
487

- 488 URL <https://lila.science/datasets/wcscameratraps>. 44, 49
489 URL <https://www.nga.gov/artworks/free-images-and-open-access>. 44, 48
490
- 491 Xiang An, Kaicheng Yang, Xiangzi Dai, Ziyong Feng, and Jiankang Deng. Multi-label cluster
492 discrimination for visual representation learning. In *European Conference on Computer Vision*, pp.
493 428–444. Springer, 2024. 68
- 494 Ali Athar, Jonathon Luiten, Paul Voigtlaender, Tarasha Khurana, Achal Dave, Bastian Leibe, and
495 Deva Ramanan. Burst: A benchmark for unifying object recognition, segmentation and tracking in
496 video. In *Proceedings of the IEEE/CVF winter conference on applications of computer vision*, pp.
497 1674–1683, 2023. 55
- 498
- 499 Hexin Bai, Wensheng Cheng, Peng Chu, Juehuan Liu, Kai Zhang, and Haibin Ling. Gmot-40: A
500 benchmark for generic multiple object tracking. In *Proceedings of the IEEE/CVF Conference on*
501 *Computer Vision and Pattern Recognition*, pp. 6719–6728, 2021. 55
- 502 Kevin Barnard, Elaine Liu, Kristine Walz, Brian Schlining, Nancy Jacobsen Stout, and Lonny
503 Lundsten. DeepSea MOT: A benchmark dataset for multi-object tracking on deep-sea video. *arXiv*
504 preprint [arXiv:2509.03499](https://arxiv.org/abs/2509.03499), 2025. doi: 10.48550/arXiv.2509.03499. 55
- 505
- 506 Philipp Bergmann, Tim Meinhardt, and Laura Leal-Taixe. Tracking without bells and whistles. In
507 *Proceedings of the IEEE/CVF international conference on computer vision*, pp. 941–951, 2019. 22
- 508 Alex Bewley, Zongyuan Ge, Lionel Ott, Fabio Ramos, and Ben Upcroft. Simple online and realtime
509 tracking. In *2016 IEEE international conference on image processing (ICIP)*, pp. 3464–3468. Ieee,
510 2016. 22
- 511
- 512 Lucas Beyer, Andreas Steiner, André Susano Pinto, Alexander Kolesnikov, Xiao Wang, Daniel
513 Salz, Maxim Neumann, Ibrahim Alabdulmohsin, Michael Tschannen, Emanuele Bugliarello, et al.
514 PaliGemma: A versatile 3B VLM for transfer. *arXiv preprint arXiv:2407.07726*, 2024. 54
- 515 Alexey Bochkovskiy, Chien-Yao Wang, and Hong-Yuan Mark Liao. Yolov4: Optimal speed and
516 accuracy of object detection, 2020. URL <https://arxiv.org/abs/2004.10934>. 34
- 517
- 518 Daniel Bolya, Chaitanya Ryali, Judy Hoffman, and Christoph Feichtenhofer. Window attention is
519 bugged: How not to interpolate position embeddings. In *International Conference on Learning*
520 *Representations*, 2024. 24, 28
- 521
- 522 Daniel Bolya, Po-Yao Huang, Peize Sun, Jang Hyun Cho, Andrea Madotto, Chen Wei, Tengyu Ma,
523 Jiale Zhi, Jathushan Rajasegaran, Hanoona Rasheed, Junke Wang, Marco Monteiro, Hu Xu, Shiyu
524 Dong, Nikhila Ravi, Daniel Li, Piotr Dollár, and Christoph Feichtenhofer. Perception encoder:
525 The best visual embeddings are not at the output of the network. *arXiv:2504.13181*, 2025. 2, 3, 23,
526 28, 29, 32, 33, 41
- 527
- 528 Zhi Cai, Songtao Liu, Guodong Wang, Zeming Li, Zheng Ge, Xiangyu Zhang, and Di Huang.
Align-detr: Enhancing end-to-end object detection with aligned loss. In *35th British Machine*
529 *Vision Conference 2024, BMVC 2024, Glasgow, UK, November 25-28, 2024*. BMVA, 2024. URL
530 <https://papers.bmvc2024.org/0211.pdf>. 3, 33
- 531
- 532 Jinkun Cao, Jiangmiao Pang, Xinshuo Weng, Rawal Khirodkar, and Kris Kitani. Observation-centric
533 sort: Rethinking sort for robust multi-object tracking. In *Proceedings of the IEEE/CVF conference*
on *computer vision and pattern recognition*, pp. 9686–9696, 2023. 22
- 534
- 535 Nicolas Carion, Francisco Massa, Gabriel Synnaeve, Nicolas Usunier, Alexander Kirillov, and Sergey
536 Zagoruyko. End-to-end object detection with transformers. In *European conference on computer*
537 *vision*, pp. 213–229. Springer, 2020. 2, 3, 22, 28, 29
- 538
- 539 Qiang Chen, Xiangbo Su, Xinyu Zhang, Jian Wang, Jiahui Chen, Yunpeng Shen, Chuchu Han,
Ziliang Chen, Weixiang Xu, Fanrong Li, et al. Lw-detr: A transformer replacement to yolo for
real-time detection. *arXiv preprint arXiv:2406.03459*, 2024a. 53

- 540 Yi-Chia Chen, Wei-Hua Li, Cheng Sun, Yu-Chiang Frank Wang, and Chu-Song Chen. Sam4mllm:
 541 Enhance multi-modal large language model for referring expression segmentation. In *European
 542 Conference on Computer Vision*, pp. 323–340. Springer, 2024b. 66, 68
- 543
- 544 Yuming Chen, Jiangyan Feng, Haodong Zhang, Lijun Gong, Feng Zhu, Rui Zhao, Qibin Hou, Ming-
 545 Ming Cheng, and Yibing Song. Re-aligning language to visual objects with an agentic workflow.
 546 In *International Conference on Learning Representations*, 2025. 67
- 547
- 548 Bowen Cheng, Alexander G. Schwing, and Alexander Kirillov. Per-pixel classification is not all you
 549 need for semantic segmentation. 2021. 3, 30
- 550 Jang Hyun Cho, Andrea Madotto, Effrosyni Mavroudi, Triantafyllos Afouras, Tushar Nagarajan,
 551 Muhammad Maaz, Yale Song, Tengyu Ma, Shuming Hu, Hanoona Rasheed, Peize Sun, Po-
 552 Yao Huang, Daniel Bolya, Suyog Jain, Miguel Martin, Huiyu Wang, Nikhila Ravi, Shashank
 553 Jain, Temmy Stark, Shane Moon, Babak Damavandi, Vivian Lee, Andrew Westbury, Salman
 554 Khan, Philipp Krähenbühl, Piotr Dollár, Lorenzo Torresani, Kristen Grauman, and Christoph
 555 Feichtenhofer. Perceptionlm: Open-access data and models for detailed visual understanding.
 556 *arXiv:2504.13180*, 2025. 41
- 557
- 558 Yung-Sung Chuang, Yang Li, Dong Wang, Ching-Feng Yeh, Kehan Lyu, Ramya Raghavendra, James
 559 Glass, Lifei Huang, Jason Weston, Luke Zettlemoyer, Xinlei Chen, Zhuang Liu, Saining Xie,
 560 Wen tau Yih, Shang-Wen Li, and Hu Xu. Meta clip 2: A worldwide scaling recipe, 2025. URL
<https://arxiv.org/abs/2507.22062>. 48
- 561
- 562 Kevin Clark, Minh-Thang Luong, Quoc V Le, and Christopher D Manning. ELECTRA: Pre-training
 563 text encoders as discriminators rather than generators. In *ICLR*, 2020. 33
- 564
- 565 Gheorghe Comanici, Eric Bieber, Mike Schaeckermann, Ice Pasupat, Noveen Sachdeva, Inderjit
 566 Dhillon, Marcel Blistein, Ori Ram, Dan Zhang, Evan Rosen, et al. Gemini 2.5: Pushing the frontier
 567 with advanced reasoning, multimodality, long context, and next generation agentic capabilities.
 568 *arXiv preprint arXiv:2507.06261*, 2025. 22, 52, 53, 54
- 569
- 570 Marius Cordts, Mohamed Omran, Sebastian Ramos, Timo Rehfeld, Markus Enzweiler, Rodrigo
 571 Benenson, Uwe Franke, Stefan Roth, and Bernt Schiele. The cityscapes dataset for semantic
 572 urban scene understanding. In *Proc. of the IEEE Conference on Computer Vision and Pattern
 573 Recognition (CVPR)*, 2016. 52
- 574
- 575 Achal Dave, Piotr Dollár, Deva Ramanan, Alexander Kirillov, and Ross Girshick. Evaluating large-
 576 vocabulary object detectors: The devil is in the details, 2022. URL <https://arxiv.org/abs/2102.01066>. 52
- 577
- 578 Matt Deitke, Christopher Clark, Sangho Lee, Rohun Tripathi, Yue Yang, Jae Sung Park, Moham-
 579 madreza Salehi, Niklas Muennighoff, Kyle Lo, Luca Soldaini, et al. Molmo and pixmo: Open
 580 weights and open data for state-of-the-art vision-language models. In *Proceedings of the Computer
 581 Vision and Pattern Recognition Conference*, pp. 91–104, 2025. 22, 54
- 582
- 583 Henghui Ding, Kaining Ying, Chang Liu, Shuting He, Xudong Jiang, Yu-Gang Jiang, Philip HS Torr,
 584 and Song Bai. Mosev2: A more challenging dataset for video object segmentation in complex
 585 scenes. *arXiv preprint arXiv:2508.05630*, 2025. 44, 56
- 586
- 587 Kexin Ding, Mu Zhou, He Wang, Olivier Gevaert, Dimitris Metaxas, and Shaoting Zhang. A
 588 large-scale synthetic pathological dataset for deep learning-enabled segmentation of breast cancer.
 589 *Scientific Data*, 10(1):231, 2023. 44, 49
- 590
- 591 Shuangrui Ding, Rui Qian, Xiaoyi Dong, Pan Zhang, Yuhang Zang, Yuhang Cao, Yuwei Guo, Dahua
 592 Lin, and Jiaqi Wang. Sam2long: Enhancing sam 2 for long video segmentation with a training-free
 593 memory tree. *arXiv preprint arXiv:2410.16268*, 2024. 56
- 594
- 595 Zheng Ding, Jieke Wang, and Zhuowen Tu. Open-vocabulary universal image segmentation with
 596 maskclip. *arXiv preprint arXiv:2208.08984*, 2022. 22

- 594 Zi-Yi Dou, Aishwarya Kamath, Zhe Gan, Pengchuan Zhang, Jianfeng Wang, Linjie Li, Zicheng
 595 Liu, Ce Liu, Yann LeCun, Nanyun Peng, Jianfeng Gao, and Lijuan Wang. Coarse-to-fine vision-
 596 language pre-training with fusion in the backbone, 2022. URL <https://arxiv.org/abs/2206.07643>. 35, 67
- 598 Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha
 599 Letman, Akhil Mathur, Alan Schelten, Amy Yang, Angela Fan, et al. The llama 3 herd of models.
 600 *arXiv e-prints*, pp. arXiv–2407, 2024. 5, 40
- 602 Christoffer Edlund, Timothy R Jackson, Nabeel Khalid, Nicola Bevan, Timothy Dale, Andreas
 603 Dengel, Sheraz Ahmed, Johan Trygg, and Rickard Sjögren. Livecell—a large-scale dataset for
 604 label-free live cell segmentation. *Nature methods*, 18(9):1038–1045, 2021. 44, 49
- 606 FFmpeg developers. FFmpeg. <https://ffmpeg.org/>. 41
- 607 Shenghao Fu, Qize Yang, Qijie Mo, Junkai Yan, Xihan Wei, Jingke Meng, Xiaohua Xie, and Wei-Shi
 608 Zheng. Llmdet: Learning strong open-vocabulary object detectors under the supervision of large
 609 language models. In *Proceedings of the IEEE/CVF conference on computer vision and pattern*
 610 *recognition*, 2025. 52, 53, 56
- 612 Jevgenij Gamper, Navid Alemi Koohbanani, Ksenija Benes, Ali Khuram, and Nasir Rajpoot. Pannuke:
 613 an open pan-cancer histology dataset for nuclei instance segmentation and classification. In
 614 *European Congress on Digital Pathology*, pp. 11–19. Springer, 2019. 44, 49
- 616 Jevgenij Gamper, Navid Alemi Koohbanani, Simon Graham, Mostafa Jahanifar, Syed Ali Khurram,
 617 Ayesha Azam, Katherine Hewitt, and Nasir Rajpoot. Pannuke dataset extension, insights and
 618 baselines. *arXiv preprint arXiv:2003.10778*, 2020. 44
- 619 Ruopeng Gao, Ji Qi, and Limin Wang. Multiple object tracking as id prediction. In *Proceedings of*
 620 *the Computer Vision and Pattern Recognition Conference*, pp. 27883–27893, 2025. 22
- 622 Kristen Grauman, Andrew Westbury, Eugene Byrne, Zachary Chavis, Antonino Furnari, Rohit Gird-
 623 har, Jackson Hamburger, Hao Jiang, Miao Liu, Xingyu Liu, Miguel Martin, Tushar Nagarajan,
 624 Ilija Radosavovic, Santhosh Kumar Ramakrishnan, Fiona Ryan, Jayant Sharma, Michael Wray,
 625 Mengmeng Xu, Eric Zhongcong Xu, Chen Zhao, Siddhant Bansal, Dhruv Batra, Vincent Cart-
 626 tillier, Sean Crane, Tien Do, Morrie Doulaty, Akshay Erapalli, Christoph Feichtenhofer, Adriano
 627 Fragomeni, Qichen Fu, Christian Fuegen, Abrham Gebreselasie, Cristina Gonzalez, James Hillis,
 628 Xuhua Huang, Yifei Huang, Wenqi Jia, Leslie Khoo, Jachym Kolar, Satwik Kottur, Anurag
 629 Kumar, Federico Landini, Chao Li, Yanghao Li, Zhenqiang Li, Karttikeya Mangalam, Raghava
 630 Modhugu, Jonathan Munro, Tullie Murrell, Takumi Nishiyasu, Will Price, Paola Ruiz Puentes,
 631 Merey Ramazanova, Leda Sari, Kiran Somasundaram, Audrey Southerland, Yusuke Sugano, Ruijie
 632 Tao, Minh Vo, Yuchen Wang, Xindi Wu, Takuma Yagi, Yunyi Zhu, Pablo Arbelaez, David Crandall,
 633 Dima Damen, Giovanni Maria Farinella, Bernard Ghanem, Vamsi Krishna Ithapu, C. V. Jawahar,
 634 Hanbyul Joo, Kris Kitani, Haizhou Li, Richard Newcombe, Aude Oliva, Hyun Soo Park, James M.
 635 Rehg, Yoichi Sato, Jianbo Shi, Mike Zheng Shou, Antonio Torralba, Lorenzo Torresani, Mingfei
 636 Yan, and Jitendra Malik. Ego4d: Around the World in 3,000 Hours of Egocentric Video. In
 637 *IEEE/CVF Computer Vision and Pattern Recognition (CVPR)*, 2022. 44, 48
- 638 Xiuye Gu, Tsung-Yi Lin, Weicheng Kuo, and Yin Cui. Open-vocabulary object detection via vision
 639 and language knowledge distillation. *arXiv preprint arXiv:2104.13921*, 2021. 22
- 640 Agrim Gupta, Piotr Dollar, and Ross Girshick. Lvis: A dataset for large vocabulary instance
 641 segmentation. In *Proceedings of the IEEE/CVF conference on computer vision and pattern*
 642 *recognition*, pp. 5356–5364, 2019. 3, 52
- 643 Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked
 644 autoencoders are scalable vision learners. In *Proceedings of the IEEE/CVF conference on computer*
 645 *vision and pattern recognition*, pp. 16000–16009, 2022. 23
- 646 Byeongho Heo, Song Park, Dongyo Han, and Sangdoo Yun. Rotary position embedding for vision
 647 transformer. *arXiv preprint arXiv:2403.13298*, 2024. 24

- 648 Matthew Honnibal, Ines Montani, Sofie Van Landeghem, and Adriane Boyd. spaCy: Industrial-
 649 strength Natural Language Processing in Python. 2020. doi: 10.5281/zenodo.1212303. 35
 650
- 651 Grant Van Horn, Oisin Mac Aodha, Yang Song, Alexander Shepard, Hartwig Adam, Pietro Perona,
 652 and Serge J. Belongie. The inaturalist challenge 2017 dataset. *CoRR*, abs/1707.06642, 2017. URL
 653 <http://arxiv.org/abs/1707.06642>. 44, 49
- 654 Zhengdong Hu, Yifan Sun, Jingdong Wang, and Yi Yang. DAC-DETR: Divide the attention layers
 655 and conquer. In *Advances in Neural Information Processing Systems*, 2023. 3, 29
 656
- 657 Jiaqi Huang, Zunnan Xu, Ting Liu, Yong Liu, Haonan Han, Kehong Yuan, and Xiu Li. Densely
 658 connected parameter-efficient tuning for referring image segmentation. In *Proceedings of the AAAI
 Conference on Artificial Intelligence*, volume 39, pp. 3653–3661, 2025. 68
 659
- 660 Ding Jia, Yuhui Yuan, Haodi He, Xiaopei Wu, Haojun Yu, Weihong Lin, Lei Sun, Chao Zhang, and
 661 Han Hu. Detrs with hybrid matching. *arXiv preprint arXiv:2207.13080*, 2022. 29
 662
- 663 Menglin Jia, Mengyun Shi, Mikhail Sirotenko, Yin Cui, Claire Cardie, Bharath Hariharan, Hartwig
 664 Adam, and Serge J. Belongie. Fashionpedia: Ontology, segmentation, and an attribute localization
 665 dataset. *CoRR*, abs/2004.12276, 2020. URL <https://arxiv.org/abs/2004.12276>. 44
 666
- Junjie Jiang, Zelin Wang, Manqi Zhao, Yin Li, and DongSheng Jiang. Sam2mot: A novel paradigm
 667 of multi-object tracking by segmentation. *arXiv preprint arXiv:2504.04519*, 2025. 22
 668
- Qing Jiang, Feng Li, Zhaoyang Zeng, Tianhe Ren, Shilong Liu, and Lei Zhang. T-rex2: Towards
 669 generic object detection via text-visual prompt synergy. In *European Conference on Computer
 670 Vision*, pp. 38–57. Springer, 2024. 22, 53
 671
- Arne Hoffhues Jonathon Luiten. Trackeval. [https://github.com/JonathonLuiten/
 672 TrackEval](https://github.com/JonathonLuiten/TrackEval), 2020. 55
 673
- Aishwarya Kamath, Mannat Singh, Yann LeCun, Gabriel Synnaeve, Ishan Misra, and Nicolas Carion.
 674 Mdetr-modulated detection for end-to-end multi-modal understanding. In *Proceedings of the
 675 IEEE/CVF international conference on computer vision*, pp. 1780–1790, 2021. 3, 22, 29
 676
- Seil Kang, Jinyeong Kim, Junhyeok Kim, and Seong Jae Hwang. Your large vision-language model
 677 only needs a few attention heads for visual grounding. In *Proceedings of the Computer Vision and
 678 Pattern Recognition Conference*, pp. 9339–9350, 2025. 66, 68
 679
- Kakani Katija, Eric C. Orenstein, Brian Schlining, Lonny Lundsten, Kevin Barnard, Giovanna Sainz,
 680 Oceane Boulais, Benjamin G. Woodward, and Katy Croff Bell. Fathomnet: A global underwater
 681 image training set for enabling artificial intelligence in the ocean. *CoRR*, abs/2109.14646, 2021.
 682 URL <https://arxiv.org/abs/2109.14646>. 44, 49
 683
- Sahar Kazemzadeh, Vicente Ordonez, Mark Matten, and Tamara Berg. Referitgame: Referring to
 684 objects in photographs of natural scenes. In *Proceedings of the 2014 conference on empirical
 685 methods in natural language processing (EMNLP)*, pp. 787–798, 2014. 9, 61, 62, 66, 68
 686
- Lei Ke, Henghui Ding, Martin Danelljan, Yu-Wing Tai, Chi-Keung Tang, and Fisher Yu. Video mask
 687 transfiner for high-quality video instance segmentation. In *European Conference on Computer
 688 Vision*, pp. 731–747. Springer, 2022. 55
 689
- Alexander Khazatsky, Karl Pertsch, Suraj Nair, Ashwin Balakrishna, Sudeep Dasari, Siddharth
 690 Karamcheti, Soroush Nasiriany, Mohan Kumar Srirama, Lawrence Yunliang Chen, Kirsty Ellis,
 691 Peter David Fagan, Joey Hejna, Masha Itkina, Marion Lepert, Yecheng Jason Ma, Patrick Tree
 692 Miller, Jimmy Wu, Suneel Belkhale, Shivin Dass, Huy Ha, Arhan Jain, Abraham Lee, Youngwoon
 693 Lee, Marius Memmel, Sungjae Park, Ilija Radosavovic, Kaiyuan Wang, Albert Zhan, Kevin Black,
 694 Cheng Chi, Kyle Beltran Hatch, Shan Lin, Jingpei Lu, Jean Mercat, Abdul Rehman, Pannag R
 695 Sanketi, Archit Sharma, Cody Simpson, Quan Vuong, Homer Rich Walke, Blake Wulfe, Ted
 696 Xiao, Jonathan Heewon Yang, Arefeh Yavary, Tony Z. Zhao, Christopher Agia, Rohan Baijal,
 697 Mateo Guaman Castro, Daphne Chen, Qiuyu Chen, Trinity Chung, Jaimyn Drake, Ethan Paul
 698 Foster, Jensen Gao, Vitor Guizilini, David Antonio Herrera, Minho Heo, Kyle Hsu, Jiaheng
 699 Hu, Muhammad Zubair Irshad, Donovon Jackson, Charlotte Le, Yunshuang Li, Kevin Lin, Roy
 700
- 701

- 702 Lin, Zehan Ma, Abhiram Maddukuri, Suvir Mirchandani, Daniel Morton, Tony Nguyen, Abigail
 703 O'Neill, Rosario Scalise, Derick Seale, Victor Son, Stephen Tian, Emi Tran, Andrew E. Wang,
 704 Yilin Wu, Annie Xie, Jingyun Yang, Patrick Yin, Yunchu Zhang, Osbert Bastani, Glen Berseth,
 705 Jeannette Bohg, Ken Goldberg, Abhinav Gupta, Abhishek Gupta, Dinesh Jayaraman, Joseph J
 706 Lim, Jitendra Malik, Roberto Martín-Martín, Subramanian Ramamoorthy, Dorsa Sadigh, Shuran
 707 Song, Jiajun Wu, Michael C. Yip, Yuke Zhu, Thomas Kollar, Sergey Levine, and Chelsea Finn.
 708 Droid: A large-scale in-the-wild robot manipulation dataset. 2024. 44, 48
- 709 Rawal Khirodkar, Timur Bagautdinov, Julieta Martinez, Su Zhaoen, Austin James, Peter Selednik,
 710 Stuart Anderson, and Shunsuke Saito. Sapiens: Foundation for human vision models, 2024. URL
 711 <https://arxiv.org/abs/2408.12569>. 44
- 712 Alexander Kirillov, Eric Mintun, Nikhila Ravi, Hanzi Mao, Chloe Rolland, Laura Gustafson, Tete
 713 Xiao, Spencer Whitehead, Alexander C Berg, Wan-Yen Lo, et al. Segment anything. In *Proceedings*
 714 *of the IEEE/CVF international conference on computer vision*, pp. 4015–4026, 2023. 1, 22, 28, 44,
 715 48
- 716 Ranjay Krishna, Yuke Zhu, Oliver Groth, Justin Johnson, Kenji Hata, Joshua Kravitz, Stephanie
 717 Chen, Yannis Kalantidis, Li-Jia Li, David A Shamma, et al. Visual genome: Connecting language
 718 and vision using crowdsourced dense image annotations. *International journal of computer vision*,
 719 123(1):32–73, 2017. 44
- 720 Alina Kuznetsova, Hassan Rom, Neil Alldrin, Jasper Uijlings, Ivan Krasin, Jordi Pont-Tuset, Shahab
 721 Kamali, Stefan Popov, Matteo Malloci, Alexander Kolesnikov, et al. The open images dataset
 722 v4: Unified image classification, object detection, and visual relationship detection at scale.
 723 *International journal of computer vision*, 128(7):1956–1981, 2020. 44
- 724 Xin Lai, Zhuotao Tian, Yukang Chen, Yanwei Li, Yuhui Yuan, Shu Liu, and Jiaya Jia. Lisa: Reasoning
 725 segmentation via large language model. In *Proceedings of the IEEE/CVF Conference on Computer
 726 Vision and Pattern Recognition*, pp. 9579–9589, 2024. 9, 22, 61, 62, 66, 68
- 727 Hoang-An Le, Partha Das, Thomas Mensink, Sezer Karaoglu, and Theo Gevers. EDEN: Multimodal
 728 Synthetic Dataset of Enclosed garDEN Scenes. In *Proceedings of the IEEE/CVF Winter Conference
 729 of Applications on Computer Vision (WACV)*, 2021. 44, 49
- 730 Chunyuan Li, Haotian Liu, Liunian Li, Pengchuan Zhang, Jyoti Aneja, Jianwei Yang, Ping Jin,
 731 Houdong Hu, Zicheng Liu, Yong Jae Lee, et al. Elevater: A benchmark and toolkit for evaluating
 732 language-augmented visual models. *Advances in Neural Information Processing Systems*, 35:
 733 9287–9301, 2022a. 53
- 734 Feng Li, Qing Jiang, Hao Zhang, Tianhe Ren, Shilong Liu, Xueyan Zou, Hu-Sheng Xu, Hongyang
 735 Li, Chun Yue Li, Jianwei Yang, Lei Zhang, and Jianfeng Gao. Visual in-context prompting. 2024
 736 *IEEE/CVF Conference on Computer Vision and Pattern Recognition (CVPR)*, pp. 12861–12871,
 737 2023a. URL <https://api.semanticscholar.org/CorpusID:265351501>. 22
- 738 Jiachen Li, Qing Xie, Renshu Gu, Jinyu Xu, Yongjian Liu, and Xiaohan Yu. Lgd: Leveraging genera-
 739 tive descriptions for zero-shot referring image segmentation. *arXiv preprint arXiv:2504.14467*,
 740 2025. 68
- 741 Junnan Li, Dongxu Li, Silvio Savarese, and Steven Hoi. Blip-2: Bootstrapping language-image
 742 pre-training with frozen image encoders and large language models, 2023b. URL <https://arxiv.org/abs/2301.12597>. 35
- 743 Liunian Li, Zi-Yi Dou, Nanyun Peng, and Kai-Wei Chang. Desco: Learning object recognition with
 744 rich language descriptions. *Advances in Neural Information Processing Systems*, 36:37511–37526,
 745 2023c. 67
- 746 Liunian Harold Li, Pengchuan Zhang, Haotian Zhang, Jianwei Yang, Chunyuan Li, Yiwu Zhong,
 747 Lijuan Wang, Lu Yuan, Lei Zhang, Jenq-Neng Hwang, et al. Grounded language-image pre-
 748 training. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*,
 749 pp. 10965–10975, 2022b. 22, 67

- 756 Siyuan Li, Martin Danelljan, Henghui Ding, Thomas E Huang, and Fisher Yu. Tracking every thing
 757 in the wild. In *European Conference on Computer Vision*, 2022c. 56
 758
- 759 Yanghao Li, Hanzi Mao, Ross Girshick, and Kaiming He. Exploring plain vision transformer
 760 backbones for object detection. In *European Conference on Computer Vision*, pp. 280–296.
 761 Springer, 2022d. 28, 30
- 762 Feng Liang, Bichen Wu, Xiaoliang Dai, Kunpeng Li, Yinan Zhao, Hang Zhang, Peizhao Zhang,
 763 Peter Vajda, and Diana Marculescu. Open-vocabulary semantic segmentation with mask-adapted
 764 clip. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp.
 765 7061–7070, 2023. 22, 66
- 766 Tsung-Yi Lin, Michael Maire, Serge Belongie, James Hays, Pietro Perona, Deva Ramanan, Piotr
 767 Dollár, and C Lawrence Zitnick. Microsoft coco: Common objects in context. In *European
 768 conference on computer vision*, pp. 740–755. Springer, 2014. 52
- 770 Yutong Lin, Yuhui Yuan, Zheng Zhang, Chen Li, Nanning Zheng, and Han Hu. Detr doesn't need
 771 multi-scale or locality design. In *Proceedings of the IEEE/CVF International Conference on
 772 Computer Vision*, pp. 6545–6554, 2023. 3, 29
- 773 Shilong Liu, Zhaoyang Zeng, Tianhe Ren, Feng Li, Hao Zhang, Jie Yang, Chun yue Li, Jianwei
 774 Yang, Hang Su, Jun-Juan Zhu, and Lei Zhang. Grounding dino: Marrying dino with grounded
 775 pre-training for open-set object detection. In *European Conference on Computer Vision*, 2023.
 776 URL <https://api.semanticscholar.org/CorpusID:257427307>. 22, 53
- 778 Shilong Liu, Zhaoyang Zeng, Tianhe Ren, Feng Li, Hao Zhang, Jie Yang, Qing Jiang, Chunyuan
 779 Li, Jianwei Yang, Hang Su, et al. Grounding dino: Marrying dino with grounded pre-training for
 780 open-set object detection. In *European conference on computer vision*, pp. 38–55. Springer, 2024a.
 781 52, 53
- 782 Ting Liu and Siyuan Li. Hybrid global-local representation with augmented spatial guidance for
 783 zero-shot referring image segmentation. In *Proceedings of the Computer Vision and Pattern
 784 Recognition Conference*, pp. 29634–29643, 2025. 68
- 785 Yong Liu, Cairong Zhang, Yitong Wang, Jiahao Wang, Yujiu Yang, and Yansong Tang. Universal
 786 segmentation at arbitrary granularity with language instruction. In *Proceedings of the IEEE/CVF
 787 Conference on Computer Vision and Pattern Recognition*, pp. 3459–3469, 2024b. 68
- 789 Yuqi Liu, Bohao Peng, Zhisheng Zhong, Zihao Yue, Fanbin Lu, Bei Yu, and Jiaya Jia. Seg-
 790 zero: Reasoning-chain guided segmentation via cognitive reinforcement. *arXiv preprint
 791 arXiv:2503.06520*, 2025. 66
- 792 Shangbang Long, Siyang Qin, Dmitry Panteleev, Alessandro Bissacco, Yasuhisa Fujii, and Michalis
 793 Raptis. Towards end-to-end unified scene text detection and layout analysis. In *Proceedings of the
 794 IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2022. 44
- 796 Shangbang Long, Siyang Qin, Dmitry Panteleev, Alessandro Bissacco, Yasuhisa Fujii, and Michalis
 797 Raptis. Icdar 2023 competition on hierarchical text detection and recognition. *arXiv preprint
 798 arXiv:2305.09750*, 2023. 44
- 799 Ilya Loshchilov and Frank Hutter. Decoupled weight decay regularization. *ICLR*, 2019. 33
 800
- 801 Yi Lu, Jiawang Cao, Yongliang Wu, Bozheng Li, Licheng Tang, Yangguang Ji, Chong Wu, Jay
 802 Wu, and Wenbo Zhu. Rsvp: Reasoning segmentation via visual prompting and multi-modal
 803 chain-of-thought. *arXiv preprint arXiv:2506.04277*, 2025. 66
- 804 Jonathon Luiten, Aljosa Osep, Patrick Dendorfer, Philip Torr, Andreas Geiger, Laura Leal-Taixé, and
 805 Bastian Leibe. Hota: A higher order metric for evaluating multi-object tracking. *International
 806 Journal of Computer Vision*, pp. 548–578, 2021. 55
- 807 Jun Ma, Zongxin Yang, Sumin Kim, Bihui Chen, Mohammed Baharoon, Adibvafa Fallahpour, Reza
 808 Asakereh, Hongwei Lyu, and Bo Wang. Medsam2: Segment anything in 3d medical images and
 809 videos. *arXiv preprint arXiv:2504.03600*, 2025. 44, 49

- Junhua Mao, Jonathan Huang, Alexander Toshev, Oana Camburu, Alan L Yuille, and Kevin Murphy.
 Generation and comprehension of unambiguous object descriptions. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pp. 11–20, 2016. [9, 66, 68](#)
- Xiaofeng Mao, Yuefeng Chen, Yao Zhu, Da Chen, Hang Su, Rong Zhang, and Hui Xue. Coco-o: A benchmark for object detectors under natural distribution shifts, 2023. URL <https://arxiv.org/abs/2307.12730>. [52](#)
- Tim Meinhardt, Alexander Kirillov, Laura Leal-Taixe, and Christoph Feichtenhofer. Trackformer: Multi-object tracking with transformers. In *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition*, pp. 8844–8854, 2022. [22](#)
- Matthias Minderer, Alexey Gritsenko, Austin Stone, Maxim Neumann, Dirk Weissenborn, Alexey Dosovitskiy, Aravindh Mahendran, Anurag Arnab, Mostafa Dehghani, Zhuoran Shen, et al. Simple open-vocabulary object detection. In *European conference on computer vision*, pp. 728–755. Springer, 2022. [22, 26](#)
- Matthias Minderer, Alexey Gritsenko, and Neil Houlsby. Scaling open-vocabulary object detection, 2024. URL <https://arxiv.org/abs/2306.09683>. [35, 52, 53](#)
- Chaitanya Mitash, Fan Wang, Shiyang Lu, Vikedo Terhuja, Tyler Garaas, Felipe Polido, and Manikantan Nambi. Armbench: An object-centric benchmark dataset for robotic manipulation. *arXiv preprint arXiv:2303.16382*, 2023. [44, 48](#)
- Sharada Prasanna Mohanty, Gaurav Singhal, Eric Antoine Scuccimarra, Djilani Kebaili, Harris Héritier, Victor Boulanger, and Marcel Salathé. The food recognition benchmark: Using deeplearning to recognize food on images, 2021. URL <https://arxiv.org/abs/2106.14977>. [44, 48](#)
- Roozbeh Mottaghi, Xianjie Chen, Xiaobai Liu, Nam-Gyu Cho, Seong-Whan Lee, Sanja Fidler, Raquel Urtasun, and Alan Yuille. The role of context for object detection and semantic segmentation in the wild. In *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2014. [52](#)
- Minheng Ni, Yabo Zhang, Kailai Feng, Xiaoming Li, Yiwen Guo, and Wangmeng Zuo. Ref-diff: Zero-shot referring image segmentation with generative models. *arXiv preprint arXiv:2308.16777*, 2023. [68](#)
- Maxime Oquab, Timothée Darcet, Theo Moutakanni, Huy Vo, Marc Szafraniec, Vasil Khalidov, Pierre Fernandez, Daniel Haziza, Francisco Massa, Alaaeldin El-Nouby, Russell Howes, Po-Yao Huang, Hu Xu, Vasu Sharma, Shang-Wen Li, Wojciech Galuba, Mike Rabbat, Mido Assran, Nicolas Ballas, Gabriel Synnaeve, Ishan Misra, Herve Jegou, Julien Mairal, Patrick Labatut, Armand Joulin, and Piotr Bojanowski. Dinov2: Learning robust visual features without supervision. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=a68SUT6zFt>. [23](#)
- Roni Paiss, Ariel Ephrat, Omer Tov, Shiran Zada, Inbar Mosseri, Michal Irani, and Tali Dekel. Teaching clip to count to ten. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, pp. 3170–3180, 2023. [54](#)
- Kwanyong Park, Kuniaki Saito, and Donghyun Kim. Weak-to-strong compositional learning from generative models for language-based object detection. In *European Conference on Computer Vision*, pp. 1–19. Springer, 2024. [67](#)
- Paul Voigtlaender, Valentin Gabeur and Rohan Doshi. Conversational image segmentation with Gemini 2.5. <https://developers.googleblog.com/en/conversational-image-segmentation-gemini-2-5/>, 2025. [52, 66](#)
- Bryan A Plummer, Kevin J Shih, Yichen Li, Ke Xu, Svetlana Lazebnik, Stan Sclaroff, and Kate Saenko. Revisiting image-language networks for open-ended phrase detection. *IEEE transactions on pattern analysis and machine intelligence*, 44(4):2155–2167, 2020. [22](#)
- Jordi Pont-Tuset, Federico Perazzi, Sergi Caelles, Pablo Arbeláez, Alex Sorkine-Hornung, and Luc Van Gool. The 2017 davis challenge on video object segmentation. *arXiv preprint arXiv:1704.00675*, 2017a. [27](#)

- 864 Jordi Pont-Tuset, Federico Perazzi, Sergi Caelles, Pablo Arbeláez, Alexander Sorkine-Hornung, and
 865 Luc Van Gool. The 2017 davis challenge on video object segmentation. *arXiv:1704.00675*, 2017b.
 866 44
- 867
- 868 PySceneDetect Developers. PySceneDetect. <https://www.scenedetect.com/>. 41
- 869
- 870 Jiyang Qi, Yan Gao, Yao Hu, Xinggang Wang, Xiaoyu Liu, Xiang Bai, Serge Belongie, Alan Yuille,
 871 Philip Torr, and Song Bai. Occluded video instance segmentation: A benchmark. *International
 872 Journal of Computer Vision*, 2022. 55
- 873 Alec Radford, Jong Wook Kim, Chris Hallacy, Aditya Ramesh, Gabriel Goh, Sandhini Agarwal,
 874 Girish Sastry, Amanda Askell, Pamela Mishkin, Jack Clark, et al. Learning transferable visual
 875 models from natural language supervision. In *International conference on machine learning*, pp.
 876 8748–8763. PMLR, 2021. 22
- 877
- 878 Vikram V. Ramaswamy, Sing Yu Lin, Dora Zhao, Aaron B. Adcock, Laurens van der Maaten, Deepti
 879 Ghadiyaram, and Olga Russakovsky. Geode: a geographically diverse evaluation dataset for object
 880 recognition. In *NeurIPS Datasets and Benchmarks*, 2023. 44, 48
- 881 Viresh Ranjan, Udbhav Sharma, Thu Nguyen, and Minh Hoai. Learning to count everything. In
 882 *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pp.
 883 3394–3403, 2021. 44
- 884
- 885 Hanoona Rasheed, Muhammad Maaz, Sahal Shaji, Abdelrahman Shaker, Salman Khan, Hisham
 886 Cholakkal, Rao M Anwer, Eric Xing, Ming-Hsuan Yang, and Fahad S Khan. Glamm: Pixel
 887 grounding large multimodal model. In *Proceedings of the IEEE/CVF Conference on Computer
 888 Vision and Pattern Recognition*, pp. 13009–13018, 2024. 22, 66, 68
- 889 Nikhila Ravi, Valentin Gabeur, Yuan-Ting Hu, Ronghang Hu, Chaitanya Ryali, Tengyu Ma, Haitham
 890 Khedr, Roman Rädle, Chloe Rolland, Laura Gustafson, et al. SAM 2: Segment anything in images
 891 and videos. *arXiv preprint arXiv:2408.00714*, 2024. 1, 22, 27, 28, 29, 33, 44, 48
- 892
- 893 Tianhe Ren, Shilong Liu, Ailing Zeng, Jing Lin, Kunchang Li, He Cao, Jiayu Chen, Xinyu Huang,
 894 Yukang Chen, Feng Yan, et al. Grounded sam: Assembling open-world models for diverse visual
 895 tasks. *arXiv 2024*. *arXiv preprint arXiv:2401.14159*. 66
- 896
- 897 Tianhe Ren, Qing Jiang, Shilong Liu, Zhaoyang Zeng, Wenlong Liu, Han Gao, Hongjie Huang,
 898 Zhengyu Ma, Xiaoke Jiang, Yihao Chen, et al. Grounding dino 1.5: Advance the “edge” of
 899 open-set object detection. *arXiv preprint arXiv:2405.10300*, 2024. 53
- 900
- 901 Tianhe Ren, Yihao Chen, Qing Jiang, Zhaoyang Zeng, Yuda Xiong, Wenlong Liu, Zhengyu Ma,
 902 Junyi Shen, Yuan Gao, Xiaoke Jiang, Xingyu Chen, Zhuheng Song, Yuhong Zhang, Hongjie
 903 Huang, Han Gao, Shilong Liu, Hao Zhang, Feng Li, Kent Yu, and Lei Zhang. Dino-x: A
 904 unified vision model for open-world object detection and understanding, 2025. URL <https://arxiv.org/abs/2411.14347>. 52, 54
- 905
- 906 Peter Robicheaux, Matvei Popov, Anish Madan, Isaac Robinson, Joseph Nelson, Deva Ramanan, and
 907 Neehar Peri. Roboflow100-vl: A multi-domain object detection benchmark for vision-language
 908 models. *arXiv preprint arXiv:2505.20612*, 2025. 53
- 909
- 910 Olga Russakovsky, Jia Deng, Hao Su, Jonathan Krause, Sanjeev Satheesh, Sean Ma, Zhiheng Huang,
 911 Andrej Karpathy, Aditya Khosla, Michael Bernstein, et al. Imagenet large scale visual recognition
 912 challenge. *International journal of computer vision*, 115(3):211–252, 2015. 44
- 913
- 914 Chaitanya Ryali, Yuan-Ting Hu, Daniel Bolya, Chen Wei, Haoqi Fan, Po-Yao Huang, Vaibhav
 915 Aggarwal, Arkabandhu Chowdhury, Omid Poursaeed, Judy Hoffman, Jitendra Malik, Yanghao
 916 Li, and Christoph Feichtenhofer. Hiera: A hierarchical vision transformer without the bells-and-
 917 whistles. *ICML*, 2023. 23, 28
- 918
- 919 Victor Sanh, Lysandre Debut, Julien Chaumond, and Thomas Wolf. Distilbert, a distilled version of
 920 bert: smaller, faster, cheaper and lighter. *ArXiv*, abs/1910.01108, 2019. 23

- 918 Samuel Schulter, Yumin Suh, Konstantinos M Dafnis, Zhixing Zhang, Shiyu Zhao, Dimitris Metaxas,
919 et al. Omnilabel: A challenging benchmark for language-based object detection. In *Proceedings of
920 the IEEE/CVF International Conference on Computer Vision*, pp. 11953–11962, 2023. 9, 67
- 921
- 922 Shuai Shao, Zeming Li, Tianyuan Zhang, Chao Peng, Gang Yu, Xiangyu Zhang, Jing Li, and Jian
923 Sun. Objects365: A large-scale, high-quality dataset for object detection. In *Proceedings of the
924 IEEE/CVF international conference on computer vision*, pp. 8430–8439, 2019. 44
- 925
- 926 Yunhang Shen, Chaoyou Fu, Peixian Chen, Mengdan Zhang, Ke Li, Xing Sun, Yunsheng Wu,
927 Shaohui Lin, and Rongrong Ji. Aligning and prompting everything all at once for universal
928 visual perception. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern
929 Recognition (CVPR)*, pp. 13193–13203, June 2024. 52
- 930
- 931 Jianlin Su, Yu Lu, Shengfeng Pan, Bo Wen, and Yunfeng Liu. Roformer: Enhanced transformer with
932 rotary position embedding. *arXiv preprint arXiv:2104.09864*, 2021. 24
- 933
- 934 Peize Sun, Jinkun Cao, Yi Jiang, Rufeng Zhang, Enze Xie, Zehuan Yuan, Changhu Wang, and Ping
935 Luo. Trantrack: Multiple object tracking with transformer. *arXiv preprint arXiv:2012.15460*,
936 2020. 22
- 937
- 938 Shuyang Sun, Runjia Li, Philip Torr, Xiuye Gu, and Siyang Li. Clip as rnn: Segment countless visual
939 concepts without training endeavor. In *Proceedings of the IEEE/CVF Conference on Computer
940 Vision and Pattern Recognition*, pp. 13171–13182, 2024. 68
- 941
- 942 Yucheng Suo, Linchao Zhu, and Yi Yang. Text augmented spatial-aware zero-shot referring image
943 segmentation. *arXiv preprint arXiv:2310.18049*, 2023. 68
- 944
- 945 Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez,
946 Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. In I. Guyon, U. Von
947 Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, and R. Garnett (eds.), *Ad-
948 vances in Neural Information Processing Systems*, volume 30. Curran Associates, Inc.,
949 2017. URL [https://proceedings.neurips.cc/paper_files/paper/2017/
950 file/3f5ee243547dee91fb053c1c4a845aa-Paper.pdf](https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fb053c1c4a845aa-Paper.pdf). 28
- 951
- 952 Denny Vrandečić and Markus Krötzsch. Wikidata: a free collaborative knowledgebase. *Commun.
953 ACM*, 57(10):78–85, September 2014. ISSN 0001-0782. doi: 10.1145/2629489. URL <https://doi.org/10.1145/2629489>. 35
- 954
- 955 Hao Wang, Limeng Qiao, Zequn Jie, Zhijian Huang, Chengjian Feng, Qingfang Zheng, Lin Ma,
956 Xiangyuan Lan, and Xiaodan Liang. X-sam: From segment anything to any segmentation. *arXiv
957 preprint arXiv:2508.04655*, 2025a. 66, 68
- 958
- 959 Haochen Wang, Cilin Yan, Shuai Wang, Xiaolong Jiang, Xu Tang, Yao Hu, Weidi Xie, and Efstratios
960 Gavves. Towards open-vocabulary video instance segmentation. In *proceedings of the IEEE/CVF
961 international conference on computer vision*, pp. 4057–4066, 2023. 55
- 962
- 963 Peng Wang, Shuai Bai, Sinan Tan, Shijie Wang, Zhihao Fan, Jinze Bai, Keqin Chen, Xuejing Liu,
964 Jialin Wang, Wenbin Ge, et al. Qwen2-vl: Enhancing vision-language model’s perception of the
965 world at any resolution. *arXiv preprint arXiv:2409.12191*, 2024. 54
- 966
- 967 Shijie Wang, Dahun Kim, Ali Taalimi, Chen Sun, and Weicheng Kuo. Learning visual grounding
968 from generative vision and language model. In *2025 IEEE/CVF Winter Conference on Applications
969 of Computer Vision (WACV)*, pp. 8057–8067. IEEE, 2025b. 68
- 970
- 971 Yuji Wang, Jingchen Ni, Yong Liu, Chun Yuan, and Yansong Tang. Iterprime: Zero-shot referring
972 image segmentation with iterative grad-cam refinement and primary word emphasis. In *Proceedings
973 of the AAAI Conference on Artificial Intelligence*, volume 39, pp. 8159–8168, 2025c. 68
- 974
- Cong Wei, Yujie Zhong, Haoxian Tan, Yong Liu, Zheng Zhao, Jie Hu, and Yujiu Yang. Hyperseg: To-
975 wards universal visual segmentation with large language model. *arXiv preprint arXiv:2411.17606*,
976 2024. 66, 68

- 972 Nicolai Wojke, Alex Bewley, and Dietrich Paulus. Simple online and realtime tracking with a
 973 deep association metric. In *2017 IEEE international conference on image processing (ICIP)*, pp.
 974 3645–3649. IEEE, 2017. [22](#), [56](#)
- 975 Junfeng Wu, Yi Jiang, Qihao Liu, Zehuan Yuan, Xiang Bai, and Song Bai. General object foundation
 976 model for images and videos at scale. In *Proceedings of the IEEE/CVF Conference on Computer*
 977 *Vision and Pattern Recognition*, pp. 3783–3795, 2024a. [22](#), [56](#), [67](#), [68](#)
- 978 Junfeng Wu, Yi Jiang, Qihao Liu, Zehuan Yuan, Xiang Bai, and Song Bai. General object foundation
 979 model for images and videos at scale. In *Proceedings of the IEEE/CVF Conference on Computer*
 980 *Vision and Pattern Recognition*, pp. 3783–3795, 2024b. [53](#)
- 981 Zhuofan Xia, Dongchen Han, Yizeng Han, Xuran Pan, Shiji Song, and Gao Huang. Gsva: Generalized
 982 segmentation via multimodal large language models. In *Proceedings of the IEEE/CVF Conference*
 983 *on Computer Vision and Pattern Recognition*, pp. 3858–3869, 2024. [68](#)
- 984 Yin Xie, Kaicheng Yang, Xiang An, Kun Wu, Yongle Zhao, Weimo Deng, Zimin Ran, Yumeng
 985 Wang, Ziyong Feng, Roy Miles, et al. Region-based cluster discrimination for visual representation
 986 learning. *arXiv preprint arXiv:2507.20025*, 2025. [68](#)
- 987 Hu Xu, Po-Yao Huang, Xiaoqing Ellen Tan, Ching-Feng Yeh, Jacob Kahn, Christine Jou, Gargi
 988 Ghosh, Omer Levy, Luke Zettlemoyer, Wen-tau Yih, et al. Altogether: Image captioning via
 989 re-aligning alt-text. *arXiv preprint arXiv:2410.17251*, 2024a. [44](#)
- 990 Hu Xu, Saining Xie, Xiaoqing Ellen Tan, Po-Yao Huang, Russell Howes, Vasu Sharma, Shang-Wen
 991 Li, Gargi Ghosh, Luke Zettlemoyer, and Christoph Feichtenhofer. Demystifying clip data, 2024b.
 992 URL <https://arxiv.org/abs/2309.16671>. [24](#), [35](#), [44](#), [48](#)
- 993 N. Xu, L. Yang, Yuchen Fan, Dingcheng Yue, Yuchen Liang, Jianchao Yang, and Thomas S. Huang.
 994 Youtube-vos: A large-scale video object segmentation benchmark. *ArXiv*, abs/1809.03327, 2018.
 995 [44](#)
- 996 Yifan Xu, Mengdan Zhang, Chaoyou Fu, Peixian Chen, Xiaoshan Yang, Ke Li, and Changsheng Xu.
 997 Multi-modal queried object detection in the wild. *Advances in Neural Information Processing*
 998 *Systems*, 36:4452–4469, 2023. [22](#), [53](#)
- 999 Cheng-Yen Yang, Hsiang-Wei Huang, Wenhao Chai, Zhongyu Jiang, and Jenq-Neng Hwang. Samurai:
 1000 Adapting segment anything model for zero-shot visual tracking with motion-aware memory. *arXiv*
 1001 *preprint arXiv:2411.11922*, 2024. [56](#)
- 1002 Jianwei Yang, Hao Zhang, Feng Li, Xueyan Zou, Chunyuan Li, and Jianfeng Gao. Set-of-mark
 1003 prompting unleashes extraordinary visual grounding in gpt-4v. *arXiv preprint arXiv:2310.11441*,
 1004 2023. [61](#)
- 1005 Heng Yin, Yuqiang Ren, Ke Yan, Shouhong Ding, and Yongtao Hao. Rod-mllm: Towards more
 1006 reliable object detection in multimodal large language models. In *Proceedings of the Computer*
 1007 *Vision and Pattern Recognition Conference*, pp. 14358–14368, 2025. [67](#)
- 1008 En Yu, Tiancai Wang, Zhuoling Li, Yuang Zhang, Xiangyu Zhang, and Wenbing Tao. Motrv3:
 1009 Release-fetch supervision for end-to-end multi-object tracking. *arXiv preprint arXiv:2305.14298*,
 1010 2023a. [22](#)
- 1011 Fisher Yu, Haofeng Chen, Xin Wang, Wenqi Xian, Yingying Chen, Fangchen Liu, Vashisht Madhavan,
 1012 and Trevor Darrell. Bdd100k: A diverse driving dataset for heterogeneous multitask learning. In
 1013 *The IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, June 2020. [44](#), [48](#),
 1014 [49](#), [55](#)
- 1015 Seonghoon Yu, Paul Hongsuck Seo, and Jeany Son. Zero-shot referring image segmentation with
 1016 global-local context features. In *Proceedings of the IEEE/CVF conference on computer vision and*
 1017 *pattern recognition*, pp. 19456–19465, 2023b. [68](#)
- 1018 Seonghoon Yu, Paul Hongsuck Seo, and Jeany Son. Pseudo-ris: Distinctive pseudo-supervision
 1019 generation for referring image segmentation. In *European Conference on Computer Vision*, pp.
 1020 18–36. Springer, 2024. [68](#)

- 1026 Rowan Zellers, Jiasen Lu, Ximing Lu, Youngjae Yu, Yanpeng Zhao, Mohammadreza Salehi, Aditya
 1027 Kusupati, Jack Hessel, Ali Farhadi, and Yejin Choi. Merlot reserve: Multimodal neural script
 1028 knowledge through vision and language and sound. In *CVPR*, 2022. 44
- 1029 Fango Zeng, Bin Dong, Yuang Zhang, Tiancai Wang, Xiangyu Zhang, and Yichen Wei. Motr:
 1030 End-to-end multiple-object tracking with transformer. In *European conference on computer vision*,
 1031 pp. 659–675. Springer, 2022. 22
- 1032 Xiaohua Zhai, Alexander Kolesnikov, Neil Houlsby, and Lucas Beyer. Scaling vision transformers,
 1033 2022. URL <https://arxiv.org/abs/2106.04560>. 33
- 1034 Hao Zhang, Feng Li, Shilong Liu, Lei Zhang, Hang Su, Jun Zhu, Lionel M Ni, and Heung-Yeung
 1035 Shum. Dino: Detr with improved denoising anchor boxes for end-to-end object detection. *arXiv*
 1036 preprint *arXiv:2203.03605*, 2022a. 29
- 1037 Haotian Zhang, Pengchuan Zhang, Xiaowei Hu, Yen-Chun Chen, Liunian Li, Xiyang Dai, Lijuan
 1038 Wang, Lu Yuan, Jenq-Neng Hwang, and Jianfeng Gao. Glipv2: Unifying localization and vision-
 1039 language understanding. *Advances in Neural Information Processing Systems*, 35:36067–36080,
 1040 2022b. 53
- 1041 Tao Zhang, Xiangtai Li, Hao Fei, Haobo Yuan, Shengqiong Wu, Shunping Ji, Chen Change Loy,
 1042 and Shuicheng Yan. Omg-llava: Bridging image-level, object-level, pixel-level reasoning and
 1043 understanding. *Advances in neural information processing systems*, 37:71737–71767, 2024a. 22
- 1044 Yifu Zhang, Peize Sun, Yi Jiang, Dongdong Yu, Fucheng Weng, Zehuan Yuan, Ping Luo, Wenyu
 1045 Liu, and Xinggang Wang. Bytetrack: Multi-object tracking by associating every detection box. In
 1046 *European conference on computer vision*, pp. 1–21. Springer, 2022c. 22, 56
- 1047 Yuxuan Zhang, Tianheng Cheng, Lianghui Zhu, Rui Hu, Lei Liu, Heng Liu, Longjin Ran, Xiaoxin
 1048 Chen, Wenyu Liu, and Xinggang Wang. Evf-sam: Early vision-language fusion for text-prompted
 1049 segment anything model. *arXiv preprint arXiv:2406.20076*, 2024b. 68
- 1050 Zheng Zhang, Yeyao Ma, Enming Zhang, and Xiang Bai. Psalm: Pixelwise segmentation with large
 1051 multi-modal model. In *European Conference on Computer Vision*, pp. 74–91. Springer, 2024c. 68
- 1052 Zhixiong Zhang, Shuangrui Ding, Xiaoyi Dong, Songxin He, Jianfan Lin, Junsong Tang, Yuhang
 1053 Zang, Yuhang Cao, Dahua Lin, and Jiaqi Wang. Sec: Advancing complex video object segmenta-
 1054 tion via progressive concept construction. *arXiv preprint arXiv:2507.15852*, 2025. 56
- 1055 Shiyu Zhao, Long Zhao, Yumin Suh, Dimitris N Metaxas, Manmohan Chandraker, Samuel Schulter,
 1056 et al. Generating enhanced negatives for training language-based object detectors. In *Proceedings*
 1057 of the IEEE/CVF Conference on Computer Vision and Pattern Recognition, pp. 13592–13602,
 1058 2024. 67
- 1059 Bolei Zhou, Hang Zhao, Xavier Puig, Tete Xiao, Sanja Fidler, Adela Barriuso, and Antonio Torralba.
 1060 Semantic understanding of scenes through the ade20k dataset. *International Journal of Computer*
 1061 *Vision*, 127(3):302–321, 2019. 52
- 1062 Yang Zhou, Shiyu Zhao, Yuxiao Chen, Zhenting Wang, Can Jin, and Dimitris N Metaxas. Led: Llm
 1063 enhanced open-vocabulary object detection without human curated data generation. *arXiv preprint*
 1064 *arXiv:2503.13794*, 2025. 67
- 1065 Pengfei Zhu, Longyin Wen, Dawei Du, Xiao Bian, Heng Fan, Qinghua Hu, and Haibin Ling.
 1066 Detection and tracking meet drones challenge. *IEEE Transactions on Pattern Analysis and*
 1067 *Machine Intelligence*, 44(11):7380–7399, 2021. 44, 49
- 1068 Xizhou Zhu, Weijie Su, Lewei Lu, Bin Li, Xiaogang Wang, and Jifeng Dai. Deformable detr:
 1069 Deformable transformers for end-to-end object detection. *arXiv preprint arXiv:2010.04159*, 2020.
 1070 3, 29
- 1071 Xueyan Zou, Jianwei Yang, Hao Zhang, Feng Li, Linjie Li, Jianfeng Wang, Lijuan Wang, Jianfeng
 1072 Gao, and Yong Jae Lee. Segment everything everywhere all at once. *Advances in neural information*
 1073 *processing systems*, 36:19769–19782, 2023. 66

1080	APPENDIX	
1081		
1082	A Related Work	22
1083	B Ablations	22
1084	B.1 Model Ablations	22
1085	B.2 Image Training Data Ablations	24
1086	B.3 Automatic Domain Adaptation via the Synthetic Data Engine	25
1087	B.4 Image Data Engine Annotation Speed	26
1088	B.5 Video Training Data Ablations	27
1089	C Limitations	28
1090	D Model Details	28
1091	D.1 Model Architecture	28
1092	D.2 Image Implementation Details	28
1093	D.3 Video Implementation Details	31
1094	D.4 Model Training	32
1095	E Data Engine Details	35
1096	E.1 Media Pool	35
1097	E.2 SA-Co Ontology	35
1098	E.3 Phase 1: Human Verification	35
1099	E.4 Phase 2: Human + AI Verification	36
1100	E.5 Phase 3: Scaling and Domain Expansion	41
1101	E.6 Phase 4: Video Annotation	41
1102	F SA-Co Dataset and Metric Details	43
1103	F.1 SA-Co Training Data	43
1104	F.2 SA-Co Evaluation Benchmark	43
1105	F.3 Metrics	44
1106	F.4 Human Performance on SA-Co	46
1107	F.5 Additional Dataset Examples	47
1108	G Additional Experiments and Details	52
1109	G.1 PCS with NPs on Images	52
1110	G.2 Visual Exemplars and Interactivity	52
1111	G.3 Few-Shot Fine-tuning	53
1112	G.4 Object Counting	54
1113	G.5 Video PCS Details	55
1114	G.6 PVS Details	56
1115	H SAM 3 Agent	60
1116	H.1 Agent Design	60
1117	H.2 Qualitative Analysis	61
1118	H.3 Full Quantitative Results	65
1119		
1120		
1121		
1122		
1123		
1124		
1125		
1126		
1127		
1128		
1129		
1130		
1131		
1132		
1133		

1134 **A RELATED WORK**

1135
 1136 **Promptable and Interactive Visual Segmentation.** SAM 1 (Kirillov et al., 2023) introduces
 1137 “promptable” image segmentation with interactive refinement. While the original task definition
 1138 included text prompts, they were not fully developed. SAM 2 (Ravi et al., 2024) extended the
 1139 promptable visual segmentation task to video, allowing refinement points on any frame. SAM 3
 1140 inherits geometry-based segmentation from SAM 1 and 2, while extending to include text and image
 1141 exemplar prompts that segment all instances of a specified concept in images and videos.

1142 **Open-Vocabulary Detection and Segmentation in Images** exhaustively labels every instance of an
 1143 open-vocabulary object category with a coarse bounding box (detection) or a fine-grained pixel mask
 1144 (segmentation). Recent open-vocabulary (OV) detection (Gu et al., 2021; Minderer et al., 2022) and
 1145 segmentation (Ding et al., 2022; Liang et al., 2023) methods leverage large-scale vision-language
 1146 encoders such as CLIP (Radford et al., 2021) to handle categories described by arbitrary text, even
 1147 those never seen during training. While DETR (Carion et al., 2020) is limited to a closed set of
 1148 categories seen during training, MDETR (Kamath et al., 2021) evolves the approach to condition
 1149 on raw text queries, such as a sentence or a caption. Image exemplars used as prompts to specify
 1150 the desired object category (e.g., DINOV (Li et al., 2023a), T-Rex2 (Jiang et al., 2024)) present a
 1151 practical alternative to text, but fall short in conveying the abstract concept of objects as effectively
 1152 as text prompts. We introduce a new benchmark for OV segmentation with $> 100\times$ more unique
 1153 concepts than prior work, and release a state-of-the-art model.

1154 **Visual Grounding** localizes a language expression referring to a region of the image with a box or
 1155 mask. (Plummer et al., 2020) introduces phrase detection as both deciding whether the phrase is
 1156 relevant to an image and localizing it. GLIP (Li et al., 2022b) and GroundingDino (Liu et al., 2023)
 1157 formulate object detection as phrase grounding, unifying both tasks during training. MQ-GLIP (Xu
 1158 et al., 2023) adds image exemplars to text as queries. Building on this trend toward models supporting
 1159 multiple tasks and modalities, GLEE (Wu et al., 2024a) allows text phrases, referring expressions, and
 1160 visual prompts for category and instance grounding in both images and videos. Unlike SAM 3, GLEE
 1161 does not support exemplars or interactive refinement. LISA (Lai et al., 2024) allows segmentation
 1162 that requires reasoning, while OMG-LLaVa (Zhang et al., 2024a) and GLaMM (Rasheed et al.,
 1163 2024) generate natural language responses interleaved with corresponding segmentation masks, with
 1164 GLaMM accepting both textual and optional image prompts as input. Some general-purpose MLLMs
 1165 can output boxes and masks (Gemini2.5 (Comanici et al., 2025)) or points (Molmo (Deitke et al.,
 1166 2025)). SAM 3 can be used as a “vision tool” in combination with an MLLM (§6).

1167 **Multi-Object Tracking (MOT) and Segmentation** methods identify and track multiple instances of
 1168 an object category in video, associating each with a unique ID. In tracking-by-detection methods,
 1169 detection is performed independently on each frame to produce boxes and confidence scores, followed
 1170 by association of boxes using motion-based and appearance-based matching as in SORT (Bewley
 1171 et al., 2016; Wojke et al., 2017), Tracktor (Bergmann et al., 2019), ByteTrack (Zhang et al., 2022c),
 1172 SAM2MOT (Jiang et al., 2025), or OC-SORT (Cao et al., 2023). An alternative is an end-to-end
 1173 trainable architecture that jointly detects and associates objects, e.g., TrackFormer (Meinhardt et al.,
 1174 2022), TransTrack (Sun et al., 2020), or MOTR (Zeng et al., 2022). TrackFormer uses a DETR-like
 1175 encoder-decoder that initializes new tracks from static *object queries* and auto-regressively follows
 1176 existing tracks with identity-preserving *track queries*. A challenge with joint models is the conflict
 1177 between detection and association (Yu et al., 2023a; Gao et al., 2025), where one needs to focus on
 1178 semantics while the other on instance identities. In SAM 3, we tackle this challenge with a video
 1179 tracker that is built on top of a strong image detector to propagate masks across frames.

1180 **B ABLATIONS**

1181 **B.1 MODEL ABLATIONS**

1182
 1183 **Presence Token.** We first ablate the impact of the presence token and the approach to its training.
 1184 The presence token is included in the decoder (discussed further in §D.2), together with the object
 1185 queries, and predicts a *concept* presence score. The presence score receives gradients *only* on the
 1186 PCS task during joint training and is *always* supervised with the presence (or absence) of the concept
 1187 in the image using a binary cross-entropy loss. Using a presence token to *decouple* presence and
 1188 localization brings large gains in performance, see Tab. 10.

		SA-Co/Gold		
	Presence token	CGF ₁	IL_MCC	pmF ₁
	×	57.6	0.77	74.7
	✓	63.3	0.82	77.1

Table 10: **Presence token.** Decoupling presence and localization brings large gains.

When used with a presence score, we found that the best supervision strategy for the box/mask object scores is to *not* receive gradients when a concept is an image-level negative, see Setting (a) in Tab. 11. Note that this is in contrast to the approach in typical DETR variants, where all individual object scores are supervised *negatively* to reflect the absence of the concept in the image, see Setting (b) in Tab. 11. We find that (b) works worse than (a) when used with the presence score. When a concept is *present* in the image, individual object queries always receive classification supervision based on Hungarian matching. Setting (a) is consistent with our recognition-localization decoupled design, where the presence score is responsible for recognition (existence in the image) and the object scores are responsible for localization (i.e., rank the best match to the *positive* ground-truth highest among all the proposals).

During inference, we use the product of the global presence score and the object score as the total object score. In Setting (c), we explored directly supervising the *total* object scores (instead of the typical object scores) as positive or negative (as determined by matching) but observed worse performance. Finally, Setting (d) in Tab. 11 investigates detaching the presence score from the computation graph while supervising the total scores, but this does not improve over (c).

	Supervise mask scores only when concept present	Supervise total score	Sup. total score, detach presence	SA-Co/Gold		
				CGF ₁	IL_MCC	pmF ₁
a.	✓	✗	✗	65.0	0.82	79.7
b.	✗	✗	✗	64.1	0.81	79.2
c.	✓	✓	✗	63.6	0.83	76.7
d.	✓	✗	✓	63.4	0.83	76.7

Table 11: **Supervision strategy for object/mask scores for a model with a presence token.** We find the best supervision strategy is to supervise mask scores only for positive concepts and to supervise the presence and mask scores separately, although their product is used as the total object score during inference.

Training with presence can be considered as a form of post-training and occurs in Stage 3 (see §D.4.1) of our training pipeline. By default, *ablations* do *not* undergo this stage unless otherwise mentioned.

Vision and Text Encoder. While SAM 2 uses an MAE (He et al., 2022) pre-trained Hiera (Ryali et al., 2023) for its strong localization capability and efficiency for the more *geometric* PVS task, SAM 3 also needs strong *semantic* and *linguistic* understanding with broad coverage. We adapted PE (Bolya et al., 2025) for the vision and text encoders of SAM 3, so that a large and diverse set of concepts is seen during Stage 1 of training, while producing *aligned* image and text encoders. In Tab. 12, we compare performance with Hiera and DINOv2 (Oquab et al., 2024); since these vision encoders lack an aligned text encoder, we use DistilRoBERTa-base (Sanh et al., 2019). We find PE to be the best overall choice of vision backbone, and using its own aligned text encoder provides further gains over PE with an unaligned text baseline. Use of PE enables strong robustness in SAM 3 (here measured by AP on COCO-O, demonstrating good object detection across various domain shifts, e.g. “sketch”, “cartoon”, “painting”, etc.).

	Encoder (patch size)	SA-Co/Gold (CGF ₁)	COCO-O (AP)
	PE-L+ (14)	50.4	42.5
	PE-L+ (14) w/ DistilRoBERTa	43.3	37.9
	DINOv2-L (14)	42.0	31.6
	Hiera-L (16)	37.0	22.0

Table 12: **Choice of encoders.** As SAM 3 needs both semantic visual and linguistic understanding, we find PE’s aligned image and text encoders work well.

Implementation Details. The image resolution was set to 1008 px, 1008 px, 1152 px for PE, DINOv2, Hiera, respectively, ensuring the same number of tokens in the detector due to their differences in

patch size. All vision encoders used global attention in only a subset of the layers, using windowed (24×24 tokens) attention otherwise. Since Hiera is a hierarchical encoder, we set the window size to 24×24 in stage 3 of the encoder, which has most of the FLOPs. Since PE is capable of using relative positional information via RoPE (Su et al., 2021; Heo et al., 2024), we include relative positional embeddings in global layers for Hiera and DINOv2 following Bolya et al. (2024). All models were trained using SA-Co/HQ viewing 5 million samples over the course of training. Recipe was separately optimized for each choice of encoders. Tokens from the respective vision encoders are downsampled by 2×2 to 1296 tokens before being passed to the fusion encoder and detector.

B.2 IMAGE TRAINING DATA ABLATIONS

Setup. We adopt a simplified, lighter model and training strategy for ablations in this section. Specifically, we use (i) a stride-28 (instead of 14) variant of SAM 3 using $4 \times$ fewer tokens in the detector, (ii) limit to 45% of the entire SA-Co/SYN dataset and adopt (iii) shorter training schedules and do not run “presence post-training” (see §B). This allows running ablations more efficiently (but results in lower absolute accuracy vs. SAM 3). We observed similar trends when training at scale.

SAM 3 Training Data. Tab. 13 analyzes the impact of various SA-Co training data subsets. Training with even with *just* SA-Co/EXT shows comparable performance with *best* external models on SA-Co/Gold (see OWLv2’s and DINO-X’s performance in Tab. 1), indicating a strong base model. Adding synthetic data SA-Co/SYN into the training mix results in significantly improved performance. The performance further increases after adding the high-quality SA-Co/HQ data due to its quality and distributional similarity with SA-Co/Gold. Although SA-Co/HQ is large-scale and in-domain with SA-Co/Gold, SA-Co/SYN shows further gains on SA-Co/Gold when added on top of SA-Co/HQ.

	SA-Co/EXT	SA-Co/SYN	SA-Co/HQ	CGF ₁	IL_MCC	pmF ₁
	✓	✗	✗	30.9	0.46	66.3
	✓	✓	✗	39.7	0.57	70.6
	✓	✗	✓	51.8	0.71	73.2
	✓	✓	✓	54.3	0.74	73.5

Table 13: **Training data effect on SA-Co/Gold.** Adding each subset improves performance.

Training data	SA-Co/Gold (All) (in-domain)			SA-Co/Gold-MetaCLIP (in-domain)			SA-Co/Gold-Wiki-Food&Drink (in domain)		
	CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁
SA-Co/EXT	30.9	0.50	66.3	28.1	0.45	62.3	30.49	0.45	67.5
+ 1% SA-Co/HQ	36.4	0.54	68.7	36.1	0.56	64.1	40.3	0.55	73.3
+ 4% SA-Co/HQ	44.2	0.62	71.1	43.2	0.65	66.8	50.2	0.66	75.6
+ 10% SA-Co/HQ	46.7	0.65	71.7	45.6	0.67	67.7	54.7	0.71	76.9
+ 20% SA-Co/HQ	49.0	0.68	72.2	46.8	0.68	68.2	57.5	0.74	77.4
+ 100% SA-Co/HQ	51.8	0.71	73.2	49.0	0.71	69.4	59.9	0.77	77.5
Teacher (Human)	[74.2, 81.4]	[0.88, 0.96]	84.7	[68.2, 76.3]	[0.86, 0.96]	79.8	[79.2, 85.0]	[0.91, 0.97]	87.3

Table 14: **SA-Co/HQ scaling.** SA-Co/EXT data alone is not enough to solve SA-Co/Gold, training on SA-Co/HQ scales well with increasing amount of data. Human performance given as an estimated range where applicable, see §F.4 for details. Ablations use a lighter model and training setting vs SAM 3.

SA-Co/HQ Scaling Law. Tab. 14 investigates scaling behavior of the SA-Co/HQ training data. For this ablation, the data mix is sampled randomly from the entire SA-Co/HQ (collected from the three phases in §4) at a fixed percentage. We also report scaling behavior on two specific subsets of SA-Co/Gold: the MetaCLIP Xu et al. (2024b) subset annotated with generic caption-derived NPs, and Wiki-Food&Drink subset annotated with fine-grained NPs from SA-Co Ontology nodes. SA-Co/HQ improves performance on both subsets as expected, since they are from the same distribution (in-domain). We also report the Teacher (Human) performance in the last row. Due to the simplified setting, the gap between SAM 3 and Human is larger than that of the best SAM 3 model.

SA-Co/SYN Scaling Law. Tab. 15 shows that SAM 3 scales well with SA-Co/SYN data on SA-Co/Gold benchmark as it benefits from the large scale concepts captured from image captions generated by Llama4 and alt-text associated with the images, for *both* the *in-domain* MetaCLIP subset

1296 1297 1298	Training data	SA-Co/Gold (All)			SA-Co/Gold-MetaCLIP (in-domain)			SA-Co/Gold-Wiki-Food&Drink (out-of-domain)		
		CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁
	SA-Co/EXT	30.9	0.46	66.3	28.1	0.45	62.3	30.49	0.45	67.5
+ 1% SA-Co/SYN	35.9	0.5	69.3	37.3	0.6	64.5	32.4	0.4	73.1	
+ 4% SA-Co/SYN	36.9	0.5	69.7	38.5	0.59	64.9	34.0	0.46	73.1	
+ 15% SA-Co/SYN	38.9	0.56	70.2	40.8	0.62	65.5	36.6	0.50	73.8	
+ 45% SA-Co/SYN	39.7	0.57	70.6	42.5	0.64	66.1	37.4	0.51	73.9	
Teacher (SAM 3 + AI verifiers)	66.1	0.84	78.5	62.5	0.83	75.7	70.0	0.87	80.8	

Table 15: **SA-Co/SYN scaling.** SAM 3 benefits from increasing SA-Co/SYN data, *both* on MetaCLIP which is *in-domain* with the synthetic data, and on Wiki concepts which are *out-of-domain* of the synthetic data. The teacher that generated the SA-Co/SYN data consists of an older version of SAM 3, and AI verifiers from the SAM 3 data engine. Ablations use a lighter model and training setting *vs* SAM 3.

and the *out-of-domain* Wiki-Food&Drink subset within the SA-Co/Gold benchmark. The last row shows the Teacher performance (an older version of SAM 3 and AI verifiers) is much better than the student, and explains why SA-Co/SYN is useful. When comparing the SA-Co/SYN in Tab. 15 and SA-Co/HQ in Tab. 14, the lower in-domain performance gap on MetaCLIP (42.5 vs 49.0) comes from the relatively weaker annotation quality of SA-Co/SYN, due to lacking of the human correction step. The gap is larger on the out-of-domain Wiki-Food&Drink set (37.4 vs 59.9), because SA-Co/SYN only covers the MetaCLIP images and noun phrases from a captioning model; see Table 28. We also show in §B.3 that with additional *in-domain* synthetic data, we can close the performance gap on SA-Co/Gold-Wiki-Food&Drink subset without *any* human involvement.

Hard Negatives. We ablate the number of hard negative noun phrases in SA-Co/HQ per image in Tab. 16. We show that increasing the number of negatives improves SAM 3 performance across all metrics, most notably IL_MCC. Hard negatives are phrases that are not present in the image but that (a previous generation of) SAM 3 predicts masks for, i.e., they are adversarial to (a previous generation of) SAM 3. Training on such difficult distractors helps improve the image-level classification performance captured by the IL_MCC metric.

SAM 3 and AI Verifiers. AI verifiers improve performance over the final SAM 3 model alone on the PCS task, as shown in Tab. 17. We first replace the presence score from SAM 3 with a presence score from the Exhaustivity Verification (EV) AI verifier (given the image and noun phrase with no objects as input, the probability of *not exhaustive*, defined in Tab. 24). This results in a +3.1 point gain in CGF₁, all due to IL_MCC. Then we apply the Mask Verification (MV) AI verifier to each mask, and remove the rejected masks. This results in a further +1.1 point gain in CGF₁, all due to pmF₁. The system closes half the gap between SAM 3 and human performance, which indicates potential further improvements of SAM 3 by scaling up the SA-Co/SYN data and SAM 3 model size.

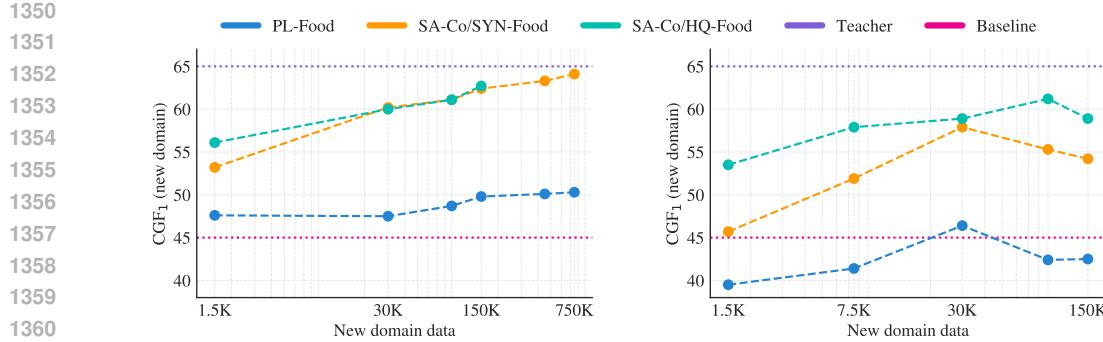
# negatives	CGF ₁	IL_MCC	pmF ₁	Model	CGF ₁	IL_MCC	pmF ₁
0	31.8	0.44	70.2	SAM 3	65.0	0.818	79.7
5	44.8	0.62	71.9	+ EV AI verifier	68.1	0.864	78.8
15	47.9	0.67	72.1	+ MV AI verifier	69.2	0.852	81.3
30	49.2	0.68	72.3	Human	[74.2, 81.4]	[0.876, 0.962]	84.7

Table 16: **Effect of hard negatives on SA-Co/Gold.** More negatives improve performance.

Table 17: **Effect of using AI verifiers** on top of SAM 3, evaluated on SA-Co/Gold.

B.3 AUTOMATIC DOMAIN ADAPTATION VIA THE SYNTHETIC DATA ENGINE

In this section, we demonstrate that with domain-specific synthetic data generated with the synthetic data engine (using SAM 3 + AI verifiers), we can significantly improve performance on a new domain *without any human annotation*. We select “Food & drink” Wiki nodes in our ontology as the new domain and mine images from MetaCLIP (refer to *Concept Selection*, *Offline Concept Indexing* and *Online Mining* steps in §E.4 for more details on data mining). For pseudo-annotating fine-grained “Food & drink” concepts, we use our ontology to identify relevant coarse-grained concepts that SAM 3 works well on and use this as the prompt to generate masks. We ensure that no data from the new domain was used in training the AI annotators (including both SAM 3 and AI verifiers) that were used to generate the data.



(a) Data scaling mixing pre-training data at a 1:1 ratio. (b) Data scaling without mixing pre-training data.

Figure 7: **Domain adaptation via synthetic data.** (a) SAM 3 + AI verifiers (teacher system) can annotate synthetic (SYN) data in new domains (e.g., fine-grained food concepts) and achieve similar scaling behavior as with *human-annotated* (HQ) data. (b) Not mixing in high-quality pre-training data can limit performance gains when fine-tuning on new domains, particularly when using synthetic data.

We study data scaling laws for three variants of “Food & drink” training data, evaluating performance on the Wiki-Food&Drink subset of the SA-Co/Gold benchmark:

- **PL-Food:** The target concepts are annotated using the data engine *without* AI verification. This data is similar to typical pseudo-labelled data used in prior work for self-training (e.g. Minderer et al. (2022)).
- **SA-Co/SYN-Food:** This data is PL-Food, but *cleaned* by AI verifiers.
- **SA-Co/HQ-Food:** Data annotated by *human* annotators.

We train the models in 2 steps to isolate the impact of the data from the new domain from other data as well as to amortize training costs. We first pre-train a base model using SA-Co/HQ *minus* SA-Co/HQ-Food to establish base capability and a common starting point. Next, we fine-tune the same base model with the three data variants in two settings: with or without mixing the pre-training data.

Fig. 7a shows the scaling law with mixing the pre-training data in a 1:1 ratio. We observe some modest improvement in performance with PL-Food compared to the baseline, but there is a large gap to the other variants due to its lower quality. SA-Co/HQ-Food and SA-Co/SYN-Food have similar data scaling behavior, but a model trained using SA-Co/SYN-Food can eventually surpass using SA-Co/HQ-Food because SA-Co/SYN-Food can be scaled up without incurring any annotation cost. The model trained on SA-Co/SYN-Food eventually gets very close to the performance of its teacher system.

In the second setting, we do not mix the pre-training data with the fine-tuning data from the new domain. All three data variants result in worse performance in Fig. 7b (vs Fig. 7a). In this setting, there is a larger gap between SA-Co/HQ-Food and SA-Co/SYN-Food reflecting the lower quality of SA-Co/SYN-Food (mainly lack of exhaustivity due to no human correction) but can eventually reach a similar level of performance as using SA-Co/HQ-Food (likely due to volume and thereby improved coverage). Comparing Fig. 7a and 7b, it is beneficial to include high-quality general-domain data when fine-tuning SAM 3 on new domains, particularly when using synthetic data, allowing us to “make up” for the lower quality in the synthetic data to some degree.

B.4 IMAGE DATA ENGINE ANNOTATION SPEED

Tab. 18 measures the speedup in the SAM 3 data engine from adding AI verifiers when collecting data on a new domain with fine-grained concepts. We use the same setup as §B.3, annotating Wiki-Food&Drink data generated with a data engine where neither SAM 3 nor AI verifiers have been trained on Wiki-Food&Drink data. We annotate the same set of image-NP pairs in four settings:

- **Human (NP Input).** A human annotator is given a single image noun-phrase pair from SA-Co/HQ-Food, and is required to manually annotate all instance masks. No mask proposals or AI-verifiers are used in the loop.
- **Human (Mask Input).** The same annotation task as “NP input” but in this setting, the human annotators starts with PL-Food, i.e., image noun-phrase pairs with mask proposals generated by SAM 3.
- **Engine (All Human)** Similar to Phase 1 in the SAM 3 data engine, humans start with PL-Food, and sequentially perform 3 tasks: Mask Verification, Exhaustivity Verification and Correction. All three tasks are performed by humans.
- **Engine (Full)** Similar to Phase 3 in the SAM 3 data engine, Mask Verification and Exhaustivity Verification tasks are completed by AI verifiers, and Correction is done by humans i.e human annotators in the manual annotation task start with SA-Co/SYN-Food.

Task	Human from NP	Human from masks	Engine - all human	Engine - full
Time for datamix (sec)	90	86	50	23
Time for positive NP (sec)	236	205	207	152
Time for negative NP (sec)	71	70	30	6

Table 18: **Data engine efficiency.** AI verifiers significantly increase throughput, allowing humans to focus on challenging cases and the manual correction task. AI verifiers allow for a 5x speed up on negative phrases and a 36% speed up for positive phrases. The time for datamix is calculated based on 88.5% negatives in SA-Co/HQ, see Tab. 26 for dataset composition. Timing is calculated based on the Wiki-Food&Drink domain. Compared to captioner-based domains, fine-grained domains require more research by annotators to understand the concept, leading to much higher annotation times for negative phrases and amortized time per mask.

Tab. 18 shows that a version of the SAM 3 model and AI verifiers that were never trained in this new domain *double* the throughput of the data engine. AI verifiers also allow verifying generated hard negative NPs at scale with close to no human-annotator involvement. As SAM 3 and AI verifiers are updated with the collected data and improve, human annotators need to manually correct fewer errors. This leads to increasingly higher throughput and the collection of more challenging data for a given amount of human annotation time.

In Tab. 25, we show that AI verifiers achieve a similar even better performance on the MV and EV tasks than human verifiers, so the quality of annotations from these four settings are similar.

B.5 VIDEO TRAINING DATA ABLATIONS

We analyze how much the SAM 3 model benefits from the videos and annotations in SA-Co/VIDEO obtained through the video data engine, which are used in Stage 4 (video-level) training (described further in §D.4.1). Specifically, we train the model with a varying amount of masklets from SA-Co/VIDEO as VOS training data, and evaluate the resulting checkpoints on SA-Co/VEval under the VOS task with the $\mathcal{J} \& \mathcal{F}$ metric. The results are shown in Tab. 19, where adding masklets collected with noun phrases through the video data engine (as additional Stage 4 training data) improves the $\mathcal{J} \& \mathcal{F}$ performance on both SA-Co/VEval and public benchmarks such as DAVIS17 (Pont-Tuset et al., 2017a) and SA-V (Ravi et al., 2024).

Stage-4 training data	SA-Co/VEval YT-1B	SA-Co/VEval SA-V	DAVIS17	SA-V	SA-V
	val $\mathcal{J} \& \mathcal{F}$	test $\mathcal{J} \& \mathcal{F}$			
using SAM 2 video data only	80.7	84.9	91.6	77.0	77.1
+ 25% SA-Co Train videos	80.9	85.3	91.3	75.2	76.9
+ 50% SA-Co Train videos	81.2	85.3	91.5	76.7	77.0
+ 75% SA-Co Train videos	81.4	85.9	91.7	76.5	78.5
+ 100% SA-Co Train videos	81.4	86.5	91.5	77.4	78.0

Table 19: Scaling analysis on SA-Co/VEval under Stage 4 (video-level) training, evaluated on multiple benchmarks through the Video Object Segmentation (VOS) task under the $\mathcal{J} \& \mathcal{F}$ metric. Note that “SA-Co/VEval YT-1B” and “SA-Co/VEval SA-V” refer to the subset of SA-Co/VEval built upon YT-Temporal-1B videos and SA-V videos respectively, while “SA-V” referred to the VOS evaluation dataset released in Ravi et al. (2024).

1458 **C LIMITATIONS**
 1459

1460 SAM 3 demonstrates strong performance on the PCS task in images and videos but has limitations in
 1461 many scenarios.

1462 SAM 3 struggles to generalize to fine-grained out-of-domain concepts (e.g., aircraft types, medical
 1463 terms) in a zero-shot manner, especially in niche visual domains (e.g., thermal imagery). Concept
 1464 generalization for PCS is inherently more challenging than the class-agnostic generalization to new
 1465 visual domains for the PVS task, with the latter being the key that enables SAM and SAM 2 to be
 1466 successfully applied zero-shot in diverse settings. Our experiments show that SAM 3 is able to quickly
 1467 adapt to new concepts and visual domains when fine-tuned on small quantities of human-annotated
 1468 data (Tab. 2). Further, we show that we can improve the performance in a new domain without any
 1469 human involvement (§B.3), using domain-specific synthetic data generated using our data engine.

1470 From our formulation of the PCS task, SAM 3 is constrained to simple noun phrase prompts and does
 1471 not support multi-attribute queries beyond one or two attributes or longer phrases including referring
 1472 expressions. We show that when combined with an MLLM, SAM 3 is able to handle more complex
 1473 phrases well (§6 and §H).

1474 In the video domain, SAM 3 tracks every object with a SAM 2 style masklet, which means the cost
 1475 of SAM 3 inference scales linearly with the number of objects being tracked. To support real-time
 1476 inference (30 FPS) on videos in practical applications (e.g., a web demo), we parallelize the inference
 1477 over multiple GPUs: up to 10 objects on 2 H200s, up to 28 objects on 4 H200s, and up to 64
 1478 objects on 8 H200s. Further, under the current architecture, there is no shared object-level contextual
 1479 information to aid in resolving ambiguities in multi-object tracking scenarios. Future developments
 1480 could address this through shared global memory across multiple objects, which would also improve
 1481 inference efficiency.

1482 Supporting concept-level interactivity for PCS, alongside instance-level interactivity for PVS, poses
 1483 several challenges. To support instance-level modifications without affecting all other instances of the
 1484 concept, we enforce a hard “mode-switch” within the model from concept to instance mode. Future
 1485 work could include more seamlessly interleaving concept and instance prompts.

1486 **D MODEL DETAILS**
 1487

1488 **D.1 MODEL ARCHITECTURE**
 1489

1490 Our architecture is broadly based on the SAM series (Ravi et al., 2024; Kirillov et al., 2023) and
 1491 DETR (Carion et al., 2020) and uses a (dual) encoder-decoder transformer architecture, see Fig. 8
 1492 for an overview. SAM 3 is a generalization of SAM 2, supporting the new Promptable Concept
 1493 Segmentation (PCS) task as well as the Promptable Visual Segmentation (PVS) task (Ravi et al.,
 1494 2024). The design supports *multimodal prompts* (e.g., text, boxes, points) and *interactivity*, in images
 1495 and videos. We first discuss the detector architecture for images and build upon it to discuss the
 1496 tracker components for video.

1497 **D.2 IMAGE IMPLEMENTATION DETAILS**
 1498

1499 The image detector is an encoder-decoder transformer architecture.
 1500

1501 **Image and Text Encoders.** The image and text encoders are Transformers (Vaswani et al., 2017)
 1502 trained using contrastive vision language training using 5.4 billion image-text pairs following
 1503 Perception Encoder (PE) (Bolya et al., 2025), see §D.4.1 for training details. As in SAM 2, the
 1504 vision encoder uses windowed attention (Ryali et al., 2023; Li et al., 2022d) and global attention
 1505 in only a small subset of layers (4 out of 32), where an image of 1008 pixels is divided into 3×3
 1506 non-overlapping windows of 336 pixels each. The vision encoder uses RoPE in each layer and
 1507 windowed absolute positional embeddings as in Bolya et al. (2024). The text encoder is causal, with
 1508 a maximum context length of 32.

1509 As in Ravi et al. (2024), we use a streaming approach, ingesting new frames as they become available.
 1510 We run the PE backbone only *once* per frame for the entire interaction, which can span multiple
 1511 forward/backward propagation steps through a video. The backbone provides unconditioned tokens

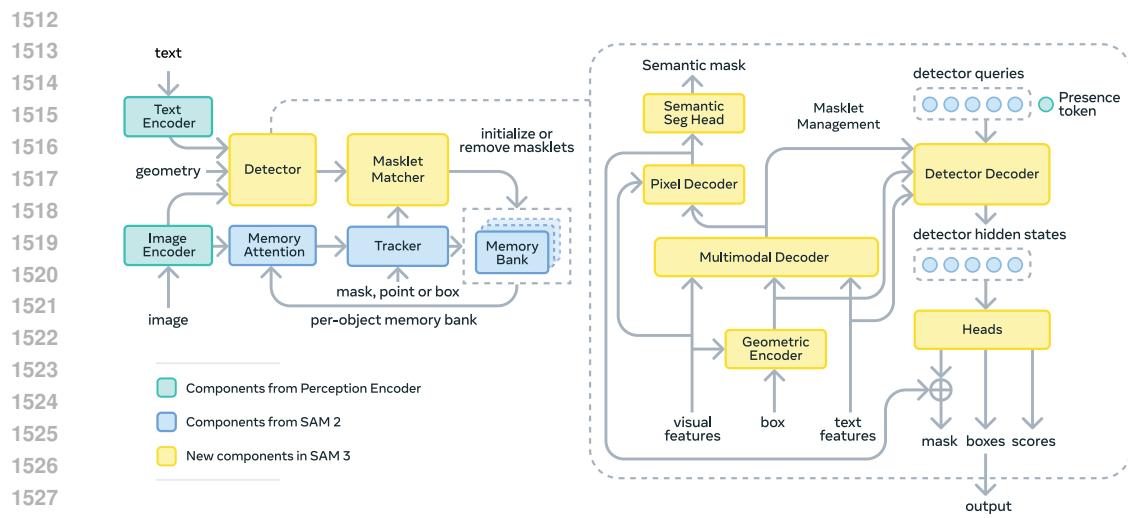


Figure 8: SAM 3 architecture. We highlight new components in yellow, SAM 2 (Ravi et al., 2024) in blue and PE (Bolya et al., 2025) in cyan.

(feature embeddings) representing each frame to the dual-encoder consisting of the fusion encoder and memory attention, described below.

Geometry and Exemplar Encoder. The geometry and exemplar encoder is *primarily* used to encode image *exemplars* (if present) for the PCS task. It is additionally also used to encode *visual prompts* for the PVS task on images but this is an *auxiliary* functionality that is primarily used to include pre-training data for the PVS task in stages-2,-3 of training (see §D.4.1) and is thereby used to enable a more modular training approach.

Each individual image exemplar is encoded using positional embedding, label embedding (positive or negative) and ROI-pooled visual features that are concatenated (comprising “exemplar tokens”) and processed by a small transformer. Visual prompts (points, boxes) for auxiliary training are encoded in a similar manner, comprising “geometry tokens”. It is possible for neither “geometry tokens” nor “exemplar tokens” to be present (e.g. when only a text prompt is used). The geometry or exemplar tokens attend to each other via self-attention and also cross-attend to the frame-embeddings of the corresponding (unconditioned) frame from the image encoder.

Fusion Encoder. The text and geometry/exemplar tokens together constitute the *prompt tokens*. The fusion encoder accepts the unconditioned frame-embeddings and conditions on prompt tokens using a stack of 6 transformer blocks with self- and cross-attention (to prompt tokens) layers followed by an MLP. We use vanilla self-attention operations. The output of the fusion encoder are the *conditioned* frame-embeddings.

Decoder. The decoder architecture follows Carion et al. (2020); Kamath et al. (2021) as a starting point and is a stack of 6 transformer blocks. K learned *object queries* (not to be confused with *prompts*) self-attend to each other and cross attend to the prompts tokens (made up of text and geometry/exemplar tokens) and conditioned frame-embeddings, followed by an MLP. We use box-to-pixel relative position bias (Lin et al., 2023) in the cross-attention layers attending to the conditioned frame-embeddings.

Following standard practice in stronger DETR variants, we use iterative box refinement (Zhu et al., 2020), look-forward-twice (Zhang et al., 2022a) and hybrid matching (Jia et al., 2022) and Divide-And-Conquer (DAC) DETR (Hu et al., 2023). By default, we use $K = 200$ object queries. Bounding boxes and scores are predicted using dedicated MLPs and accept the object queries as input.

Presence Head. Classifying each object in isolation is often difficult, due to insufficient information, and may require contextual information from the rest of the image. Forcing each object query to acquire such global awareness is however detrimental, and can conflict with the localization objectives that are by nature very local. To address this, we propose decomposing the classification problem into

1566 two complementary components: a global-level classification that determines object presence within
 1567 the entire image, and a local-level localization that functions as foreground-background segmentation
 1568 while preventing duplicate detections. Formally, we add the following probabilistic structure: instead
 1569 of predicting $p(\text{query}_i \text{ matches NP})$ directly, we break it down as follows:

$$1570 \quad p(\text{query}_i \text{ matches NP}) = p(\text{query}_i \text{ matches NP} \mid \text{NP appears in image}) \cdot p(\text{NP appears in image}).$$

1572

1573 To compute $p(\text{NP appears in image})$, we use a *presence token*, which is added to our decoder and
 1574 then fed through an MLP classification head. Crucially, the presence score is shared by all object
 1575 queries. The per-query classification loss is kept as usual, but to account for the decomposition, we
 1576 only compute it when the NP is present in the image (see §B.1 for ablations on supervision strategy).
 1577 The same decomposition is applied to the semantic segmentation head, where we reuse the same
 1578 presence score, and train the binary mask head only on the positive examples.

1579

Besides being more robust to false positives, decomposing the prediction in this manner is also more
 1580 *flexible*, e.g. in typical counting tasks, we already know the NP is present in the image and instead want
 1581 to know how many instances are present - in this case we can simply set $p(\text{NP is present in frame}) =$
 1582 1. The presence token is concatenated with the object queries in all operations, but is excluded from
 1583 DAC.

1584

We also learn 4 *geometric* queries. Their function is similar to the 4 geometric queries in SAM
 1585 1 and 2 (where they were called “output tokens”) and are used to perform the PVS on individual
 1586 image or video frames during the stags-2,-3 of training, see §D.4.1. The prompts are provided by the
 1587 “geometry tokens” in the form of visual prompts. The presence score is set to 1 when performing the
 1588 PVS task on a single frame, as the target is *known* to be present in the frame.

1589

Segmentation Head. The segmentation head is adapted from MaskFormer (Cheng et al., 2021).
 1590 Semantic segmentation and instance segmentation share the same segmentation head. The *conditioned*
 1591 features from the fusion encoder are used to produce semantic segmentation masks, while instance
 1592 segmentation additionally uses the decoder’s output object queries. “Multi-scale” features are
 1593 provided to the segmentation head using SimpleFPN (Li et al., 2022d), since the vision encoder is a
 1594 (single-scale) ViT.

1595

Handling Ambiguity. Experimentally, if we train a SAM 3 model without handling ambiguities as
 1596 described in §2 in any way, we observe that the model tends to predict several valid but conflicting
 1597 interpretations of the phrase. This is expected; if in our training dataset a given phrase has two
 1598 distinct interpretations, and roughly half the data is annotated assuming the first one, while the other
 1599 half follows the second one, then the solution that minimizes the training loss is to output both
 1600 interpretations with 50% confidence. However, this behavior is undesirable for end-users, because it
 1601 produces conflicting, sometimes overlapping masks.

1602

To address this issue, we add an ambiguity head to our model. Similar to SAM 1 and 2, this head is a
 1603 mixture of experts, where we train in parallel K experts, and only supervise the expert that gets the
 1604 lowest loss (winner-takes-all). We find that $K = 2$ performs the best and that it is more difficult to
 1605 train $K > 3$ experts due to mode collapse.

1606

For a mixture of K experts, each producing an output y_k with loss \mathcal{L}_k , the mixture loss is a weighted
 1607 average:

1608

1609

$$\text{Loss: } \mathcal{L}_{\text{MoE}} = \sum_{k=1}^K p_k \mathcal{L}_k \quad \text{Gradient: } \frac{\partial \mathcal{L}_{\text{MoE}}}{\partial \theta_j} = p_j \frac{\partial \mathcal{L}_j}{\partial \theta_j}.$$

1610

In our winner-takes-all variant, only the expert with the lowest loss receives gradient:

1611

1612

1613

1614

1615

1616

1617

1618

1619

$$\text{Loss: } k^* = \arg \min_k \mathcal{L}_k, \quad \mathcal{L}_{\text{WTA}} = \mathcal{L}_{k^*}$$

$$\text{Gradient: } \frac{\partial \mathcal{L}_{\text{WTA}}}{\partial \theta_j} = \begin{cases} \frac{\partial \mathcal{L}_{k^*}}{\partial \theta_j}, & \text{if } j = k^*, \\ 0, & \text{otherwise.} \end{cases}$$

1620 Backpropagating the loss only through the expert which received the minimal loss allows each expert
 1621 to specialize to one kind of interpretation. This behavior is illustrated in Fig. 9.
 1622



1623 (a) Original image (b) Prediction by Expert 1 (c) Prediction by Expert 2
 1624
 1625
 1626
 1627
 1628
 1629
 1630

1631 Figure 9: Two interpretations of the noun phrase “large circular shape” learned by two Experts (SA-1B image).
 1632
 1633

1634 While this strategy allows experts to specialize, it does not explicitly select which expert should be
 1635 used at inference time. To resolve this, we train a classification head that predicts the expert that has
 1636 the highest probability of being correct. The classification head is trained in a supervised fashion
 1637 with a cross entropy loss, by predicting which expert obtained the minimal loss during training. The
 1638 Ambiguity head adjusts only the classification logits, leaving masks, boxes, and presence scores
 1639 unchanged. We train it on top of a frozen SAM 3 model.

1640 Finally, to detect overlapping instances, we compute the Intersection-over-Minimum (IoM) between
 1641 masks. IoM is more effective than Intersection-over-Union (IoU) for identifying nested instances.
 1642 With the ambiguity head, we obtain a *15% reduction* in overlapping instances.
 1643

1644 D.3 VIDEO IMPLEMENTATION DETAILS

1645 As illustrated in §3, tracking in videos can suffer from ambiguities in mask propagation, false
 1646 predictions from the detector, or limitations of IoU-based matching in crowded scenes with highly
 1647 overlapping objects. In this section, we present the details of the temporal disambiguation strategies
 1648 used to address these challenges. We begin by introducing the notation used throughout this section.

1649 Let \mathcal{D}_τ and $\hat{\mathcal{M}}_\tau$ denote the set of detector outputs and the set of tracker’s predicted masks on frame
 1650 τ respectively. We define a frame-wise matching function $\Delta_i(\tau)$ for a masklet i on frame τ as
 1651

$$\Delta_i(\tau) = \begin{cases} +1, & \text{if } \exists d \in \mathcal{D}_\tau \text{ s.t. } \text{IoU}(d, \hat{\mathcal{M}}_\tau^i) > \text{iou_threshold} \\ -1, & \text{otherwise,} \end{cases}$$

1652 where $\hat{\mathcal{M}}_\tau^i$ is the predicted output mask of object i on frame τ . In addition, we define a Masklet
 1653 Detection Score (MDS) over an interval $[t, t']$ as $S_i(t, t') = \sum_{\tau=t}^{t'} \Delta_i(\tau)$. This score measures how
 1654 a masklet is consistently matched to a detection within a temporal window. We also denote the first
 1655 frame in which object i appears as t_{first}^i .
 1656

1657 D.3.1 TRACK CONFIRMATION DELAY

1658 To reduce spurious and duplicate masklets, we adopt a delayed output confirmation strategy. Specifically,
 1659 the output at frame τ is finalized only after observing frames up to frame $\tau + T$. This delay
 1660 provides temporal context for validating candidate masklets before outputting their masks. We find
 1661 that $T = 15$ achieves good accuracy. During this delay period, we apply the two following criteria to
 1662 remove unconfirmed or duplicate masklets.
 1663

1664 **Removal of Unconfirmed Masklets.** Candidate masklets are considered unconfirmed within the
 1665 confirmation window $[t, t + T]$ if their MDS is below a threshold, $S_i(t, t + T) < K$, and the masklet
 1666 first appears within the window $t_{first}^i \geq t$. If both conditions are satisfied within the confirmation
 1667 delay, we remove the masklet from the tracker’s state. We choose $K = 0$, requiring that the masklet
 1668 has to be matched to a detection for at least half of the frames within the confirmation delay period to
 1669 be confirmed. This strategy helps reject some false positive detections and not track them.
 1670
 1671
 1672
 1673

Removal of Duplicate Masklets. If the tracker temporarily fails to predict a mask for an object in some frames, but the detector continues to detect the object during those frames, this can lead to the creation of a new masklet for the same object. As a result, two masklets may end up tracking the same object: the original (older) masklet, and a new masklet that is initiated during the period when the tracker missed the object. To resolve this issue, during the confirmation delay period, if two masklets consistently overlap with the same detection, we remove the one that started later. Specifically, two masklets i, j are considered duplicates on frame τ if there exists a detection $d \in \mathcal{D}_\tau$ such that $\text{IoU}(\hat{\mathcal{M}}_\tau^i, d) \geq \text{iou.threshold}$ and $\text{IoU}(\hat{\mathcal{M}}_\tau^j, d) \geq \text{iou.threshold}$. If the two masklets i and j are found to be duplicates for at least $\lceil T/2 \rceil$ frames, we remove the one with the latest first appearance t_{first} only if it first appeared within the confirmation window $[t, t + T]$. Empirically, we find that using $\text{iou.threshold} = 0.1$ gives the best results.

D.3.2 MASKLET SUPPRESSION

For confirmed masklets that were not removed during the confirmation delay, we apply an additional suppression step: if a masklet’s MDS over its entire lifetime falls below zero at any frame τ (i.e. $S_i(t_{first}, \tau) < 0$), we suppress its output by zeroing out its mask. However, we retain the masklet in the tracker’s state, allowing for the possibility that the object may be confirmed in future frames. This strategy primarily addresses ambiguous detections, such as objects entering the scene near the boundary. For example, if only a person’s hands are visible as they enter the frame, the detector may be unable to determine whether the object matches the text prompt (e.g., impossible to distinguish between a man and a woman). In such cases, if the detector subsequently fails to detect the object after it fully enters the scene, the masklet suppression criterion ensures that these masklets are suppressed, unless they are consistently matched with new detections.

D.3.3 RE-PROMPTING

Periodic Re-Prompting. In challenging scenarios involving occlusions or visually similar distractor objects, the tracker may lose track of the target object. To address such tracking failures, we periodically re-prompt the tracker using the latest detection outputs. Specifically, on every N -th frame τ , we compare each detection $d \in \mathcal{D}_\tau$ with the tracker’s current predictions $\hat{\mathcal{M}}_\tau$. If a detection d has a high overlap with the tracker’s prediction (i.e., $\text{IoU}(d, \hat{\mathcal{M}}_\tau^i) \geq 0.8$) and both the detection score and the masklet prediction score exceed a confidence threshold of 0.8, we re-initialize the tracker for that object using the detection output mask. We observed that re-prompting is most effective on frames where objects are not occluded and are fully visible, which motivates our choice of high confidence thresholds. In our experiments, we set $N = 16$ by default. This periodic re-prompting helps the tracker recover from temporary failures and maintain accurate object tracking throughout the video.

Detection-Guided Re-Prompting. In cases where the tracker’s predictions may drift and its predicted masks become leaky, we employ the detectors’ outputs. For each frame τ , we compare every detection $d \in \mathcal{D}_\tau$ with the tracker’s current predictions $\hat{\mathcal{M}}_\tau$. If the highest-matching detection d has a low bounding box IoU (i.e., $\text{IoU}_{bbox}(d, \hat{\mathcal{M}}_\tau^i) < 0.85$) with the corresponding tracker prediction $\hat{\mathcal{M}}_\tau^i$, we recondition the tracker for that object using the latest detector output. This approach ensures that the tracker remains synchronized with reliable detection results.

The impact of these strategies is ablated in Tab. 38, and they show quantitative improvements across our evaluation sets.

D.4 MODEL TRAINING

D.4.1 TRAINING STAGES

SAM 3 is trained in 4 stages, with each stage introducing new capabilities or refining existing capabilities. Tab. 20 lists the data used in each stage.

Stage 1. Perception Encoder (PE) pre-training (Bolya et al., 2025), which pre-trains the image and text encoders with 5.4 billion image-text pairs. In addition to broad concept coverage, this stage is key for robustness (see §B.1). Since the vision encoder has to support multiple tasks (while also not being too large) we opt for an “L+” size; The vision and text encoders are transformers with 450m

1728 and 300m parameters respectively. We largely follow [Bolya et al. \(2025\)](#), but do not use distillation
 1729 and do not perform video fine-tuning in this stage.
 1730

1731 **Stage 2.** This stage is for detector pre-training and trains the (image-level) detector as well as the
 1732 vision and text encoders with large-scale image segmentation data (including video frames as images).
 1733 This stage uses both psuedo-labelled and human-annotated data, see Tab. 20. The main goal of this
 1734 stage is broad concept coverage of (image, noun phrase, masks) tuples. At the end of this stage, the
 1735 model is able to do open-vocabulary object detection, instance and semantic segmentation across
 1736 many domains fairly well.
 1737

1738 An additional goal of this stage is to prepare the base model for tasks in subsequent stages. To prepare
 1739 for the PCS task, (image, noun phrase) pairs are randomly ($p = 0.2$) converted into visual queries
 1740 (i.e. noun phrase is dropped) or augmented with input bounding boxes ($p = 0.2$).
 1741

1742 Besides training for the PCS task, in this stage, the model is also pre-trained on the *geometric*, PVS
 1743 task. This is done by adding 4 decoder queries specific to this task following the design of SAM 1,
 1744 SAM 2. Training data includes images (e.g., SA-1B) and videos frames (e.g, SA-V), see Tab. 20; the
 1745 number of interactivity steps is restricted to 4 for efficiency. We largely follow the settings from [Ravi
 1746 et al. \(2024\)](#), but use the Align loss ([Cai et al., 2024](#)) in lieu of the IoU prediction loss, co-opting the
 1747 classification head for object queries for this task.
 1748

1749 We train for $\sim 95k$ iterations with a batch size of 896 with 5k warm up and cooldown steps using
 1750 AdamW ([Loshchilov & Hutter, 2019](#)). We apply layer-wise learning rate decay ([Clark et al., 2020](#)) of
 1751 0.9 to the vision encoder. We use a reciprocal square-root schedule ([Zhai et al., 2022](#)) and weight
 1752 decay of 0.1. We use an initial learning rate of $5e-4$, $1e-4$ for vision and text encoder and $1e-3$ for all
 1753 other components. For boxes, we use L_1 and gIoU losses with weights of 5 and 2. Classification loss
 1754 uses a weight of 100 and focal and dice losses use weights of 200 and 10 respectively. The encoder
 1755 and decoder use a dropout of 0.1.
 1756

1757 **Stage 3.** This stage further trains the model with the highest-quality human annotated image
 1758 segmentation data, expands the interactivity capabilities and introduces post-training to improve
 1759 detection performance.
 1760

1761 Specifically, in terms of interactivity, (a) in the PVS task, the number of interactivity steps is increased
 1762 to 7 and (b) interactivity is introduced into the PCS task, where positive or negative exemplars are
 1763 provided based on model error. We iteratively sample box prompts to mimic the real user policy.
 1764 Positive boxes are sampled from false negative errors, and we prompt their corresponding ground-
 1765 truth boxes. Negative boxes are sampled from high-confidence false positive predictions that do not
 1766 have significant overlap with ground truths. At each iteration, the box inputs are added on top of the
 1767 previous ones. If both a valid positive and negative box exist, we randomly select one; if no valid
 1768 candidates are available, no additional prompt is given. The process is repeated for 5 iterations.
 1769

1770 The expanded interactivity in the PCS and PVS in this stage significantly slow down training, but the
 1771 extensive pretraining with limited interactivity for the PVS and no interactivity for PCS (but using
 1772 image exemplars together with text prompts) prepares the model well to ensure that a short stage 3 is
 1773 sufficient.
 1774

1775 This stage retains only the highest quality, exhaustivity verified data (e.g., SA-Co/SYN is dropped)
 1776 and introduces a presence token (and presence loss) to better model *presence* of target segments and
 1777 their location *location* greatly increasing the precision of the model. The presence loss is a binary
 1778 cross-entropy loss with weight of 20. All learning rates are lowered by a factor of 0.025. We train for
 1779 $\sim 5k$ iterations with a batch size of 512, with other settings identical to stage 2.
 1780

1781 **Stage 4.** The tracker decoder is trained on top of the frozen backbone. Freezing the backbone at
 1782 this stage is made possible by pre-training on VOS data in previous stages at the video frame level.
 1783 This stage retains the strong spatial grounding of the previous stage and focuses on spatial-temporal
 1784 tracking without degrading other capabilities. As in [Ravi et al. \(2024\)](#), we further fine-tune the tracker
 1785 with a longer temporal context using 16-frame and 32-frame videos. We largely follow [Ravi et al.
 1786 \(2024\)](#)'s recipe. We use a batch size of 512, training for $\sim 190k$ iterations using a cosine schedule
 1787 with a linear warmup of 1k iterations.
 1788

	Dataset	Ingested As	Train	Test	Stage 1	Stage 2	Stage 3	Stage 4
<i>Promptable Visual Segmentation- In Images</i>								
1784	SA-1B	Image	✓		✓	✓	✓	
1785	SA-Co/VIDEO	Frames	✓		✓	✓	✓	
1786	SA-Co/VIDEO-EXT	Frames	✓		✓	✓	✓	
1787	SA-37	Image		✓	✓	✓	✓	
<i>Promptable Visual Segmentation- In Videos</i>								
1789	SA-Co/VIDEO	Video	✓				✓	
1790	SA-Co/VIDEO-EXT	Video	✓				✓	
1791	SA-V val	Video		✓			✓	
<i>Promptable Concept Segmentation- In Images</i>								
1793	SA-Co/SYN	Image	✓		✓			
1794	SA-Co/HQ	Image	✓		✓	✓	✓	
1795	SA-Co/EXT	Image	✓		✓	✓	✓	
1796	SA-Co/VIDEO	Frames	✓		✓	✓	✓	
1797	SA-Co/Gold	Image		✓	✓	✓	✓	
1798	SA-Co/Silver	Image		✓	✓	✓	✓	
1799	SA-Co/Bronze, SA-Co/Bio	Image		✓	✓	✓	✓	
<i>Promptable Concept Segmentation- In Videos</i>								
1800	SA-Co/VEval	Video		✓			✓	

Table 20: Dataset usage across different tasks and training stages.

D.4.2 ADDITIONAL TRAINING SETTINGS

Data augmentation. For the PCS task, we apply the following transformations:

- **Geometric:** We use some cropping and resizing to vary the aspect ratios and help with small objects. The input resolution of our model is always a fixed square (usually 1008×1008). During evaluation, the images are resized to this size, without respecting their aspect ratio. During training, we apply our augmentations, and pad if the resulting size is smaller than 1008×1008 . We found it important to randomly distribute the padding on all sides, to avoid creating biases towards one particular region of the image. If the dataset does not contain notions of left and right, we also apply random horizontal flips.
- **Semantic:** When training on datasets with a closed vocabulary, we leverage our mapping to wikidata to further enhance the training. There are three main ways we can leverage the ontology: to sample synonyms, which expand the vocabulary of the model, to sample negatives (typically, if the dataset is exhaustively annotated, we can sample any node in the graph that corresponds in to a category and is not present in the image) and to ensure the hierarchy closure of the concepts (for example, if we have some annotations for “canoe” and “boat” in the same image, we need to make sure that all the “canoe” objects are also labeled as “boat” since a canoe is a type of boat).
- **Safety:** To prevent the model from randomly making predictions for unsafe concepts, we randomly sample some of them at train time and add them as negatives. These concepts mainly include slurs of all kinds. We also try to prevent the model from making predictions for subjective and non-visual adjectives, especially when applied to a person. This includes flattering ones (such as “a smart person”) as well as derogatory ones (such as “a dull person”)
- **Mosaics:** On some datasets, we further increase the complexity of the images by doing mosaics (Bochkovskiy et al., 2020). The maximal grid size of our mosaics is 3×3 , and we sample any configuration that is at most that, including irregular ones, as long as the constituents are still square. For example, in a 3×3 regular grid, we can have a large image that effectively covers a 2×2 area, and use 1×1 for the remaining 5 slots. Unifying different images can be tricky in an open vocabulary setting, since there is no guarantee that concepts are exhaustively annotated. For example, if one image has a car annotated, but the second does not (neither as positive nor negative), then we do not know if the second image has a car or not, and thus could create some labeling noise. To avoid this, we only mosaic datasets that have low chance of such missing annotations (either closed vocabulary ones, or some created with specific mining patterns). To merge annotations, we again rely on the wikidata mapping if available, otherwise rely on plain-text queries to merge appropriately.

1836 E DATA ENGINE DETAILS

1838 The overview of the SAM 3 data engine’s components is shown in Fig. 4. In this section we provide
 1839 further details of how each component is implemented in the image (phases 1-3) and video (phase 4)
 1840 versions of the engine. The datasets collected in each phase and the improvements in the model’s
 1841 performance are summarized in Tab. 21.

1843 E.1 MEDIA POOL

1845 The media (image and video) pool consists of many sources with varying visual domains, from web-
 1846 scraped data to datasets collected for specialized domains such as art, food, or driving. Tab. 28 lists
 1847 the datasets used to mine media for each subset of the SA-Co training data. The web-scraped images
 1848 and alt captions are sourced from MetaCLIP (Xu et al., 2024b), a curated version of CommonCrawl.
 1849 We further expand coverage by mining media from a large pool with the help of a curated ontology.
 1850 Compared to previous works such as OWL-V2 (Minderer et al. (2024)) which mainly rely on
 1851 uncurated web-scraped data, our target mining strategy resulted in coverage of 12 media domains.

1853 E.2 SA-CO ONTOLOGY

1854 To track and improve the coverage and overall distribution of concepts in our data, we build a custom
 1855 SA-Co ontology of visual concepts from Wikidata (Vrandečić & Krötzsch, 2014), which covers a
 1856 comprehensive set of entities and offers hierarchical information with its graph data structure. We
 1857 manually select high-level Wikidata nodes (e.g., Human, Mammals) and recursively include all of
 1858 their descendants. The resulting 22.4 million nodes are classified into 17 top-level categories (e.g.
 1859 animal, furnishing & home) and 72 sub-categories (e.g., birds, home appliance). The full list of
 1860 categories and Wikidata node counts are shown in Tab. 22. We further develop a mapping process
 1861 that can map an arbitrary NP to a node in the SA-Co ontology by leveraging a retrieval model
 1862 (Sentence-BERT) to source candidate nodes and an AI annotator as judge (Llama 3.2) to select the
 1863 closest match. This mapping is used to track the distribution of nodes in the dataset (see Fig. 10) as
 1864 well as to create negative phrases (see below for details).

1865 E.3 PHASE 1: HUMAN VERIFICATION

1866 **Data Mining.** During this phase, we randomly sample images from MetaCLIP.

1867 **Proposing NPs.** We generate image-level captions using the BLIP-2 captioner (Li et al., 2023b)
 1868 followed by the spaCy parser (Honnibal et al., 2020) to parse the caption into NPs.

1869 **Proposing Masks.** We prompt an off-the-shelf open-vocabulary detector, FIBER (Dou et al., 2022)
 1870 or OWL-V2 (Minderer et al., 2024) with the noun phrase and use the resulting boxes to prompt
 1871 SAM 2 to generate mask proposals.

1872 **Verification (Human).** Verification of mask proposals consists of two tasks which can be performed
 1873 by human or AI annotators: mask quality verification and mask exhaustivity verification. In Phase 1,
 1874 verification is done by humans only. Each human verifier works exclusively on one task type.

- 1875 • **Mask Verification (MV).** Given a triplet of an image, a noun phrase and a set of candidate masks for
 1876 that phrase, the task is to accept or reject each of the masks. A mask is accepted if it matches the
 1877 given noun phrase and is high quality (no holes, coverage issues, etc.) If the mask is unrelated to
 1878 the phrase, or low quality, it is rejected.

	SA-Co/HQ			SA-Co/SYN			SA-Co/EXT			SA-Co/VIDEO			SAM 3 performance		
	# Images	# Image-NPs	# Annotation Domains	# Images	# Image-NPs	# Images	# Image-NPs	# Videos	# Video-NPs	SA-Co/Gold (CGF ₁)	SA-Co/Silver	SA-Co/VEval (val pDetA)			
Phase 1	1.2M	4.3M	1	0	0	0	0	0	0	-	-	-			
Phase 2	2.4M	122.2M	5	0	0	0	0	0	0	38.4	18.5	-			
Phase 3	1.6M	19.5M	15	39.4M	1.7B	9.3M	136.6M	-	-	65.0	59.0	45.7			
Phase 4	-	-	-	-	-	-	-	52.5K	134.3K	65.0	59.0	48.1			

1889 Table 21: Data engine phases and SAM 3 progress.

1890	1. animals	2.3M	6. electronics	10.2K	12. object parts	101.9K
1891	insects & crustaceans	1.7M	electronics	6.9K	body parts	75.8K
1892	molluscs	188.4K	cameras	3.3K	other object parts	26.1K
1893	other animals	166.5K	7. equipments	14.9K	13. other products	3.5K
1894	fish & other chordates	85.7K	military equipments	10.2K	other products	2.7K
1895	birds	52.4K	sport equipments	2.0K	celebration supplies	384
1896	mammals	38.3K	safety equipments	1.2K	animal-related products	359
1897	reptiles	28.2K	medical equipments	1.1K	tobacco products	51
1898	echinoderms	23.0K	agricultural machinery	458	14. patterns & material	896.6K
1899	amphibians	14.2K	8. fashion & beauty	7.5K	material	885.6K
1900	artworks	3.1M	fashion	3.9K	patterns & shapes	10.9K
1901	collectibles	10.2K	beauty & healthcare products	3.7K	15. plants & fungi	1.5M
1902	religious objects	9.1K	9. food & drinks	33.1K	plants	1.1M
1903	flags	8.3K	dishes	12.9K	fungi	376.4K
1904	musical instruments	4.9K	other food	6.7K	16. tools & appliances	14.5K
1905	gemstones	526	fruits	6.6K	other appliances	6.5K
1906	art material	438	drinks	6.3K	toys	3.0K
1907	3. buildings & locations	2.7M	vegetables	621	tools	1.8K
1908	places	2.4M	10. furnishing & home	2.7K	kitchenware	1.6K
1909	geographical features	343.2K	furnishing	1.3K	containers	945
1910	4. celestial	9.2K	home appliances	486	light sources	672
1911	meteorological phenomena	5.3K	stationery	472	17. transportation	258.1K
1912	space related	3.1K	household supplies	417	watercraft	178.9K
1913	light related	734	11. human	11.3M	land vehicles	41.6K
1914	5. documents & ocr	201.3K	humans	11.0M	aircraft	27.4K
1915	glyphs	173.4K	occupations	140.8K	other vehicles	6.4K
1916	logos	21.0K	fictional characters	87.8K	transport infrastructures	3.7K
1917	documents	6.0K	gestures & expressions	158	construction machines	100
1918	cards	435				
1919	infographics	324				
1920	GUI & layout elements	135				
1921	maps	23				

Table 22: SA-Co ontology top-level categories and sub-categories with corresponding node counts in Wikidata.

- *Exhaustivity Verification (EV)*. All accepted masks from the verification task are sent to an exhaustivity check. Given an image, noun phrase, and any accepted masks that passed the previous mask verification for that phrase, the task is to decide whether or not the accepted masks (if any) exhaustively cover all instances of the phrase in the image. If there are unmasked instances of the phrase, annotators decide whether or not at least one of the remaining instances is separable, or if the remaining instances are too crowded together to separate. Phrases that are annotated as non-exhaustive from this step are sent to the correction task. Phrases that are annotated as exhaustive are directly sent to final annotations.

Correction. Human annotators are given the same input as the exhaustivity task: an image, noun phrase, and any (0 or more) accepted masks from the mask verification task. Annotators manually add individual masks for the unmasked instances of the noun phrase by prompting SAM 1 with clicks in a browser based tool. If there are non-separable occurrences of the phrase, annotators use special group masks to indicate that the mask covers more than a single instance. The output of the task is a complete set of instance and/or group masks covering all pixels in the image corresponding to the noun phrase. Noun phrases that are not present are submitted with no masks. If it is not possible to reach a complete set of masks due to mask complexity, the annotator rejects the job.

In each task, annotators are given the ability to reject the image-NP pairing if they decide the phrase is un-maskable as a set of objects (e.g “it”, “blue”) or if after research they are still unsure if it is present (e.g., fine-grained species of animals). Filtering out vague phrases and allowing annotators to be unsure increases the consistency and agreement in the resulting annotations.

E.4 PHASE 2: HUMAN + AI VERIFICATION

Data Mining. We use a retrieval model (including Perception Encoder, DINoV2, and MetaCLIPv2) for mining concepts that are challenging and not prevalent in the caption NPs from Phase 1. We

1944					
1945					
1946					
1947					
1948					
1949	<table border="1" style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th style="text-align: left; padding: 5px;">Input</th> <th style="padding: 5px;"></th> </tr> </thead> <tbody> <tr> <td style="text-align: center; padding: 10px;">Image input for AI verifier (left) and UI for human annotation (right)</td> <td style="text-align: center; padding: 10px;">  </td> </tr> </tbody> </table>	Input		Image input for AI verifier (left) and UI for human annotation (right)	
Input					
Image input for AI verifier (left) and UI for human annotation (right)					
1950					
1951					
1952					
1953					
1954					
1955					
1956					
1957					
1958					
1959					
1960	<table border="1" style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th style="text-align: left; padding: 5px;">AI Verifier Instructions</th> <th style="padding: 5px;"></th> </tr> </thead> <tbody> <tr> <td style="text-align: center; padding: 10px;"> <p>You are an expert annotator of object segmentation masks. For an image and a pre-defined label, you are given a mask and asked to evaluate the quality of the mask. Follow the following rules when rating the mask.</p> <ol style="list-style-type: none"> 1. Rate the mask as “Accept” when the label accurately describes the masked object and the mask covers the object with good boundaries. We do not need masks to be pixel-perfect for this task. However, the mask should cover all important parts of the object. 2. Rate the mask as “Accept as text” when the mask covers text and the label exactly matches the masked text. The mask should cover all important parts of the text (all specified letters/punctuation/etc.) 3. Rate the mask as “Flag label” when the label corresponds to Race, Ethnicity, Sexual orientation, Religion, Socio-economic status, Medical conditions, Disabilities, Derogatory terms/profanity and Animal phrases for a person. 4. Rate the mask as “Whole image” when the label corresponds to the entire image. Description that refers to the whole image may describe setting (e.g., inside, outside, day, night), type of media (e.g., flier, screenshot, photo), type of image (e.g., close up, an aerial view) and location (e.g., an airport, the woods, a bedroom). 5. Otherwise, rate the mask as “Reject”. <p>Please give your rating directly without any explanation. Now let’s start. In the given figure, the left half shows a fuchsia box highlighting the region of interest in the original image, and the right half shows a fuchsia mask overlaid on a zoom-in view of that region.</p> <p>Rate the fuchsia mask for the label “a computer monitor”: (A). Accept as text. (B). Flag label. (C). Reject. (D). Accept. (E). Whole image.</p> </td> <td style="text-align: center; padding: 10px;"></td> </tr> </tbody> </table>	AI Verifier Instructions		<p>You are an expert annotator of object segmentation masks. For an image and a pre-defined label, you are given a mask and asked to evaluate the quality of the mask. Follow the following rules when rating the mask.</p> <ol style="list-style-type: none"> 1. Rate the mask as “Accept” when the label accurately describes the masked object and the mask covers the object with good boundaries. We do not need masks to be pixel-perfect for this task. However, the mask should cover all important parts of the object. 2. Rate the mask as “Accept as text” when the mask covers text and the label exactly matches the masked text. The mask should cover all important parts of the text (all specified letters/punctuation/etc.) 3. Rate the mask as “Flag label” when the label corresponds to Race, Ethnicity, Sexual orientation, Religion, Socio-economic status, Medical conditions, Disabilities, Derogatory terms/profanity and Animal phrases for a person. 4. Rate the mask as “Whole image” when the label corresponds to the entire image. Description that refers to the whole image may describe setting (e.g., inside, outside, day, night), type of media (e.g., flier, screenshot, photo), type of image (e.g., close up, an aerial view) and location (e.g., an airport, the woods, a bedroom). 5. Otherwise, rate the mask as “Reject”. <p>Please give your rating directly without any explanation. Now let’s start. In the given figure, the left half shows a fuchsia box highlighting the region of interest in the original image, and the right half shows a fuchsia mask overlaid on a zoom-in view of that region.</p> <p>Rate the fuchsia mask for the label “a computer monitor”: (A). Accept as text. (B). Flag label. (C). Reject. (D). Accept. (E). Whole image.</p>	
AI Verifier Instructions					
<p>You are an expert annotator of object segmentation masks. For an image and a pre-defined label, you are given a mask and asked to evaluate the quality of the mask. Follow the following rules when rating the mask.</p> <ol style="list-style-type: none"> 1. Rate the mask as “Accept” when the label accurately describes the masked object and the mask covers the object with good boundaries. We do not need masks to be pixel-perfect for this task. However, the mask should cover all important parts of the object. 2. Rate the mask as “Accept as text” when the mask covers text and the label exactly matches the masked text. The mask should cover all important parts of the text (all specified letters/punctuation/etc.) 3. Rate the mask as “Flag label” when the label corresponds to Race, Ethnicity, Sexual orientation, Religion, Socio-economic status, Medical conditions, Disabilities, Derogatory terms/profanity and Animal phrases for a person. 4. Rate the mask as “Whole image” when the label corresponds to the entire image. Description that refers to the whole image may describe setting (e.g., inside, outside, day, night), type of media (e.g., flier, screenshot, photo), type of image (e.g., close up, an aerial view) and location (e.g., an airport, the woods, a bedroom). 5. Otherwise, rate the mask as “Reject”. <p>Please give your rating directly without any explanation. Now let’s start. In the given figure, the left half shows a fuchsia box highlighting the region of interest in the original image, and the right half shows a fuchsia mask overlaid on a zoom-in view of that region.</p> <p>Rate the fuchsia mask for the label “a computer monitor”: (A). Accept as text. (B). Flag label. (C). Reject. (D). Accept. (E). Whole image.</p>					
1970					
1971					
1972					
1973					
1974					
1975					
1976					
1977					
1978					
1979					
1980					
1981					
1982					
1983					
1984	<table border="1" style="width: 100%; border-collapse: collapse;"> <thead> <tr> <th style="text-align: left; padding: 5px;">Mask Verification Result</th> <th style="padding: 5px;"></th> </tr> </thead> <tbody> <tr> <td style="text-align: center; padding: 10px;">(D). Accept.</td> <td style="text-align: center; padding: 10px;"></td> </tr> </tbody> </table>	Mask Verification Result		(D). Accept.	
Mask Verification Result					
(D). Accept.					
1985					
1986					

Table 23: **An example data point of Mask Verification (either human or AI verifier).** The AI-verifier is given two crops of the image, a zoomed-out view where the object is highlighted via a box and a zoomed-in view where the mask is highlighted. This allows better visualization of small objects, and avoids color confusion from mask overlay. The AI-verifier instructions are a condensed version of the annotation guidelines given to human annotators. Human annotators are also given options to reject the phrase at the image-level as vague. For AI verifier, we use the model output logits of the choice tokens (e.g., A/B/C/D/E) as the logits for the corresponding labels, and soft-max over the logits to get their probabilities.

1993
1994
1995
1996
1997

1998	
1999	
2000	
2001	
2002	
2003	<p>Input</p>  <p>Image input for AI verifier (left) and UI for human annotation (right)</p>
2004	
2005	
2006	
2007	
2008	
2009	
2010	
2011	
2012	
2013	
2014	<p>AI Verifier Instructions</p> <p>You are an expert annotator of object detection. For an image and a pre-defined label, you are given some boxes and asked to evaluate if the boxes are exhaustive. Follow the following rules when rating the boxes.</p> <ol style="list-style-type: none"> 1. Rate the boxes as “Accept” when there are no undetected objects that match the label or the label is vague (i.e., “rent”, “it”) or malformed (heavily misspelled). 2. Rate the boxes as “Reject” when there are undetected objects and at least some undetected instances can be separated. 3. Rate the boxes as “Reject but unseparated” when there are undetected objects but cannot be separated. This is common with groups where it is not possible to tell the undetected instances apart or distinguish them from others. 4. Rate the boxes as “Flag label” when the label corresponds to Race, Ethnicity, Sexual orientation, Religion, Socio-economic status, Medical conditions, Disabilities, Derogatory terms/profanity and Animal phrases for a person. 5. Rate the boxes as “Whole image” when the label corresponds to the entire image. Description that refers to the whole image may describe setting (e.g., inside, outside, day, night), type of media (e.g., flier, screenshot, photo), type of image (e.g., close up, an aerial view) and location (e.g., an airport, the woods, a bedroom). 6. Rate the boxes as “Ungroundable / Vague / Unsure” when the label is an ungroundable or vague concept (e.g., “a corner”, “gap”) or when you are unsure whether or not the phrase is in the image. <p>Now let’s start. The given figure shows one or multiple red boxes highlighting the region of interest in the original image, evaluate the red box for the label “a computer monitor” and select your answer from the following options: (A). Reject but unseparated. (B). Whole image. (C). Reject. (D). Ungroundable / Vague / Unsure. (E). Accept. (F). Flag label.</p>
2015	
2016	
2017	
2018	
2019	
2020	
2021	
2022	
2023	
2024	
2025	
2026	
2027	
2028	
2029	
2030	
2031	
2032	
2033	
2034	
2035	
2036	
2037	
2038	<p>Exhaustivity Verification Result</p> <p>(E). Accept.</p>
2039	
2040	
2041	<p>Table 24: An example data point of Exhaustivity verification (either human or AI verifier). For AI verifier, objects to be highlighted via boxes in the image. If there are no candidate objects, the original image is used and “one or multiple red boxes” is replaced by “zero red box” in the text prompt. For AI verifier, we use the model output logits of the choice tokens (e.g., A/B/C/D/E/F) as the logits for the corresponding labels, and soft-max over the logits to get their probabilities. The presence score from the EV AI verifier is defined as $1 - \text{Prob}(\text{Accept} \text{no boxes as input})$. $\text{Prob}(\text{Accept} \text{no masks as input})$ is the probability of Accept (no missing objects) given zero detections as input, which is equivalent to the probability of NO presence.</p>
2042	
2043	
2044	
2045	
2046	
2047	
2048	
2049	
2050	
2051	

2052 leverage our SA-Co ontology to determine the list of candidate concepts, followed by offline concept
 2053 indexing and online mining from MetaCLIP.
 2054

- 2055 • *Concept Selection.* We use a taxonomy-guided mining strategy to balance the overall ontological
 2056 distribution, expand concept coverage and enhance performance on long-tail and fine-grained
 2057 phrases. Two groups of concepts are selected from the SA-Co Ontology for targeted mining:
 2058 *Wiki-Common* are nodes judged by an LLM to be common concepts, *Wiki-FG* are all nodes from
 2059 the “sports equipment” and “food and drink” sub-graphs, chosen to test the model’s ability to
 2060 generalize to very fine-grained concepts like “kefir”, “pastille”, “kettlebell”.
- 2061 • *Offline Concept Indexing.* For every new concept, we collect reference images from Wikimedia and
 2062 compute their K-dimensional embedding offline. We aggregate the embeddings from all reference
 2063 images resulting in a single embedding per concept. We repeat the process across all N concepts
 2064 resulting in an $N \times K$ dimensional offline index.
- 2065 • *Online Mining.* Relevant images for each concept are retrieved using both image and text based
 2066 mining. With image-based retrieval, we compute the embedding on every image, run KNN on the
 2067 offline concept index followed by top-k sampling, and apply a threshold before mapping it to a
 2068 specific concept. With text-based retrieval, we compute CLIP based similarity scores between the
 2069 text embedding from input concepts and image embeddings from the corpus and apply a threshold
 2070 before mapping the image to a specific concept.

2071 The following additional mining strategies are used to further refine the selection.
 2072

- 2073 • *Image-Type Balancing.* Web datasets are usually dominated by a few types of images such as ads
 2074 or product photos. To avoid over-representation of certain image types, we use a MLLM (Llama
 2075 3.2) and prompt it zero-shot to classify an image into image types (such as ads, product photos,
 2076 indoor and outdoor scenes, infographics), and sample based on a type-agnostic probability.

2077 **Proposing NPs.** We improve this step to generate higher-quality and more diverse noun phrases.
 2078

- 2079 • *Image-Level Captioner and Parser.* We use an image captioning model (Llama 3.2) to generates
 2080 image-level captions and a phrase parser (Llama 3.1) that proposes noun phrases given the caption.
 2081 The Llama 3.2 captioning model improved concept recall compared to BLIP-2 from Phase 1. The
 2082 phrase parser is fine-tuned for this task and significantly outperforms its zero-shot model variant
 2083 and spaCy parser.
- 2084 • *Removing Non-Groundable Phrases.* The parser can generate non-specific phrases such as “it”,
 2085 “them” or hard to segment phrases such as “middle”. To address this, we use another AI verifier
 2086 (MLLM) that is fine-tuned to classify such cases and remove them from the rest of the pipeline.
- 2087 • *NP Balancing.* We employ heuristics to avoid collecting too many frequent or easy objects. We
 2088 remove NPs if the data engine has already annotated enough instances, if the SAM 3 has high
 2089 accuracy when prompted with the NP, and based on a fixed list (e.g. that occur frequently, are
 2090 harmful). From Phase 3 we rely on AI verifiers to remove easy cases.
- 2091 • *Cleaning NPs.* We singularize noun phrases, deduplicate nearly-identical ones, and remove
 2092 possessives.
- 2093 • *Hard Negative Proposal.* A hard negative phrase generator proposes image-level negative phrases,
 2094 i.e. those that do not exist in the image and are adversarial to SAM 3. Given verified positive NPs
 2095 (i.e. that exist in the image), negative NPs are *proposed* and then *checked for adversariality*.
- 2096 – *Proposal.* The proposal of hard negatives is done in two ways. The first approach maps every
 2097 positive NP to a node in the SA-Co ontology, then navigates the ontology graph to find sibling,
 2098 cousin, or uncle nodes corresponding to different but related concepts. For example, the noun
 2099 phrase “gray Siamese cat” maps to the node “Siamese cat”, which could result in negative
 2100 candidates like “tabby cat” (sibling), “dog” (uncle), or “Chihuahua” (cousin). The second
 2101 approach relies on an MLLM (Llama 4), which proposes visually similar negatives for every
 2102 positive NP.
- 2103 – *Check for Adversariality.* Once the negative NPs are proposed, they are filtered to retain only
 2104 those adversarial to the current SAM 3 version. For each negative NP candidate, predictions
 2105 from SAM 3 are obtained. If the set of predictions is empty, the candidate is discarded. If the
 model predicts one or more objects, these predictions are compared to the original segmentation

2106 masks of the corresponding positive NP. If the overlap between the negative NP predictions and
 2107 the positive NP annotations exceeds a certain threshold, the negative NP is retained as a hard
 2108 negative. This final check is necessary because initial proposals may not be true negatives and
 2109 instead may be only negatives relative to the existing positive NPs (i.e. the object could still be
 2110 present somewhere else in the image).

2111
Proposing Masks. We prompt SAM 3 with the set of positive and negative phrases to produce
 2112 candidate instance and semantic masks for the image. For pseudo-annotating domains with fine-
 2113 grained concepts that SAM 3 fails on (e.g., *Zanclus cornutus*), we identify the relevant coarse-grained
 2114 concept that SAM 3 works well on (e.g., *frog*), and use this as the prompts to generate masks. We
 2115 deduplicate masks generated per NP based on a IoU metric. These noisy pseudo-labels undergo
 2116 further cleaning by both human and AI annotators.
 2117

2118 **Verification (Human+AI).** We train “AI verifiers” to perform the mask verification (MV) and
 2119 exhaustivity verification (EV) tasks. More specifically, we fine-tune Llama 3.2 [Dubey et al. \(2024\)](#)
 2120 on human annotated data collected during Phase 1 of the data engine for both tasks.

- 2121 • *Task Formulation.* Tab. 23 provides an example data point of the mask verification task: given an
 2122 (image, phrase, mask) triplet, we render the mask on top of the image as the image prompt, provide
 2123 the task guidance as text prompt, and use the human annotation (1 out of 5 choices) as output.
 2124 Each mask’s quality is evaluated independently from other masks for the same image-phrase pair.
 2125 Rendering tricks are used to better visualize small objects, and to avoid color confusion from mask
 2126 overlay. Tab. 24 provides an example data point of the exhaustivity verification task: given the
 2127 (image, phrase, masks) triplet, we render the bounding boxes of the masks on top of the image
 2128 and use this as the image prompt, provide the task guidance as the text prompt, and use the human
 2129 annotation (1 out of 6 choices) as the output.
- 2130 • *Evaluation.* We construct test sets for “AI verifiers” from jobs that were reviewed by multiple
 2131 human annotators for all SA-Co test sets. We leave one human annotation as human prediction,
 2132 and use the majority vote of the remaining human annotations as ground truth. This allows us to
 2133 compare human and AI verifiers’ accuracy.
- 2134 • *Training.* The training data of each task comes from not only the task itself, but also from the
 2135 Correction task. For example, each manually added mask is a good data point in the mask
 2136 verification task. Each exhaustively finished job in the Correction task results in a good data point
 2137 in exhaustivity verification task. We merge all training data for these two tasks together (over 200M
 2138 image-text pairs) to pre-train a foundational AI verifier, and then only use high quality human
 2139 annotated data from the task itself (around 10M scale) to fine-tune two AI verifiers, one for each
 2140 task.
- 2141 • *Result.* Thanks to the simplicity of these two tasks (MCQ tasks on image-text pairs) and the large
 2142 volume of training data from Phase 1, AI verifiers reach and even surpass human performance
 2143 on these two tasks, as shown in Tab. 25. We also evaluate the system of SAM 3 and AI verifiers
 2144 end-to-end on the PCS task, and the system always performs better than the single SAM 3 model,
 2145 as shown in Tab. 17.
- 2146 • *Generalization to new domains.* We also study the generalization ability of AI verifiers. For a
 2147 given new domain, the MV AI verifier is typically on par with human verifiers without any domain
 2148 specific data; the EV AI annotator is typically worse than human in a zero-shot evaluation, but can
 2149 reach human performance with only thousands of domain specific data points.

	Attributes	Crowded	Food&Drinks	Sports	Equip.	MetaCLIP	SA-1B	Wiki-Common	Average				
	MV	EV											
Human	72.3	81.2	72.9	82.4	76.8	76.7	79.2	87.3	72.4	79.5	72.3	73.8	
AI verifier	77.1	82.6	74.6	81.3	79.4	75.1	80.1	84.7	75.3	78.8	75.9	76.8	81.3

2156 Table 25: Human/AI verifier performance on mask verification (MV) and exhaustivity verification (EV) tasks
 2157

2158 As discussed in §B.4, using AI verifiers is effective and allows human annotators to focus on the
 2159 most challenging data points, i.e. those that have poor mask quality or missing masks. This approach
 more than doubles the throughput of the SAM 3 data engine. As both SAM 3 and AI verifier models

improve, more data can be exhaustively annotated using only SAM 3 and AI verifiers. This leads to increasingly higher throughput and ensures that human annotators only work on SAM 3 failure cases.

Correction. We perform manual correction wherever needed as described in phase 1.

E.5 PHASE 3: SCALING AND DOMAIN EXPANSION

Data Mining. We continue the data mining approaches from Phase 2 and scale to more novel domains. In addition, we target cases that are rare in web datasets and challenging for the model: crowded scenes with high object counts and images with very small objects. To mine such images, we rely on the SA-1B dataset with mask annotations and compute the “crowdedness” metric i.e. calculate IoU between pair of masks and then aggregate it over all pairs of masks. We also use statistics of the number of masks and mask area to identify images with high object counts and very small objects.

Proposing NPs. We continue leveraging the approach from phase 2. We also expand concept coverage to long-tail, fine-grained concepts by extracting NPs from each image’s alt-text where available and by mining concepts from the SA-Co ontology.

Proposing Masks. Unchanged from Phase 2.

Verification (Human+AI). We continue to use both human and AI verifiers as described in Phases 1 and 2 respectively, but primarily rely on AI verifiers to increase the data engine throughput.

Correction (Human). We perform manual correction wherever needed, as described in Phase 1. Annotators are asked to correctly mask all occurrences of the given concept in the image.

E.6 PHASE 4: VIDEO ANNOTATION

In Phase 4, we extend the data engine to video. We use the same high-level stages as the image version, but with video-specific implementation details which are described next.

Media Pool. We curate a pool of O(1M) hours of video from SA-V, SA-V internal, YouTube-1B and SA-FARI (wildlife cameras) datasets that covers diverse domains and a range of video quality.

Data Mining. To efficiently utilize human annotation resources, we developed aggressive data mining filters and selected only videos that presented the most challenging object tracking scenarios. The mining pipeline finds challenging single-shot video clips that are 5-30s long. Focusing on single-shot clips largely reduces annotation time and ambiguity originating from attempting to track objects across camera shots in edited videos. The mining pipeline consists of the following steps:

- *Scene and Motion Filters.* First, we leverage scene boundary detection and VMAF motion scores from FFmpeg ([FFmpeg developers](#)) to identify non-static single-shot camera clips from the video pool. To further improve the precision of single-shot clip selection, we also use Shot Boundary Detection from the PySceneDetect ([PySceneDetect Developers](#)) library;
- *Content Balancing.* We use a video-specific ontology to balance content distribution. We build the taxonomy by combining 1) frequent NPs annotated in the image data engine that tend to be associated with higher motion scores, and 2) a taxonomy that emphasizes human activities, animals and transportation. We then generate a set of text queries based on the video ontology and leverage PE [Bolya et al. \(2025\)](#) embeddings to retrieve video candidates for each text query. We propose text queries that elicit grouped objects and crowded scenes, for example “group of dogs” is a text query based on the concept “dog”;
- *Challenging Track Filter.* We use an MLLM (PLM ([Cho et al., 2025](#))) as a judge to select videos with challenging tracking scenarios. This is achieved by performing video-QA on a set of questions regarding the existence of various difficult scenarios, and selecting videos that receive more affirmative responses to these questions;
- *Targeted Semantic Search.* Lastly, we enhance the search for challenging scenarios by performing a video similarity search (using PE embeddings) using known challenging videos identified in human annotation as seeds.

Proposing NPs. We obtain candidate noun phrases for objects in the video.

- 2214 • *Frame-level captioner and parser.* We apply the Phase 3 captioner and parser on each video frame,
 2215 as opposed to video level, to maximize the diversity and volume of candidate noun phrases.
 2216
- 2217 • *NP Filtering.* To keep only relevant phrases, we implement a series of filters. First, we filter out
 2218 noun phrases that correspond to the overall scene, such as *room*, using a fine-tuned Llama 3.1 model.
 2219 Similarly, we filter out noun phrases that are too ambiguous to be masked, using the previously
 2220 trained EV AI Verifier, which has been trained to classify such cases. Next, we remove noun
 2221 phrases if they are present in a given list of restricted noun phrases. This list contains 1) phrases
 2222 that have been annotated as non-maskable in previous annotation rounds, 2) phrases for which we
 2223 already have a lot of annotations, and 3) phrases that correspond to “background” concepts, as our
 2224 focus is on challenging moving objects. Next, we optionally filter out phrases that do not belong to
 2225 certain pre-specified super-categories, such as “animal” or “vehicle” to further focus on moving
 2226 objects. We determine the super-category of a given noun phrase using a Llama 3.1 model.
 2227
- 2228 • *NP Cleaning.* The same cleaning is applied as in previous phases.

2229 **Proposing Masklets.** We use the latest iteration of SAM 3 to generate instance masklets by prompting
 2230 it with the proposed noun phrases.

- 2231 • *Masklet Generation.* Initially, we use SAM 3 at the image level to process frames independently,
 2232 and then propagate the masks using SAM 2. If masks detected in non-propagated frames are not
 2233 encompassed by the propagated masklets, they are used as starting points for new SAM 2 masklet
 2234 propagations. Once SAM 3 video performance surpassed the decoupled system, the pipeline was
 2235 updated to use SAM 3 alone.
 2236
- 2237 • *Masklet Deduplication.* After the masklets are obtained, we deduplicate them based on their IoU.
 2238
- 2239 • *Masklet Filtering.* We filter out the noun phrases that result in masklets containing the whole scene.
 2240
- 2241 • *Filtering Out Easy Cases.* We target challenging multi-object scenarios, namely videos that are
 2242 relatively crowded and contain multiple objects of the same category. The last step of the pseudo-
 2243 labeling pipeline filters out all noun phrases with fewer than N=3 objects, and videos that contain
 2244 fewer than M=2 such noun phrases.

2245 **Verification and Correction (Human).**

- 2246 • *Verification.* Human annotators check if the video is well pre-processed, e.g., no scene cuts, split
 2247 screen, or explicit content. Then they check if the noun phrase is groundable throughout the video,
 2248 e.g., there are no comparison or size attributes that might be unclear, and no action attributes which
 2249 might change across the timeline. Finally, they check that the masklet is challenging to track
 2250 yet possible to annotate i.e. focus on fast motion and highly occluded objects but which are still
 2251 identifiable by human annotators and not too blurry to annotate properly.
 2252
- 2253 • *Correction.* Another annotator reviews the proposed masklets, removing those that are incorrect
 2254 (improving precision), and using online SAM 2 in the loop to correct those that can be improved.
 2255 Next, they check for any missing masklets, and use SAM 2 to add them if needed (improving
 2256 recall). This annotation task results in two types of data: fully exhaustive tracking data where every
 2257 object that matches the noun phrase is annotated, or partially exhaustive tracking data, where some
 2258 masklets might be missing because they are impossible to annotate (e.g., inseparable background
 2259 objects that match the noun phrase).
 2260
- 2261 • *Exhaustivity Confirmation.* To ensure data quality, a final round of exhaustivity checking is
 2262 performed. If there are any remaining missing masklets, they are added as necessary.

2263 **Sampling Frame Annotations.** To increase the diversity and volume of the annotated video data, we
 2264 also sample video frames and annotate them using the image data engine (Phase 3), where they are
 2265 treated the same way as other images. The sampling follows two separate strategies. The first one is
 2266 just random sampling of a frame within a video. This guarantees we cover the distribution of frames.
 2267 The second strategy consists of first running the video data engine pipeline, and using the results to
 2268 determine frames that contain many objects.

2268 F SA-CO DATASET AND METRIC DETAILS
 2269

2270 F.1 SA-CO TRAINING DATA
 2271

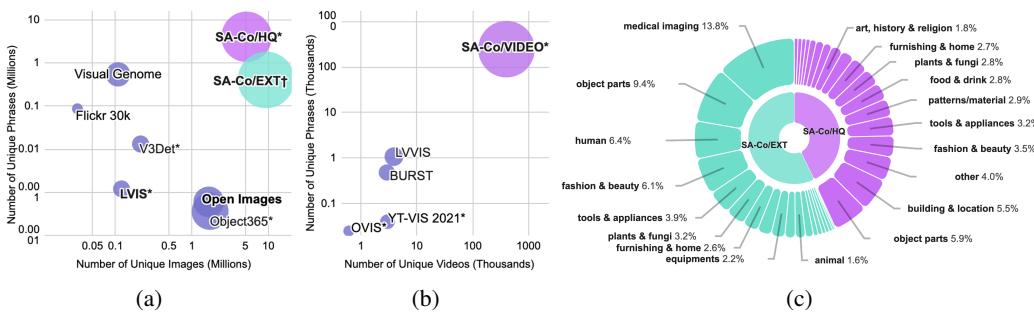
2272 SAM 3 training data includes images and videos from many diverse sources, including existing
 2273 datasets with box or mask annotations. The training data consists of three image datasets and one
 2274 video dataset. Fig. 10 visualizes statistics on these subsets in comparison with existing open-source
 2275 image and video detection and instance segmentation datasets as well as the distribution of top-level
 2276 SA-Co ontology categories on image datasets. More detailed statistics for each subset and comparison
 2277 with open-source datasets are shown in Tab. 26 and Tab. 27. The original dataset sources by subset
 2278 are listed in Tab. 28. .

2279 **SA-Co/HQ: High quality.** This image dataset is generated by the data engine in Phases 1-3 with
 2280 high quality annotations verified either by human annotators or by AI verifiers that have accuracy on
 2281 par with humans.

2282 **SA-Co/SYN: Synthetic.** We generate this synthetic dataset via the data engine in Phase 3, relying
 2283 only on AI annotators. We use MetaCLIP images as the media pool and extract NPs from two
 2284 sources: 1) alt-text captions associated with the images, 2) captions generated by Llama4. We prompt
 2285 SAM 3 using the extracted NPs to generate mask proposals. The image-NP-mask proposals are
 2286 then verified by MV and EV AI verifiers resulting in high-quality synthetic data. We also generate
 2287 hard negatives proposals (§E.4) and verify them using the EV AI verifier resulting in exhaustive
 2288 image-level negatives. This scalable system enabled large-scale synthetic data generation, resulting
 2289 in 39M images, 1.7B image-NPs and 1.4B masks.

2290 **SA-Co/EXT: External.** This dataset includes eighteen external datasets with existing instance mask
 2291 or bounding boxes annotations. For datasets with only bounding boxes, we generate instance masks
 2292 with SAM 2. We further enrich these external datasets by mapping the original label to SA-Co
 2293 ontology and propose additional negative labels using Wikidata hierarchy.

2294 **SA-Co/VIDEO: Video.** The video dataset is collected via the data engine in phase 4 with high
 2295 quality annotations. All the data in SA-Co/VIDEO is verified by human annotators.
 2296



2297 Figure 10: (a) SAM 3 image training data statistics and comparison with existing open-source image detection
 2298 and instance segmentation datasets. Bubble size denotes total number of NP-mask/bbox pairs. Bolded datasets
 2299 are annotated with masks, others are bboxes only. Datasets with * are exhaustively annotated, datasets with
 2300 † are partially exhaustively annotated. (b) SAM 3 video training data statistics and comparison with existing
 2301 open-source video instance segmentation datasets. Bubble size denotes total number of NP-masklet pairs.
 2302 Datasets with * are exhaustively annotated. (c) Instance masks distribution among SA-Co ontology top-level
 2303 categories in SA-Co/HQ and SA-Co/EXT. SA-Co/EXT incorporates several medical and microscopy image
 2304 datasets, we categorize them under medical imaging in addition to categories in Tab. 22.
 2305

2316 F.2 SA-CO EVALUATION BENCHMARK
 2317

2318 We release the Segment Anything with Concepts (SA-Co) Benchmark for promptable concept
 2319 segmentation in images and videos. Our benchmark contains images and videos paired with text
 2320 labels, each annotated exhaustively with masks on all object instances that match the label. The dataset
 2321 is federated, meaning that not all labels are annotated for all images, but only a handful of positive
 2322 and negative labels are verified as ground-truth per image. We add large volume of challenging

Dataset	# NPs	# Images	# Image-NP	% Negatives	# NP-bbox	# NP-mask	# masks per pair
Flickr 30k	86.4K	30.1K	193.0K	-	312.2K	-	-
LVIS*	1.2K	120.0K	1.6M	72.7%	1.5M	1.5M	3.51
V3Det*	13.2K	213.2K	737.7K	-	1.6M	-	-
Visual Genome	542.6K	108.1K	4.3M	-	6.3M	-	-
Open Images	600	1.7M	4.1M	-	13.3M	2.7M	2.79
Object365*	365	1.7M	10.1M	-	22.9M	-	-
SA-Co/HQ*	4.0M	5.2M	146.1M	88.5%	52.3M	52.3M	3.10
SA-Co/EXT[†]	497.4K	9.3M	136.6M	71.8%	70.5M	70.5M	1.83
SA-Co/SYN*	38.0M	39.4M	1.7B	74.0%	1.4B	1.4B	3.17

Dataset	# NPs	# Videos	# Video-NP	% Negatives	# NP-masklet	# masklets per pair
OVIS*	25	607	886	-	3.6K	4.04
YTVIS 2021*	40	3.0K	3.9K	-	6.3K	1.61
BURST	482	2.9K	6.9K	-	16.0K	2.33
LVVIS	1.2K	3.9K	13.8K	-	19.7K	1.40
SA-Co/VIDEO*	24.8K	52.5K	134.3K	26.7%	467.1K	4.75

Table 26: Detailed statistics for image training datasets and comparison with existing open-source image detection and instance segmentation datasets. Datasets with * are exhaustively annotated, datasets with † are partially exhaustively annotated. % Negatives denotes percentage of Image-NPs that are negatives.

Table 27: Detailed statistics for video training dataset and comparison with existing open-source video instance segmentation datasets. Datasets with * are exhaustively annotated. % Negatives denotes percentage of Video-NPs that are negatives.

hard negative label annotations to test model ability to handle large vocabularies. In particular, the SA-Co/Gold benchmark has $\sim 50\times$ more unique phrases compared to existing exhaustively annotated mask dataset LVIS-test. SA-Co benchmark covers a diverse array of sub-domains including common objects, fine-grained concepts, food, art, robotics, etc. See Tab. 29 for detailed benchmark statistics and Tab. 30 for the list of sub-domains.

F.3 METRICS

We introduce the *classification-gated F1* (CGF_1) to evaluate the PCS task on images. The traditional AP (Average Precision) metric designed for closed-vocabulary detection tasks (e.g., COCO), breaks down when applied to open-vocabulary detection with very large label spaces. While averaging AP over 80 classes is feasible, with tens of thousands most appear just once in the test set and the average is dominated by noise. Computing full precision-recall curves for all labels is also computationally infeasible and unnecessary for practical use cases. AP also does not account for the model calibration, which means that high-scoring models can be difficult to use in practice. F1 at a fixed confidence threshold presents a good alternative, however it is sensitive to high ratios of negative annotations: no extra credit is given for correctly predicting nothing for a negative, but the score is lowered by predicting false positives.

SA-Co/HQ	SA-Co/SYN	SA-Co/EXT	SA-Co/VIDEO	SA-Co/VIDEO-EXT	
MetaCLIP (Xu et al., 2024b;a)	MetaCLIP	Objects365 (Shao et al., 2019) OpenImages (Kuznetsova et al., 2020) ImageNet (Russakovsky et al., 2015) VisualGenome (Krishna et al., 2017) Sapiens Body-Parts (Khirodkar et al., 2024) EDEN (Le et al., 2021) Fashionpedia (Ji et al., 2020) Fathomnet (Katja et al., 2021) iNaturalist-2017 (Horn et al., 2017) BDD100k (Yu et al., 2020) Livecell (Edlund et al., 2021) PanNuke (Gamper et al., 2019; 2020) MedSAM2 (Ma et al., 2025) SNOW (Ding et al., 2023) Visdrone (Zhu et al., 2021) WCS Camera Traps (WCS) HierText (Long et al., 2023; 2022) FSC-147 (Ranjan et al., 2021)	SA-V (Ravi et al., 2024) SA-V internal (Ravi et al., 2024) YT-Temporal-1B (Zellers et al., 2022) SA-FARI	DAVIS2017 (Pont-Tuset et al., 2017b) MOSEv2 (Ding et al., 2025) YTVOS2019 (Xu et al., 2018)	
SA-1B (Kirillov et al., 2023)					
Armbench (Mitash et al., 2023)					
National Gallery of Art (nga)					
Ego4d (Grauman et al., 2022)					
MyFoodRepo-273 (Mohanty et al., 2021)					
GeoDE (Ramaswamy et al., 2023)					
DROID (Khazatsky et al., 2024)					
BDD100k (Yu et al., 2020)					
SA-V (Ravi et al., 2024)					
SA-V internal (Ravi et al., 2024)					
YT-Temporal-1B (Zellers et al., 2022)					

Table 28: Media pool used to construct each SA-Co train subset. See Figs. 11 to 15 for examples of each domain and annotations.

2376	Dataset	# NPs	# Images	# Image-NP	% Negatives	# NP-masks	% 0-shot NPs
2377	LVIS test	1.2K	19.8K	-	-	-	-
2378	COCO test2017	80	40.7K	-	-	-	-
2379	ODinW-35 test	290	15.6K	26.1K	-	131.1K	-
2380	SA-Co/Gold	59.2K	20.4K	215.4K	85.1%	157.5K	6.65%
2381	SA-Co/Silver	54.6K	66.1K	1.8M	94.0%	219.8K	8.00%
2382	SA-Co/Bronze	105.3K	32.5K	1.0M	84.9%	261.5K	57.25%
2383	SA-Co/Bio	166	5.4K	35.9K	71.8%	264.6K	-
2384	(a)						
2385	Dataset	# NPs	# Videos	# Video-NP	% Negatives	# NP-masklets	% 0-shot NPs
2386	LVVIS test	1.2K	908	-	-	5.7K	-
2387	BURST test	482	1.4K	3.4K	-	8.0K	-
2388	SA-Co/VEval	6.2K	1.7K	12.3K	71.5%	18.2K	7.62%
2389	(b)						

Table 29: (a) Summary statistics of SA-Co image benchmark by subsets and comparison with existing image instance segmentation benchmarks. (b) Summary statistics of SA-Co/VEval benchmark and comparison video instance segmentation benchmarks. % Negatives denotes percentage of Image-NPs or Video-NPs that are negative. Percentages of zero-shot NPs in each subset are reported. A zero-shot NP is defined as a phrase that has not been seen in the combined set of SA-Co/HQ, SA-Co/EXT and SA-Co/VIDEO.

2396	SA-Co/Gold	SA-Co/Silver	SA-Co/Bronze	SA-Co/Bio	SA-Co/VEval
2397	MetaCLIP	BDD100k	BDD100k	Livecell	SA-V
2398	SA-1B	DROID	EDEN	PanNuke	YT-Temporal-1B
2399	Attributes	Ego4D	Fashionpedia	MedSAM2	SmartGlasses
2400	Crowded Scenes	MyFoodRepo-273	Visdrone	SNOW	
2401	Wiki-Common1K	GeoDE	WCS Camera Traps		
2402	Wiki-Food/Drink	iNaturalist-2017			
2403	Wiki-Sports Equipment	National Gallery of Art			
2404		SA-V			
2405		YT-Temporal-1B			
2406		Fathomnet			

Table 30: Domains in each SA-Co test subset. See Figs. 11 to 15 for examples of each domain and annotations.

To remedy these issues we design new metrics for the PCS task. Given datapoints consisting of predicted and ground truth (media, phrase, masks) triplets we compute the following metrics to measure localization and classification separately:

- *Localization*. We measure this only on *positive* datapoints with at least one ground-truth mask and one predicted mask. For one sample, assume we have N predicted masks m_1, \dots, m_N and M ground-truth masks $\hat{m}_1, \dots, \hat{m}_M$. We compute the IoU matrix $\text{iou}_{i,j} = \text{iou}(m_i, \hat{m}_j)$, then deduce the optimal bipartite matching $\hat{\sigma} = \text{argmax}_{\sigma} \sum_i \text{iou}_{i,\sigma(i)}$. We fix an IoU threshold τ , then for every prediction i , if it is matched and $\text{iou}_{i,\sigma(i)} \geq \tau$, then it is counted as TP (true positive), otherwise FP (false positive). Unmatched ground truths are counted as FN (false negative). We measure $F_1^{\tau} = \frac{2\text{TP}}{2\text{TP} + \text{FP} + \text{FN}}$ for each datapoint, then average over all datapoints to obtain pmF_1^{τ} (“positive macro F1”). We compute pmF_1^{τ} for all $\tau \in [0.5, 0.95]$ with increments of 0.05, then average to obtain the final pmF_1 .
- *Classification*. This metric between $[-1, 1]$ computes the ability of the model to predict one or several masks, *if and only if* the datapoint is positive. This can be seen as a binary prediction task at the image level (“is the object present or not?”), and crucially, in this metric we do not care about the quality of the predicted masks. If the datapoint is *positive*, and if the model has predicted *any* mask (with confidence greater than 0.5), then it is an IL_TP (image level TP), otherwise an IL_FN. If the datapoint is *negative*, and if the model has predicted *any* mask, then it is an IL_FP, otherwise an IL_TN. We summarize this confusion matrix into a single metric, and measure potential imbalances with the Matthews Correlation Coefficient (MCC) as:

$$\text{IL_MCC} = \frac{\text{IL_TP} \cdot \text{IL_TN} - \text{IL_FP} \cdot \text{IL_FN}}{\sqrt{(\text{IL_TP} + \text{IL_FP}) \cdot (\text{IL_TP} + \text{IL_FN}) \cdot (\text{IL_TN} + \text{IL_FP}) \cdot (\text{IL_TN} + \text{IL_FN})}}$$

Evaluation Protocol	SA-Co/Gold			SA-Co/VEval				
	3-Reviewed Subset			YT-Temporal-1B test				
	CGF ₁	IL_MCC	pmF ₁	CGF ₁	pHOTA	pDetA	pAssA	
Best Pair	83.7	0.99	84.7	82.4	85.8	80.9	91.4	73.3
Random Pair	63.3	0.87	72.6	71.8	78.3	68.3	90.1	52.6

Table 31: Human instance segmentation performance comparison between the Best and Random Pair protocols on SA-Co benchmark. The comparison is shown on subsets of the benchmark with three human annotations.

As our main metric, we combine these two metrics to compute CGF₁ (“classification-gated F1”), defined as $\text{CGF}_1 = 100 \cdot \text{pmF}_1 \cdot \text{IL_MCC}$.

The PCS task is quite ambiguous in many cases, and to alleviate this issue our SA-Co/Gold subset contains three independent ground-truth annotations for each datapoint. To adapt our metric, we use an *oracle* setting, where we compare the model’s predictions to each ground-truth for each datapoint, and select the one that yields the best score.

F.4 HUMAN PERFORMANCE ON SA-CO

As described in §2, the PCS task is intrinsically ambiguous. Given an image-NP or video-NP pair, even trained annotators can have different interpretations that are all valid. When the phrase is vague, annotators can even disagree on the presence of the NP. Hence when evaluating on the SA-Co benchmark, disagreement with ground truth does not necessarily mean the prediction is wrong. To this end, it is important to study the human-level performance (i.e. the agreement among skilled annotators) on the PCS task to facilitate interpreting model performance.

Human Performance on SA-Co/Gold. On the image benchmark, we provide three sets of annotations by different annotators. These annotations are done from scratch, meaning that the annotators create masks (using SAM 1) without seeing any SAM 3 model interpretations. We define the “Best Pair” metric as follows to measure the upper bound of human performance. For each image-NP, the best pair (out of all three pairs of annotations) is selected by maximizing the local F1 score or minimizing the sum of false negatives (FN) and false positives (FP) when there is a tie in local F1 scores. We then report the CGF₁ metric based on these selected best pairs using one annotation as ground truth and the other as prediction. To make the model performance comparable, the “Best Pair” model performance is calculated by comparing model predictions to all three annotations and selecting the best pairs.

Alternative to the “Best Pair” protocol, human performance can also be measured on randomly selected pairs. Specifically, we adopt the following protocol to compute “Random Pair” human performance on SA-Co benchmark with three sets of annotations: 1) randomly choosing a pair of annotations for each image/video-NP, then aggregate over all image/video-NPs to get the metric values, 2) repeating the process a thousand times and reporting the 0.5 quantile for each metric. As shown in Tab. 31, there is a noticeable gap between Best Pair and Random Pair performance on both image and video benchmarks, suggesting that the PCS task is inherently ambiguous.

The image benchmark has a large portion of hard negatives. These phrases go through human verification, but as it is costly to collect three sets of human annotations on the entire dataset due to the large volume, the negative noun phrases only have one ground-truth label. The human performance on these phrases is estimated by collecting additional human annotations on a subsample of phrases and comparing them with the initial annotation (i.e., the ground truth). Specifically, we collect additional human annotations on about one thousand image-NPs for each domain in SA-Co/Gold. Since the ground truths are all negatives, these phrases only contribute to the IL_MCC metric. We compute counts of IL_TN and IL_FP on these samples, and then extrapolate these results to estimate the corresponding counts for the entire set of hard negatives. These estimated counts are then combined with image-level counts from the rest of the benchmark where NPs have three annotations to get the final IL_MCC.

During the additional human review, some hard negative NPs are marked as ambiguous by annotators. To reflect this uncertainty in the final IL_MCC metric, we estimated a lower bound by counting all ambiguous responses as IL_FP (i.e. incorrect), and an upper bound by counting them as IL_TN (i.e.

2484 correct). Hence we report human performance on image PCS task as an interval defined by the lower
 2485 and upper bounds, capturing the uncertainty introduced by ambiguous phases.
 2486

Human Performance on SA-Co/VEval. Annotating videos is much more expensive than static
 2487 images, so we collect only one set of annotations per NP on the video benchmark. To guarantee
 2488 annotation quality, these ground-truth annotations undergo multiple rounds of human refinement.
 2489 To measure human performance in a way that is directly comparable to model evaluation in video
 2490 PCS, we collect one additional from-scratch human annotation for every NP in the test set across all
 2491 sub-domains. Human performance on video PCS task is then reported by comparing this additional
 2492 annotation to the ground truth, using the same metrics as for model evaluation (CGF₁ and pHOTA).
 2493

2494 Additionally, to study the gap between the Random Pair and the Best Pair protocols, we collect two
 2495 further human annotations (for a total of three) on the YT-Temporal-1B and SmartGlasses test splits
 2496 of the SA-Co/VEval dataset. This allows us to verify that the gap observed in the image domain also
 2497 exists in the video domain (see Tab. 31).

2498 F.5 ADDITIONAL DATASET EXAMPLES

2500 Figs. 11 to 15 show examples of each visual domain in our image and video datasets. Fig. 16 show an
 2501 example image from our synthetic data set SA-Co/SYN, with its positive noun phrases in the figure
 2502 and negative noun phrases in the caption.

2503
 2504
 2505
 2506
 2507
 2508
 2509
 2510
 2511
 2512
 2513
 2514
 2515
 2516
 2517
 2518
 2519
 2520
 2521
 2522
 2523
 2524
 2525
 2526
 2527
 2528
 2529
 2530
 2531
 2532
 2533
 2534
 2535
 2536
 2537

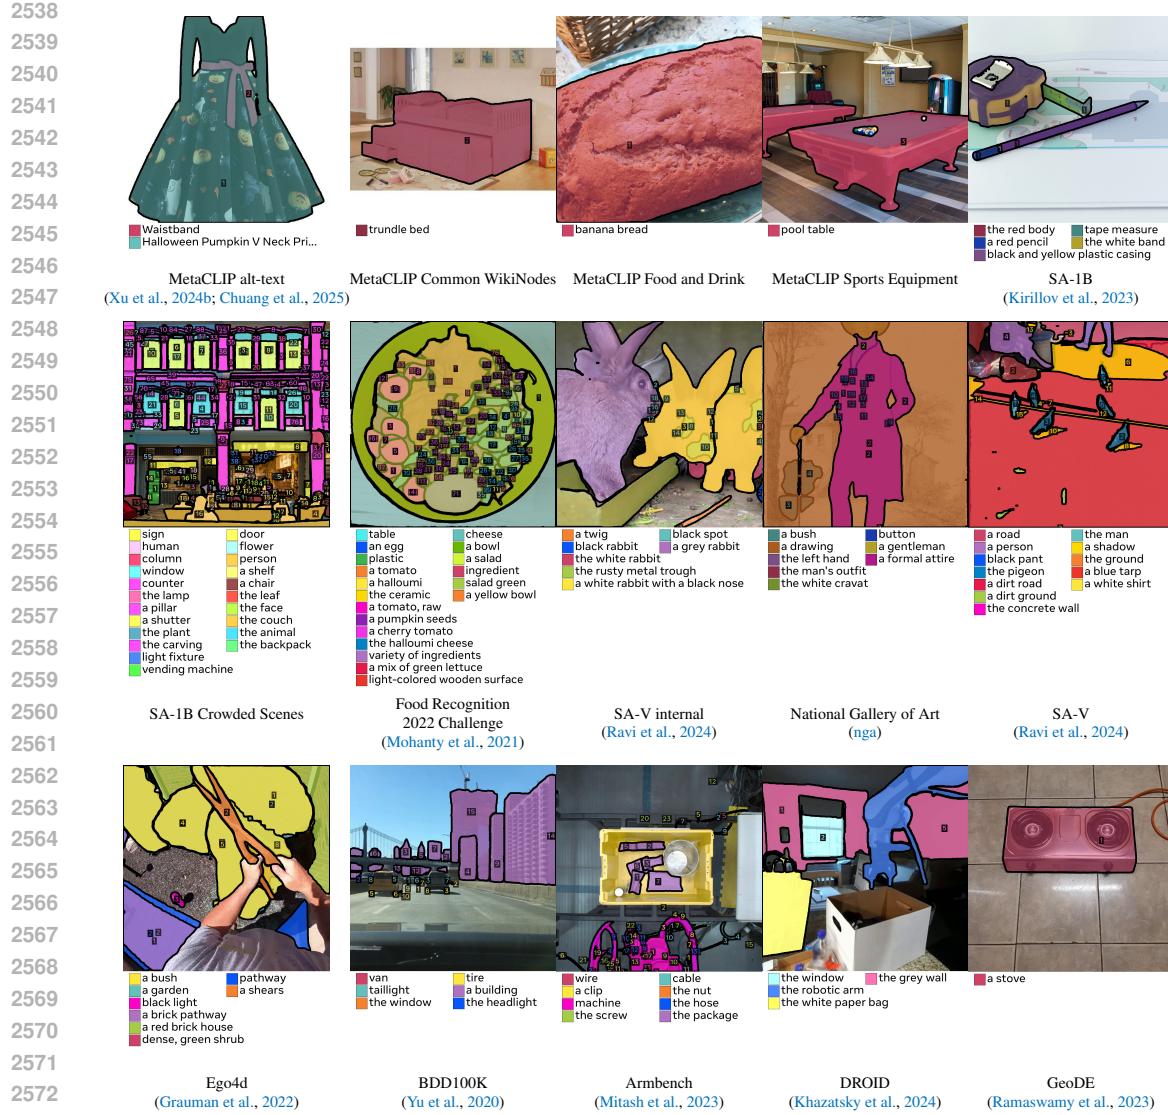
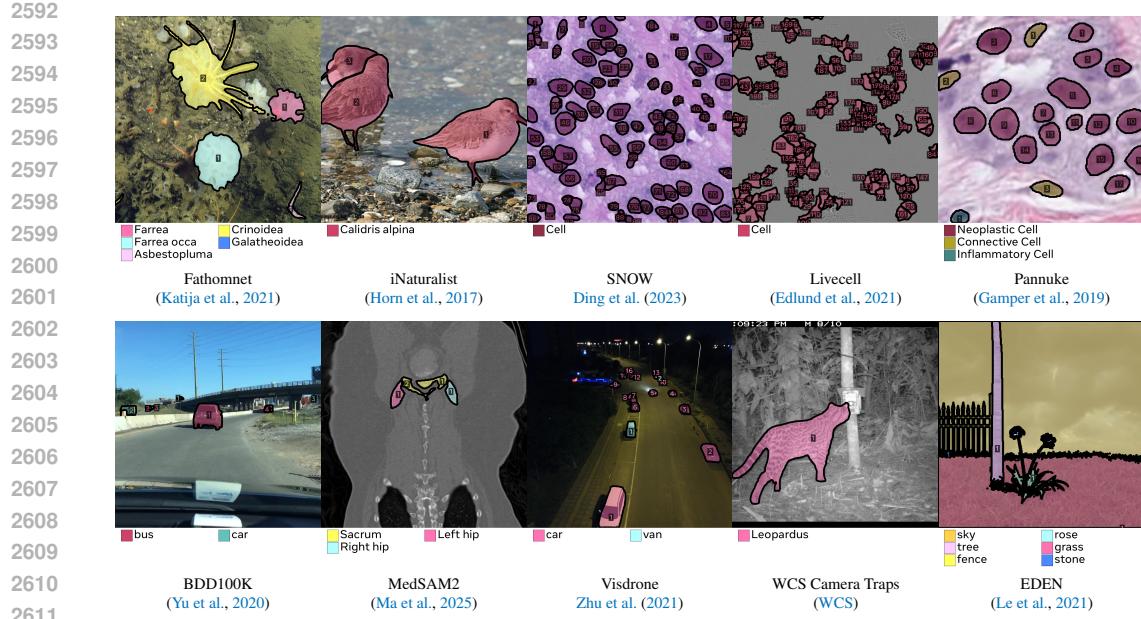
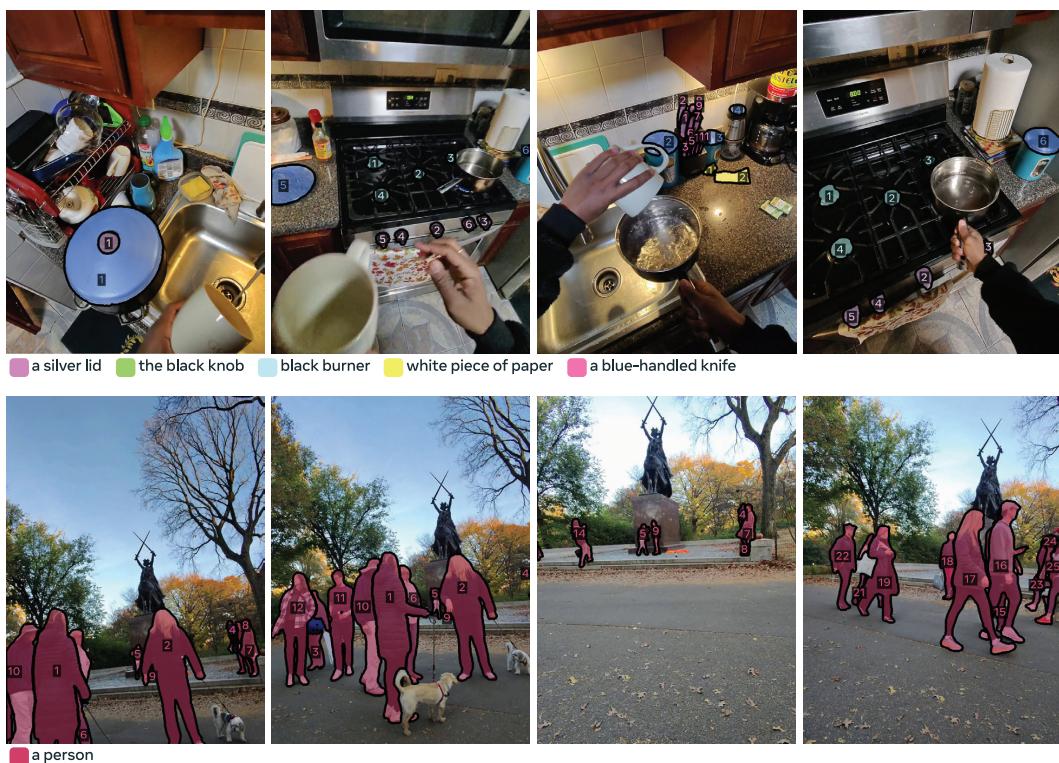


Figure 11: Per-domain examples in SA-Co/HQ. Shown with annotated phrases and masks overlaid.

2612 Figure 12: Per-domain examples in SA-Co/EXT. Shown with annotated phrases and masks overlaid.
2613
2614
2615
2616
2617
2618
2619
2620

2612 Figure 13: Example annotations from the SmartGlasses media in the SA-Co/VEval dataset.

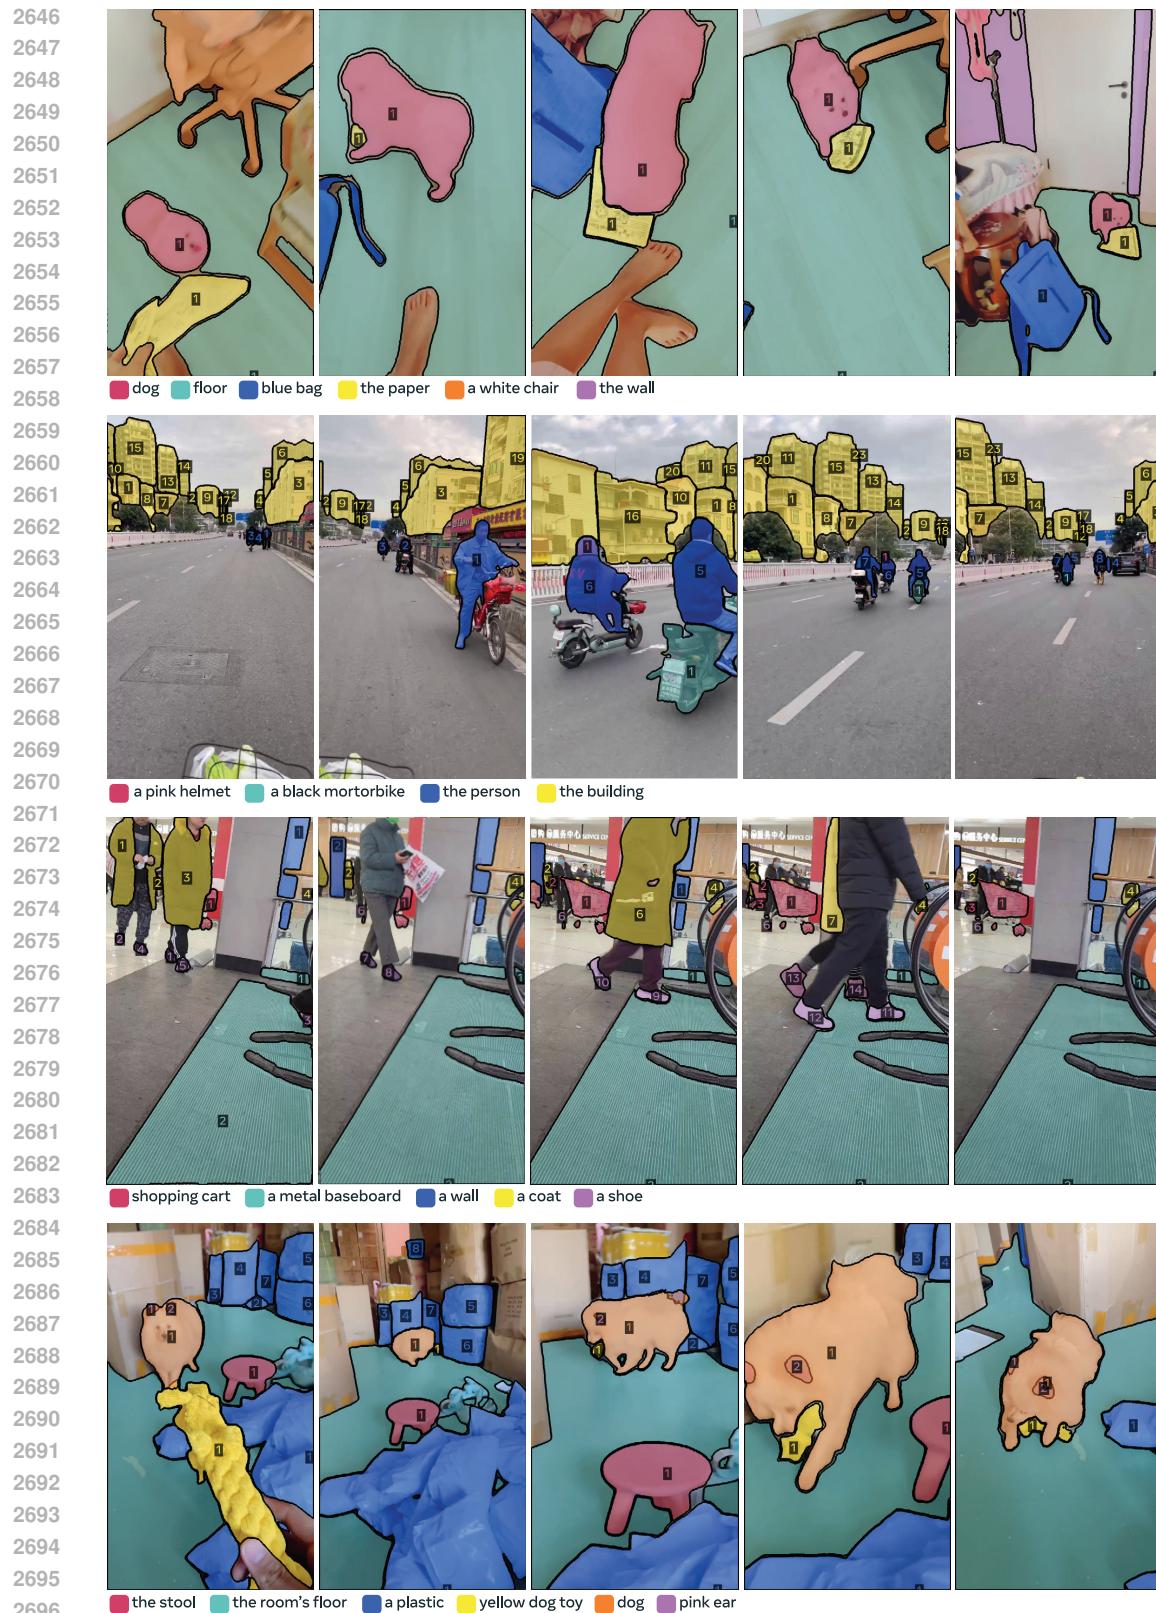


Figure 14: Example annotations from the SA-V media in the SA-Co/VIDEO and SA-Co/VEval datasets.



Figure 15: Example annotations from the SA-FARI dataset.



Figure 16: Example annotations on one image from the SA-Co/SYN dataset. There are 19 positive noun phrases for this image, which we visualize in 4 groups for better visualization quality. AI verifier assigned an exhaustivity label for each noun phrase: noun phrases in **green box** means the masks for that noun phrase is exhaustive, noun phrases in **pink box** means not exhaustive. There are 20 negative noun phrases that are verified and confirmed by AI verifier for this image: *a black and white post, a blue and yellow post, a large red and white truck, an orange and black post, large orange and white boat, the green and red post, the large orange and white motorcycle, the orange and white bollard, the purple and yellow post, yellow sash, the Malaysian Air Boeing 737-800, display vehicle, Zamboni ice resurfacer, boot, the small folded paper package, long yellow vehicle, the small gift-wrapped package, small trailer, bunker gear, two-seat model*.

2754 G ADDITIONAL EXPERIMENTS AND DETAILS

2756 G.1 PCS WITH NPs ON IMAGES

2758 This section describes the experiments in Tab. 1 in detail. We compare to OWLv2 (Minderer et al.,
 2759 GroundingDino (Liu et al., 2024a) and LLMDet (Fu et al., 2025). Since they only produce
 2760 bounding boxes, we convert them to masks using SAM 1 to evaluate segmentation. We also compare
 2761 to APE (Shen et al., 2024) and DINO-X (Ren et al., 2025), two SoTA segmentation models, and
 2762 finally Gemini 2.5 Flash (Comanici et al., 2025), a generalist LLM.

2763 We report performance on LVIS (Gupta et al., 2019), COCO (Lin et al., 2014), COCO-O (Mao et al.,
 2764 2023), Cityscapes (Cordts et al., 2016), ADE (Zhou et al., 2019), and Pascal context (Mottaghi et al.,
 2765 2014), reporting their official metrics. For LVIS, we report AP-fixed (Dave et al., 2022). On our new
 2766 SA-Co benchmark, we report the average across every split. We report CGF₁, except for SA-Co/Bio
 2767 where we report pmF₁ (this split does not have negatives, so only localization is meaningful). On
 2768 SA-Co/Gold we have three ground-truth annotations per datapoint, so we report the oracle metric and
 2769 estimated human performance (human performance measurement detailed in §F.4).

2770 We employ the following Hugging Face model checkpoints: “google/owlv2-large-patch14” for
 2771 OWLv2, “google/owlv2-large-patch14-ensemble” for OWLv2*, “IDEA-Research/grounding-dino-
 2772 tiny” for gDino-T, and “iSEE-Laboratory/llmdet_large” for LLMDet-L. OWLv2* utilizes an ensemble
 2773 of checkpoint weights after self-training and after fine-tuning the model on LVIS base, demonstrating
 2774 improved open-world generalization compared to fine-tuning alone (Minderer et al., 2024). In
 2775 Tab. 32, we provide per-domain performance of instance segmentation on SA-Co/Gold for all
 2776 baselines, SAM 3, and human in Tab. 1. We also provide the per-domain performance for AI
 2777 verifier ablation study in Tab. 9d. In Tab. 33, we provide comparisons with additional baselines
 2778 using “IDEA-Research/grounding-dino-base” for gDino-B and “iSEE-Laboratory/llmdet_base” for
 2779 LLMDet-B.

2780 For OWLv2, GroundingDino, and LLMDet, we swept over the detection threshold at 0.1 intervals
 2781 and determined the best threshold using the LVIS CGF₁ metric for the box detection task. Then, we
 2782 applied this threshold to compute CGF₁ on the remaining datasets for the box detection and instance
 2783 segmentation tasks. The detection threshold is set to 0.4 for LLMDet-L, LLMDet-B, and gDino-T;
 2784 0.3 for OWLv2*; and 0.2 for OWLv2 and gDino-B. For DINO-X, we find the detection threshold
 2785 0.5 gives the best CGF₁ metric. Additionally, we found that prompting multiple noun phrases at
 2786 once for a given image greatly improved performance for GroundingDino and LLMDet, compared to
 2787 prompting one noun phrase at a time. For example, we prompted GroundingDino and LLMDet with
 2788 30 prompts for SA-Co/Gold and 20 prompts for SA-Co/Silver, SA-Co/Bronze, and SA-Co/Bio.

2789 For Gemini 2.5 Flash, we run inference via the Gemini API. For each (image, text query) pair,
 2790 we prompt Gemini 2.5 using the same prompt template that is used in Gemini Flash 2.5 image
 2791 segmentation demo (Paul Voigtlaender, Valentin Gabeur and Rohan Doshi, 2025) with the same
 2792 generation settings. In addition, we prompt the model multiple times if there are any errors in
 2793 generation, or parsing the result into a set of masks and bounding boxes.

	Average		Metaclip		SA-IB		Crowded		Food&Drink		Sports Equip.		Attributes		Wiki-Common									
	CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁	CGF ₁	IL_MCC	pmF ₁									
gDino-T	9.1	.15	57.7	13.6	.21	64.0	12.4	.20	62.1	3.8	.09	42.3	4.64	.09	51.0	5.4	.09	59.9	20.8	.28	73.9	.3	.06	50.5
OWLv2*	34.3	.58	58.9	32.1	.53	60.9	25.3	.51	50.0	23.8	.53	45.1	40.7	.65	62.8	42.7	.64	66.9	45.0	.63	71.1	30.4	.55	55.7
OWLv2	22.9	.46	49.1	21.1	.40	52.4	17.3	.46	37.4	15.1	.38	39.8	27.2	.51	53.5	29.8	.52	56.9	31.3	.53	58.6	18.7	.42	44.8
LLMDet-L	12.9	.21	59.5	14.3	.23	61.9	13.7	.24	58.2	7.6	.17	44.9	11.9	.2	60.0	10.3	.17	61.7	29.6	.39	76.4	2.8	.05	53.2
DINO-X	27.7	.38	71.8	24.9	.35	70	27.2	.43	62.3	15.8	.33	47.7	38.6	.51	75.2	36.6	.44	82.8	37.8	.43	87.5	13.3	.17	77.4
Gemini 2.5	18.7	.29	65.8	16.5	.29	57.3	22.5	.41	54.5	12.3	.28	43.6	24.5	.33	74.0	21.1	.28	74.2	24.5	.3	82.7	9.7	.13	74.1
SAM 3	.65	.82	79.7	63.0	.81	78.1	63.6	.85	75.0	63.0	.90	69.9	67.9	.81	83.7	73.5	.89	82.3	66.5	.76	87.3	57.6	.70	81.8
SAM 3 +EV	68.1	.86	78.8	64.1	.85	75.6	65.1	.89	73.1	61.1	.88	69.3	71.3	.86	82.5	73.3	.90	81.5	80.5	.91	88.3	61.6	.76	81.5
SAM 3 +EV&MV	69.2	.85	81.3	66.2	.85	78	65	.88	76.7	63.2	.89	71.2	73.3	.87	84.3	74.2	.89	83.4	81.1	.89	90.8	61.4	.73	84.4
Human-Lower	74.2	.88	84.6	68.2	.86	79.8	73.1	.93	78.9	64.7	.85	76.4	79.2	.91	87.3	83.1	.94	88.8	84.8	.93	91.2	.66	.73	90.2
Human-Upper	81.4	.96	84.6	76.3	.96	79.8	76.4	.97	78.9	73.3	.96	75.4	85	.97	87.3	86.8	.98	88.8	87.9	.96	91.2	84.2	.93	90.2

2801 Table 32: Per-domain results of instance segmentation on SA-Co/Gold. *partially trained on LVIS.
 2802
 2803

2804 G.2 VISUAL EXEMPLARS AND INTERACTIVITY

2805 In Tab. 34, visual exemplar experiments, we report performance in 3 settings: (1) text prompt only,
 2806 (2) visual prompt only, and (3) both text and visual prompt. We note that (2) is quite ambiguous. For
 2807 example, given a visual example of a dog, one could want to detect all dogs, or only dogs of the same

Model	Box Detection							
	LVIS		COCO		SA-Co			
	CGF ₁	AP	AP	AP ₀	Gold	Silver	Bronze	Bio
Human _{min}	—	—	—	—	76.2	—	—	—
Human _{max}	—	—	—	—	83.6	—	—	—
OWLv2 (Minderer et al., 2024)	35.2	35.2	38.2	42.4	23.3	11.7	8.5	1.6
OWLv2* (Minderer et al., 2024)	47.4	45.5	46.1	23.9	35.4	19.2	19.4	0.3
gDino-B (Liu et al., 2024a)	31.6	25.7	52.5	45.5	7.6	5.6	13.9	0.9
gDino-T (Liu et al., 2024a)	33.8	20.5	45.7	35.3	9.4	6.8	12.3	0.7
LLMDet-L (Fu et al., 2025)	53.7	42.0	55.6	49.8	13.5	11.5	20.6	0.6
LLMDet-B (Fu et al., 2025)	51.9	37.8	54.2	39.4	9.9	11.0	21.2	0.9
Gemini 2.5 (Comanici et al., 2025)	23.7	—	—	—	18.7	12.2	10.1	12.0
SAM 3	57.5	51.7	53.5	55.5	67.7	58.0	53.1	60.0

Table 33: Additional evaluation on image concept segmentation tasks. AP₀ corresponds to performance on the COCO-O dataset. Models with a star were partially trained on LVIS.

Model	SA-Co/Gold		
	pmF ₁ T	pmF ₁ I	pmF ₁ T+I
	T-Rex2 (Jiang et al., 2024)	—	57.6
SAM 3	63.2	69.6	73.5

Table 34: Visual prompting on SA-Co/Gold. We report pmF₁ metric in different prompt types: T (text-only), I (image-only), and T+I (combined text and image).

color or breed. As a result, SAM 3 performs worse on SA-Co/Gold in setting (2) compared to (1). Therefore, setting (3) is better suited, as the text lifts most of the ambiguity, and the additional input box gives a hint for unfamiliar concepts.

G.3 FEW-SHOT FINE-TUNING

We evaluate SAM 3’s object detection capabilities on real-world data through comprehensive zero-shot and few-shot experiments using two established benchmarks: OdinW13 Li et al. (2022a) and Roboflow-100VL Robicheaux et al. (2025). These benchmarks encompass 13 and 100 diverse object detection datasets, respectively, capturing a wide range of real-world scenarios with standardized train and test splits that enable fair comparison with existing methods.

Few-shot training and evaluation. For OdinW13 few-shot experiments, we train on all three official few-shot training splits and report mean performance with standard deviation on the test split. For Roboflow-100VL, we utilize the official FSOD training splits provided by the benchmark and report numbers on the test split. We treat few-shot fine-tuning runs similarly to traditional training runs, but with some differences. We train for 40 epochs a reduced learning rate that is one-tenth of the standard value on a batch size of 2. Since these benchmarks focus exclusively on object detection without mask annotations, we disable all mask-specific components and losses during training.

OdinW13 results. Fig. 17a presents our few-shot performance on OdinW13, comparing SAM 3 against previous state-of-the-art methods Ren et al. (2024); Wu et al. (2024b); Xu et al. (2023); Zhang et al. (2022b). We report mean BoxAP averaged across all 13 datasets, with SAM 3 consistently achieving superior performance and establishing new state-of-the-art results. Complete dataset-specific results for each OdinW13 dataset are provided in Fig. 17b.

Roboflow-100VL results. Tab. 35 summarizes our comprehensive evaluation across zero-shot, few-shot, and full fine-tuning settings on Roboflow-100VL, with results averaged across all 100 datasets. While SAM 3 underperforms the current state-of-the-art Liu et al. (2023) in zero-shot evaluation, it surpasses leading methods Liu et al. (2023); Chen et al. (2024a) in both few-shot and full fine-tuning scenarios. This demonstrates SAM 3’s strong visual generalization capabilities when provided with task-specific training data. We attribute the zero-shot performance gap to the use of specialized, dataset-specific prompts that may lack broad generalizability in Roboflow-100VL. However, even minimal fine-tuning closes this gap and enables substantial performance improvements. Roboflow-100VL also categorizes its 100 datasets into seven dataset types; we report averages per each such dataset type in Tab. 36.

2862	Model	zero-shot	1-Shot	3-Shot	5-Shot	10-Shot	All
2863				AP			
2864	GLIPv2-H	55.5	61.7 ± 0.5	64.1 ± 0.8	64.4 ± 0.6	65.9 ± 0.3	70.4
2865	GLEE-Pro	53.4	59.4 ± 1.5	61.7 ± 0.5	64.3 ± 1.3	65.6 ± 0.4	69.0
2866	MQ-GLIP-L	54.1	62.4	64.2	65.4	66.6	71.3
2867	Grounding DINO 1.5 Pro	58.7	62.4 ± 1.1	66.3 ± 1.0	66.9 ± 0.2	67.9 ± 0.3	72.4
2868	SAM 3	59.9	63.0 ± 2.4	67.3 ± 0.7	69.1 ± 1.1	71.6 ± 0.2	75.6

2869	(a) Comparison of different models under few-shot settings on ODinW13.								
2870	Aerial	maritime	drone(l)	Aquarium	Rabbits	Egohands(g)	NAMushrooms	Packages	PascalVOC
2871				AP					
2872	0-Shot	22.3		41.7	82.9	69.8	89.4	72.9	64.0
2873	1-Shot	29.6 ± 0.4		46.0 ± 2.8	81.6 ± 1.9	70.6 ± 1.3	92.3 ± 8.4	84.2 ± 3.8	67.1 ± 0.5
2874	3-Shot	34.5 ± 0.3		51.5 ± 3.7	82.8 ± 3.0	70.5 ± 2.9	98.7 ± 1.3	88.2 ± 4.2	68.0 ± 0.5
2875	5-Shot	36.9 ± 2.1		53.2 ± 0.7	83.2 ± 2.0	71.1 ± 1.2	99.1 ± 1.0	89.2 ± 4.4	68.2 ± 0.2
2876	10-Shot	40.2 ± 1.9		55.0 ± 0.8	84.1 ± 1.5	72.1 ± 0.7	100.0 ± 0.0	93.3 ± 0.3	68.6 ± 0.3
2877	All	39.7		62.1	83.3	79.8	100.0	98.0	76.3

2878	Raccoon	Shellfish	Vehicles	Pistols	Pothole	ThermalDP	13-Average
2879			AP				
2880	0-Shot	70.5	58.8	65.6	59.6	27.5	53.4
2881	1-Shot	62.6 ± 22.0	49.4 ± 8.5	59.9 ± 3.9	68.8 ± 2.1	36.4 ± 2.5	70.6 ± 5.8
2882	3-Shot	78.9 ± 2.8	51.3 ± 2.4	65.8 ± 2.2	69.9 ± 0.7	42.0 ± 3.8	73.5 ± 6.1
2883	5-Shot	81.1 ± 1.1	52.6 ± 3.1	64.7 ± 1.0	71.3 ± 2.7	44.4 ± 2.7	82.7 ± 7.2
2884	10-Shot	86.0 ± 1.3	56.7 ± 0.6	68.0 ± 2.0	72.3 ± 1.3	47.4 ± 1.8	87.4 ± 1.4
2885	All	87.4	58.4	71.6	78.5	58.1	89.2
2886							75.6

(b) ODinW13 per-dataset results for SAM 3.

Figure 17: Zero-shot and few-shot results on ODinW13.

G.4 OBJECT COUNTING

We evaluate on object counting benchmarks CountBench (Paiss et al., 2023) and PixMo-Count (Deitke et al., 2025) to compare with MLLMs (Wang et al., 2024; Deitke et al., 2025; Comanici et al., 2025) and detection expert models (Ren et al., 2025). See Tab. 37 for results. The metrics include Accuracy(%) and Mean Absolute Error (MAE). CountBench (Paiss et al., 2023) contains 540 images and their captions, with 2-10 objects in each image. By removing images with unavailable links, we test on 487 images. PixMo-Count (Deitke et al., 2025) contains 540 images and their text descriptions in the form of simple noun phrases, with 2 to 10 objects in each image. By removing images with unavailable links, we test on 529 images.

To evaluate MLLMs on CountBench, we apply the same question set as Molmo (Deitke et al., 2025), which is inherited from PaliGemma (Beyer et al., 2024). When evaluating SAM 3 on CountBench, we modify the question sentence to the simple noun phrase. To evaluate MLLMs on PixMo-Count, we construct the question as “How many {} are there in this image” or “Count the {}”, where {} is the simple noun phrase provided by PixMo-Count annotations.

We find that presence token does not help SAM 3 on counting tasks, so we do not use it. For a group of objects, we find that SAM 3 outputs predictions for both each individual and the group as a whole, which contradicts the counting task.

As a post-processing step, we perform Non-Maximal Suppression (NMS) to remove duplicate detections. Instead of the usual Intersection-over-Union (IoU) criterion, we use Interaction over

2910	Model	0-Shot	10-Shot	All	Aver-								
		AP			Aerial	Docu-	Flora-	Indus-	Medi-	Other	Sports	Age	
2911	Grounding Dino	15.7	33.7	—	0-Shot	20.7	11.9	23.4	8.2	2.0	15.7	16.4	14.3
2912	LW-DETRm	—	—	59.8	10-Shot	33.9	34.0	38.8	39.7	24.9	34.7	40.3	35.7
2913	SAM 3	14.3	35.7	61.0	All	56.0	63.0	59.8	67.4	51.5	63.4	61.6	61.0

Table 35: Comparison on Roboflow100-VL.

Table 36: SAM 3 Roboflow100-VL results by dataset type.

Minimum (IoM), where the area of overlap is divided by the area of the smaller mask, rather than by the area of the union. By doing this, we can detect whole-vs-part situations: if a mask is fully covered by another, the IoM will be high, even if the covering mask is much bigger (which would lead to low IoU). We set the IoM threshold to 0.5 in our NMS process. Finally, we select the predictions with confidence higher than 0.5 as the final predictions and count the number of predictions as the counting result.

Model	Presence	IoM	CountBench		PixMo-Count	
			MAE ↓	Acc (%) ↑	MAE ↓	Acc (%) ↑
SAM 3	✗	✗	0.36	81.9	0.36	75.4
SAM 3	✓	✗	0.46	84.1	0.43	75.6
SAM 3	✗	✓	0.11	95.6	0.22	87.3
SAM 3	✓	✓	0.30	92.6	0.30	86.7

Table 37: Ablation on counting results.

G.5 VIDEO PCS DETAILS

In this section, we provide additional details for the video PCS evaluation (in §6 and Tab. 5).

Benchmarks. We evaluate the video PCS capabilities of the SAM 3 model based on an input text prompt (similar to the open-vocabulary video instance segmentation task (Wang et al., 2023)) on both our collected video benchmark SA-Co/VEval and public benchmarks. For SA-Co/VEval, we evaluate separately on each subset (SA-V, YT-Temporal-1B, and SmartGlasses) based on their data sources, and report classification-gated F1 (CGF₁), phrase-based HOTA (pHOTA), and Track Every Thing Accuracy (TETA). The SA-Co/VEval benchmarks contain a large number of noun phrases (5.1K in SA-V and YT-Temporal-1B subsets and 4.9K in SmartGlasses), and provides each video with a list of noun phrases as text prompts. During evaluation, for each evaluation video, we prompt SAM 3 with the list of noun phrases provided for that video, as shown in Tab. 38 (a, b, c).

For public benchmarks, we evaluate on LVVIS (Wang et al., 2023), BURST (Athar et al., 2023), YTVIS (Ke et al., 2022), OVIS (Qi et al., 2022), BDD100K (Yu et al., 2020), GMOT40 (Bai et al., 2021), and DeepSeaMOT (Barnard et al., 2025), and report the official metrics on each dataset (for DeepSeaMOT, we report the average performance over its 4 subsets). These public benchmarks are often based on a set of categories, with a relatively large vocabulary size in LVVIS and BURST (1196 categories in LVVIS and 482 categories in BURST) and much smaller numbers of categories in other datasets. We use the category name as the text prompt, and prompt SAM 3 with all category names in the dataset on every evaluation video, as shown in Tab. 38 (d).

Video PCS Metrics. Similar to its definition in the image domain in §F.3, we define the classification-gated F1 (CGF₁) metric on videos as the multiplication between the video-level Matthews correlation coefficient (VL_MCC) on whether the noun phrase exists in the video and the localization positive macro F1 (pmF₁) on positive noun phrases. To decide whether a predicted masklet matches a ground-truth masklet, we measure their volume intersection-over-union (IoU), defined by their total intersection volume divided by their total union volume over the video. When computing pmF₁, we averaged the results over multiple volume IoU thresholds from 0.5 to 0.95 with increments of 0.05, similar to how it is computed on images.

We also evaluate the phrase-based HOTA (pHOTA) metric, where we compute the Higher Order Tracking Accuracy (HOTA) metric (Luiten et al., 2021) over all video-NP pairs along with their breakdown into phrase-based detection accuracy (pDetA) and phrase-based association accuracy (pAssA). As the HOTA metric was originally designed for category-based evaluation, to get its phrase-based variant pHOTA for open-vocabulary prompts, we remap each video-NP pair in the evaluation benchmark as a new unique video ID and then set all ground-truth annotations and predictions to have the same category ID (i.e., the total number of video IDs after remapping equals the total number of video-NP pairs in the evaluation benchmark). That is, each video-NP pair in the benchmark is treated as an isolated sample for prediction and evaluation, and the results are aggregated over all video-NP pairs in a class-agnostic manner. More specifically, we save the remapped ground-truth annotations and the predictions as JSON files under the YTVIS format, and use the TrackEval package (Jonathon Luiten, 2020) to obtain the mask HOTA statistics on this remapped dataset (using the YTVIS dataset wrapper in TrackEval along with its default parameters), and report their results as

2970 pHOta, pDetA, and pAssA. Similarly, we also evaluate the Track Every Thing Accuracy (**TETA**)
 2971 metric (Li et al., 2022c) over the masklet predictions on these datasets.

2972 **Baselines.** We compare the SAM 3 model with several baselines, including GLEE (Wu et al., 2024a)
 2973 (a previous work on open-vocabulary image and video segmentation), “LLMDet as detector + SAM 3
 2974 Tracker”, by replacing the Detector component in SAM 3 with a state-of-the-art open-vocabulary
 2975 model LLMDet (Fu et al., 2025), and “SAM 3 Detector + T-by-D as tracker”, by replacing the Tracker
 2976 component in SAM 3 with a simple association module commonly used in tracking-by-detection
 2977 approaches (Wojke et al., 2017; Zhang et al., 2022c).

2978 For GLEE (Wu et al., 2024a), we follow its official implementation. Since GLEE supports taking as
 2979 inputs multiple text prompt simultaneously, we evaluate it in two ways: a) prompting it with all the
 2980 noun phrases from an evaluation video at once, denoted as “GLEE (prompted w/ all NPs at once)” in
 2981 Tab. 38, and b) looping over each noun phrase in the evaluation video and prompting GLEE with
 2982 one noun phrase at a time, denoted as “GLEE (prompted w/ one NP at a time)”. We find that for
 2983 open-vocabulary segmentation on videos, it is usually better to prompt GLEE with one noun phrase
 2984 at a time instead of prompting it with all noun phrases at once.

2985 For “LLMDet as Detector + SAM 3 Tracker”, we replace the detection outputs from the SAM 3
 2986 detector with LLMDet (Fu et al., 2025) bounding box outputs, and obtain the mask output by
 2987 prompting it with the SAM 3 component. Then we apply the SAM 3 Tracker similar to how it
 2988 is applied over the SAM 3 Detector output. We also note that GLEE and LLMDet have not been
 2989 trained on the noun phrases in the SA-Co dataset, so their results should be seen as zero-shot on the
 2990 SA-Co/VEval benchmark.

2991 For “SAM 3 Detector + T-by-D as tracker”, we replace the SAM 3 Tracker with a detection-to-masklet
 2992 association module as commonly used in the tracking-by-detection paradigm, e.g. Wojke et al. (2017);
 2993 Zhang et al. (2022c). The detection-to-masklet association module tries to match the masklets already
 2994 tracked in previous frames with detected objects in the current frame, based on a dot product between
 2995 the visual features of each detected object and the visual features of the past 16 frames of a masklet.
 2996 If a high-confidence detection isn’t matched to any existing masklet, we add it as a new object and
 2997 start a new masklet for it. The association module is trained on the SA-Co dataset.

2998 **Results.** As shown in Tab. 38, SAM 3 largely outperforms these baselines across the benchmarks.
 2999 On SA-Co/VEval with a very large number of noun phrases, SAM 3 excels in both frame-level
 3000 detection (pDetA) and cross-frame association (pAssA). Comparisons with “LLMDet as Detector
 3001 + SAM 3 Tracker” and “SAM 3 Detector + T-by-D as tracker” demonstrate that both the Detector
 3002 module and the Tracker module in SAM 3 play a critical role in the final video performance. In public
 3003 benchmarks, SAM 3 also achieves a strong performance, including new state-of-the-art results on
 3004 LViS and OViS. We also note that GLEE and LLMDet have not been trained on the SA-Co dataset,
 3005 so their results should be seen as zero-shot on SA-Co/VEval. In addition, the SmartGlasses subset in
 3006 SA-Co/VEval contains many egocentric videos, which might be out of the training distribution for
 3007 GLEE and LLMDet.

3008 **Strategies on Temporal Disambiguation.** As described in §D.3, SAM 3 adopts several strategies to
 3009 address ambiguities in videos. In Tab. 38, we also report two other settings where we turn off all
 3010 these temporal disambiguation strategies (“SAM 3 w/o any temporal disambiguation”) or only turn
 3011 off re-prompting (“SAM 3 w/o re-prompting”). The results show that the disambiguation strategies
 3012 boost the video PCS performance (especially under the pHOta metric). We also find that these
 3013 disambiguation strategies (especially re-prompting) notably improve the qualitative outputs on videos.

3014 3015 3016 G.6 PVS DETAILS

3017 We evaluate SAM 3 on a range of geometric tasks. The standard semi-supervised Video Object
 3018 Segmentation (VOS) task requires tracking an object throughout a video given an input segmentation
 3019 mask. In Tab. 6,

3020 **Baselines.** We compare SAM 3 with recent state-of-the-art models on the VOS task, including
 3021 SAMURAI (Yang et al., 2024), SAM2Long (Ding et al., 2024), and SeC (Zhang et al., 2025). SAM 3
 3022 brings gains on all datasets, including the challenging MOSEv2 benchmark (Ding et al., 2025).

SA-Co/VEval SA-V val (2.0K NPs)											SA-Co/VEval SA-V test (2.0K NPs)						
Model	CGF ₁	VL_MCC	p _{mF₁}	pHOTA	p _{DetA}	p _{AssA}	TETA	CGF ₁	VL_MCC	p _{mF₁}	pHOTA	p _{DetA}	p _{AssA}	TETA			
Human	52.5	0.76	68.6	68.1	50.9	91.6	63.4	53.2	0.78	68.0	68.0	51.3	90.7	66.0			
GLEE [†] (prompted w/ all NPs at once)	0.1	0.13	0.6	7.3	2.3	28.4	5.7	0.2	0.14	1.2	8.8	2.6	33.5	6.6			
GLEE [†] (prompted w/ one NP at a time)	0.1	0.02	3.5	12.7	3.7	45.1	15.7	0.0	0.01	4.2	12.3	3.5	46.0	14.9			
LLMDet [†] as detector + SAM 3 Tracker	3.0	0.11	25.9	33.2	13.7	80.9	34.4	3.6	0.13	27.3	31.2	12.2	80.5	32.3			
SAM 3 Detector + T-by-D as tracker	20.1	0.58	34.4	50.8	35.6	73.2	44.7	22.2	0.61	36.5	49.0	32.4	74.9	45.0			
SAM 3 w/o any temporal disambiguation	24.3	0.51	47.6	54.6	36.2	83.0	52.4	26.5	0.56	47.6	53.9	34.9	84.0	52.0			
SAM 3 w/o re-prompting	26.0	0.55	47.1	55.3	37.4	82.6	48.8	27.8	0.59	47.2	54.0	35.1	83.9	49.0			
SAM 3	26.0	0.55	47.2	55.4	37.4	82.9	48.6	27.8	0.59	47.2	53.9	35.1	83.8	48.7			

(a) Results on SA-Co/VEval SA-V val and test

SA-Co/VEval YT-Temporal-1B val (1.6K NPs)											SA-Co/VEval YT-Temporal-1B test (1.7K NPs)						
Model	CGF ₁	VL_MCC	p _{mF₁}	pHOTA	p _{DetA}	p _{AssA}	TETA	CGF ₁	VL_MCC	p _{mF₁}	pHOTA	p _{DetA}	p _{AssA}	TETA			
Human	72.9	0.95	77.1	78.1	68.0	90.0	82.2	73.8	0.97	76.2	79.2	70.0	90.0	82.4			
GLEE [†] (prompted w/ all NPs at once)	1.9	0.32	6.0	18.7	9.3	39.4	15.6	2.1	0.32	6.4	18.9	9.4	39.1	15.4			
GLEE [†] (prompted w/ one NP at a time)	3.4	0.29	11.7	22.8	13.1	40.4	22.8	3.2	0.26	11.9	22.6	12.4	42.3	22.0			
LLMDet [†] as detector + SAM 3 Tracker	11.6	0.38	30.3	43.2	22.9	81.6	41.5	9.2	0.31	29.3	41.5	21.6	79.6	38.3			
SAM 3 Detector + T-by-D as tracker	46.5	0.94	49.6	66.4	58.6	75.6	68.0	44.6	0.94	47.6	64.4	56.8	73.5	65.4			
SAM 3 w/o any temporal disambiguation	50.5	0.90	55.9	69.5	59.9	81.1	69.8	50.2	0.91	55.1	68.4	58.8	80.1	68.5			
SAM 3 w/o re-prompting	52.0	0.93	56.0	69.8	60.1	81.6	70.0	51.2	0.93	55.1	68.7	59.1	80.4	68.1			
SAM 3	52.4	0.93	56.5	70.2	60.3	82.3	70.2	51.7	0.93	55.5	69.2	59.4	81.2	68.7			

(b) Results on SA-Co/VEval YT-Temporal-1B val and test

SA-Co/VEval SmartGlasses val (2.5K NPs)											SA-Co/VEval SmartGlasses test (2.5K NPs)						
Model	CGF ₁	VL_MCC	p _{mF₁}	pHOTA	p _{DetA}	p _{AssA}	TETA	CGF ₁	VL_MCC	p _{mF₁}	pHOTA	p _{DetA}	p _{AssA}	TETA			
Human	53.0	0.88	60.5	66.3	49.9	88.4	69.8	57.6	0.89	64.8	70.3	55.3	89.7	74.2			
GLEE [†] (prompted w/ all NPs at once)	0.0	0.20	0.2	4.9	2.2	12.0	7.9	0.1	0.19	0.4	5.5	2.2	15.0	8.7			
GLEE [†] (prompted w/ one NP at a time)	0.1	0.19	0.3	5.7	2.4	14.4	13.0	0.1	0.16	0.4	6.4	2.3	18.5	14.1			
LLMDet [†] as detector + SAM 3 Tracker	0.0	0.00	1.8	7.5	0.8	72.4	4.3	0.0	0.00	1.8	7.4	0.8	75.7	4.5			
SAM 3 Detector + T-by-D as tracker	27.7	0.85	32.5	54.9	45.3	67.2	57.7	29.3	0.85	34.4	57.1	47.1	69.8	59.4			
SAM 3 w/o any temporal disambiguation	34.2	0.77	44.2	59.8	45.7	78.8	63.5	35.7	0.78	45.5	61.3	46.9	80.7	64.8			
SAM 3 w/o re-prompting	36.0	0.82	44.0	60.7	46.9	79.1	63.3	37.8	0.82	45.9	62.4	48.6	80.7	64.5			
SAM 3	36.0	0.82	44.0	60.8	47.0	79.3	63.4	38.2	0.82	46.5	62.9	48.8	81.7	64.8			

(c) Results on SA-Co/VEval SmartGlasses val and test

LVVIS test (1.2K NPs)							BURST test (482 NPs)			YTVis21 val (40 NPs)		OVIS val (25 NPs)	BDD100K val (8 NPs)	GMOT40 (10 NPs)	DeepSeaMOT (29 NPs)
Model	mAP	mAP _{base}	mAP _{novel}	HOTA	HOTA _{common}	HOTA _{uncommon}	mAP	mAP	TETA	HOTA	HOTA				
GLEE (prompted w/ all NPs at once)	20.8	24.0	18.4	28.4	46.4	24.8	62.2	38.7	18.0	36.8	6.8				
GLEE (prompted w/ one NP at a time)	9.3	13.9	5.9	20.2	42.9	15.7	56.5	32.4	14.9	29.9	22.9				
LLMDet as detector + SAM 3 Tracker	15.5	15.6	15.5	35.1	49.8	32.1	32.4	47.5	15.2	26.8	25.5				
SAM 3 Detector + T-by-D as tracker	37.3	33.5	40.1	40.7	60.2	36.9	57.3	54.9	49.6	57.8	35.3				
SAM 3	38.2	33.4	41.9	45.9	63.6	42.4	56.9	59.9	48.2	59.9	37.0				

(d) Results on public benchmarks

Table 38: More details on video PCS from a text prompt (open-vocabulary video instance segmentation) on SA-Co/VEval and public benchmarks (# NPs of each benchmark in parentheses). SAM 3 excels in both frame-level detection (p_{DetA}) and cross-frame association (p_{AssA}) and largely outperforms the baselines, especially on benchmarks with a very large number of noun phrases. [†]: GLEE and LLMDet have not been trained on the SA-Co dataset, so their results should be seen as zero-shot on SA-Co/VEval.

For the interactive image segmentation task, we evaluate SAM 3 on the 37-dataset benchmark introduced in SAM 2. As shown in Tab. 7, SAM 3 outperforms SAM 1 and SAM 2 on average mIoU, producing more accurate segmentation masks when prompted with 1 or 5 clicks.

3062
3063
3064
3065
3066
3067
3068
3069
3070
3071
3072
3073
3074
3075
3076
3077

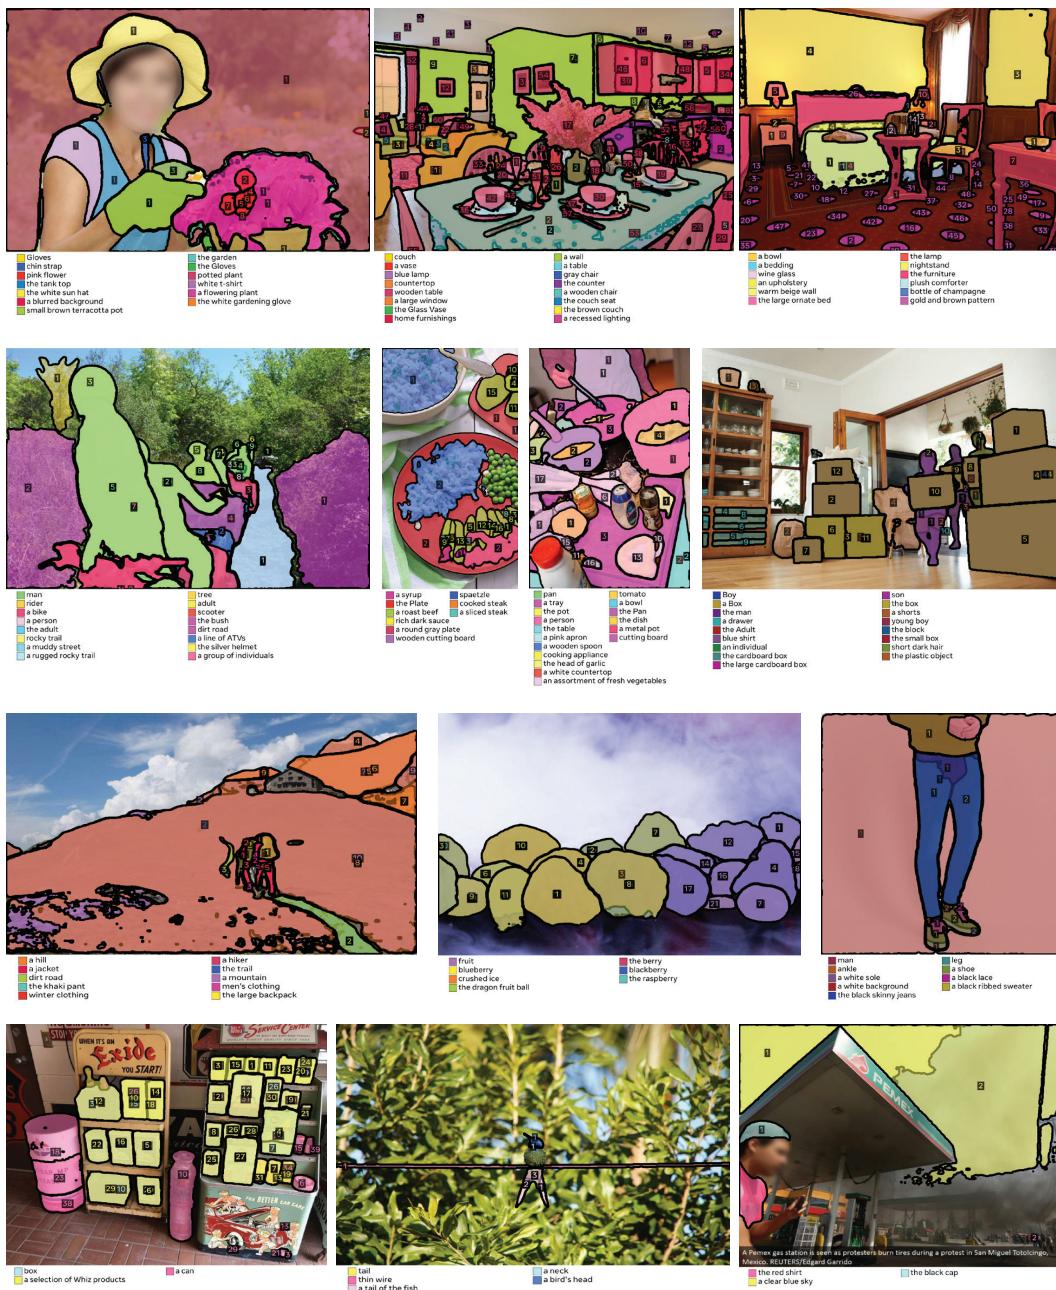
3078 G.6.1 ADDITIONAL MODEL OUTPUTS FOR DIFFERENT TASKS
3079

Figure 18: Example predictions on the SACO Gold dataset.

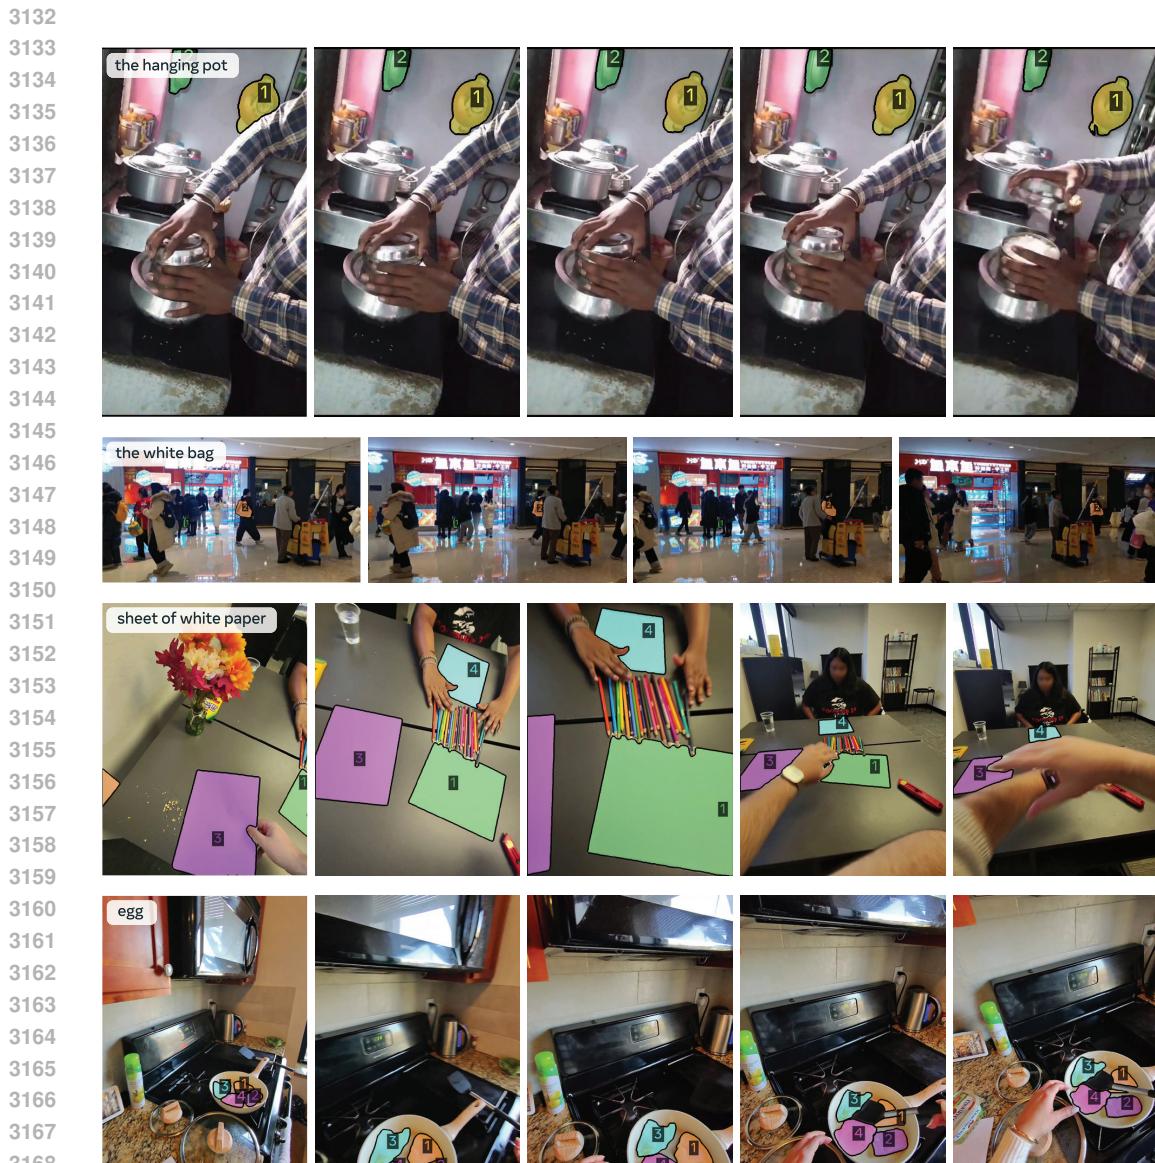


Figure 19: Example video concept segmentation predictions on the SA-Co/VEval SA-V test set (top two rows) and the SA-Co/VEval SmartGlasses test set (bottom two rows).



Figure 20: Example predictions on the countbench dataset.

3186 **H SAM 3 AGENT**

3187

3188

3189

3190 **H.1 AGENT DESIGN**

3191

3192 In this section, we introduce SAM 3 Agent, an agentic pipeline that turns natural-language seg-
 3193 mentation requests into precise masks through iteratively prompting a multimodal LLM (MLLM)
 3194 and SAM 3. Given an input image and a user request, an MLLM acts as a planner/controller: it
 3195 analyzes the image, devises a step-by-step plan, invokes SAM 3 to generate masks, inspects the
 3196 results, and finalizes candidate masks. After each action, the agent receives visual and textual feed-
 3197 back describing the updated environment state, enabling it to revise the plan and decide the next
 3198 action. This perception-action loop continues until the agent is confident it has satisfied the goal (or
 3199 determines that no valid mask exists), at which point it returns a final set of masks. The resulting
 3200 pipeline handles queries far more complex than simple noun phrases which require understanding
 3201 relationships between scene elements and visual common sense.

3202 Each action consists of calling one of several “tools”. We define the following four basic tools for the
 3203 MLLM to call: *segment_phrase*, *examine_each_mask*, *select_masks_and_return*, and *report_no_mask*.
 3204 Among these four tools, *select_masks_and_return* and *report_no_mask* are return tools, which will
 3205 trigger a return function and end the current task. The other two functions: *segment_phrase* and
 3206 *examine_each_mask*, are intermediate tools, which will either call the SAM 3 model on a noun phrase
 3207 or trigger an iterative process for the MLLM to examine each generated mask.

3210 **Tool #1: Segment Phrase**

3211

3212 (1) **Definition:** Use the Segment Anything 3 model to ground all instances of a simple noun phrase
 3213 by generating segmentation mask(s) that cover those instances on the raw input image. At the
 3214 same time, all previously generated mask(s) will be deleted and cannot be referred to in future
 3215 messages.

3216 (2) **Use cases:** Given a simple, direct, and singular noun phrase (not a referring expression that re-
 3217 quires additional understanding/reasoning), *segment_phrase* will try to locate all object instance(s)
 3218 on the raw input image that match the simple noun phrase you provided. The tool will also render
 3219 all of the generated segmentation mask(s) onto the image for you to examine and decide the next
 3220 step.

3221 (3) **Parameters:** {"type": "object", "properties": {"text_prompt": {"type": "string", "description":
 3222 "A short and simple noun phrase, e.g., rope, bird beak, speed monitor, brown handbag, person
 3223 torso"}}, "required": ["text_prompt"]}

3224 (4) **Return type:** A new image with differently colored segmentation mask(s) rendered on it, and a
 3225 text message indicating the number of mask(s) generated by the Segment Anything 3 model for
 3226 this “text_prompt” only.

3227 **Tool #2: Examine Each Mask**

3228

3229 (1) **Definition:** Use this tool when the *segment_phrase* tool generates multiple small or overlapping
 3230 mask(s), making it difficult to distinguish the correct mask(s). *examine_each_mask* allows you to
 3231 render and examine each mask independently to see small mask(s) clearly and avoid confusing
 3232 overlapping mask(s).

3233 (2) **Use cases:** Sometimes there are multiple small mask(s) or overlapping mask(s) rendered on an
 3234 image, making it difficult to distinguish each mask from others. In this case, you should call the
 3235 *examine_each_mask* tool to individually verify each mask and filter out incorrect mask(s).

3236 (3) **Parameters:** None

3237 (4) **Return type:** A new image with colored segmentation mask(s) accepted by the *exam-
 3238 ine_each_mask* tool, and a text message indicating how many masks were accepted.

3240

3241

3242

3243

3244

3245

3246

3247

3248

3249

3250

3251

3252

3253

3254

3255

3256

3257

3258

3259

3260

3261

3262

3263

3264

3265

3266

3267

3268

3269

3270

3271

3272

3273

3274

3275

3276

3277

3278

3279

3280

3281

3282

3283

3284

3285

3286

3287

3288

3289

3290

3291

3292

3293

Tool #3: Select Masks And Return

- (1) **Definition:** Call this tool to select a subset of or all of the mask(s) rendered on the most recent image as your final output. When calling `select_masks_and_return`, you cannot select any mask(s) generated by previous rounds other than the most recent round in your `"final_answer_masks"`. You can only use mask(s) from the most recent image in your message history.
- (2) **Use cases:** Given an image with one or more segmentation mask(s) already rendered on it, `select_masks_and_return` returns the set of mask(s) you select as the final output.
- (3) **Parameters:** `{"type": "object", "properties": {"final_answer_masks": {"type": "array", "description": "An array of integers representing the selected mask(s) you want to choose as your final output, e.g., [1, 4, 5]"}, "required": ["final_answer_masks"]}`
- (4) **Return type:** None (End of Conversation)

Tool #4: Report No Mask

- (1) **Definition:** Call this tool when you are absolutely sure that there are no object(s) in the image that match or answer the initial user input query.
- (2) **Use cases:** Reporting that the given image does not contain any target object(s) that match or answer the initial user input query.
- (3) **Parameters:** None
- (4) **Return type:** None (End of Conversation)

After each intermediate tool call has been executed, the system will provide the MLLM with the following two pieces of information:

- The user input image with all generated and currently available segmentation masks rendered on it in a Set-of-Marks (Yang et al., 2023) manner. The masks are randomly colored and numbered from 1 to N based on SAM 3 confidence scores received at the time of mask generation. This defines the environment state of the SAM 3 Agent at the current step, which includes the original user input image and all available masks.
- An automatically generated text message stating all changes from the previous environment state (e.g. how many masks have been generated by the `segment_phrase` tool, or how many masks were removed by the `examine_each_mask` tool).

After analyzing the updated image with available masks (current environment state) along with the initial user query (goal), the MLLM must update its tool-calling plan and generate its next tool call (current action). We allow the MLLM to call each intermediate tool as many times as it needs, before arriving at a set of segmentation masks (final state) that it is satisfied with.

Empirically, we observe that for especially challenging queries, SAM 3 Agent may produce as many as 60 steps of trial and error before being satisfied with its grounding outcome and calling a return tool. This results in an extremely long environment-state context history with each step containing a new image, pushing both the context limit and multi-image reasoning capability of even current state-of-the-art MLLMs.

To resolve this issue, we propose an aggressive context engineering mechanism that prunes all intermediate trial-and-error states between the initial user text query and the most recent agent call to the `segment_phrase` tool. We also replace all previously generated masks after each new MLLM call to the `segment_phrase` tool with newly generated masks. Empirically we found this design to be much simpler and more performant compared to allowing previously generated masks to accumulate and clutter the newly rendered images. To avoid losing important failure experience from pruned steps, we provide a continuously updated list of all previously used (and discarded) SAM 3 noun phrase prompts for the model to note.

H.2 QUALITATIVE ANALYSIS

In this section, we provide success 21 and failure 22 examples of SAM 3 Agent on the ReasonSeg (Lai et al., 2024) and RefCOCOg (Kazemzadeh et al., 2014) datasets, as they are currently the most

3294

3295

3296

3297

3298

3299

3300

3301

3302

3303



In the animal kingdom, the males of certain species have distinctive features that set them apart from females. What characteristic in this image shows the animal is male?

3304

3305

3306

3307

3308

3309

3310

3311

3312

3313

3314

3315

3316

3317

3318

3319

the camera lens that is more suitable for photographing nearby objects



the transportation vehicle that does not require electricity or gasoline



If the person in the picture is the owner of the blue car and wants to drive it, which car door is least likely to be the one he enters?



What part of this picture is funny and out of place?



If the police wanted to identify this vehicle, which part of this picture would contain useful information?

3340

3341

3342

3343

3344

3345

3346

3347



Places where bees can suck nectar from



The white apple computer monitor displaying two web pages



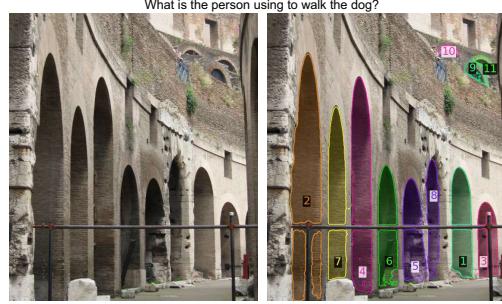
What object in this picture might reflect this person's marital status?



I want to take a trip around the world, but I need some transportation to help me cross the oceans. What type of transportation in the picture can fulfill this requirement?



What is the person using to walk the dog?



In historical architecture, buildings often have entrances that consist of large openings with a curved or pointed top. What feature in the picture resembles such an entrance?

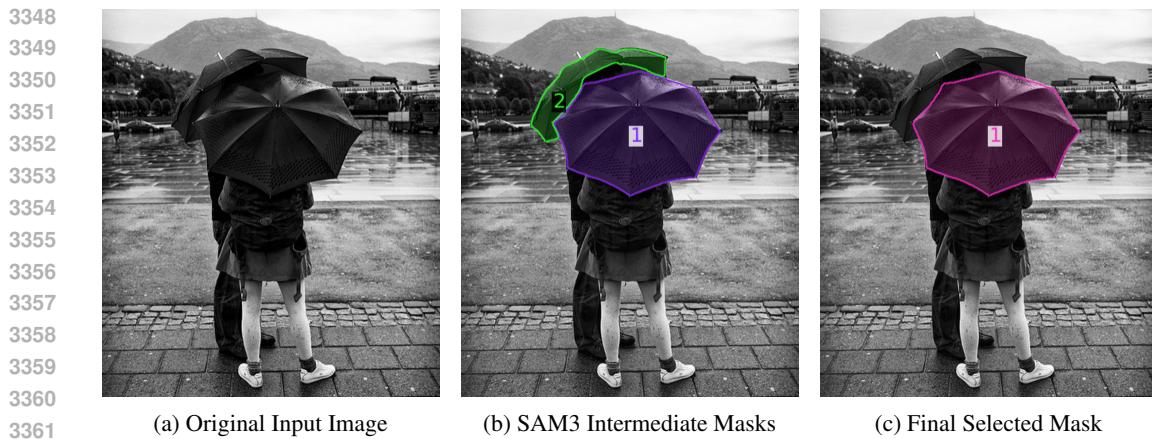


Figure 22: Error example of SAM 3 Agent (Qwen2.5-VL 72B) including the original input image and SAM 3 intermediate output masks. The textual user query is “A black object that protects you from the rain, being held by a person in jeans”. The agent was able to reason and find a suitable grounding target for calling SAM 3 (“black umbrella”). However, it failed to select the correct mask during final mask selection due to a visual reasoning error by the MLLM.

challenging and widely used reasoning segmentation and referring expression segmentation datasets. We also provide a complete reasoning trace example of SAM 3 Agent, demonstrating how SAM 3 Agent solves complex reasoning segmentation queries by leveraging SAM 3 for precise grounding and MLLM for visual reasoning.

Empirically, we observe that SAM 3 Agent is able to handle free-form textual queries of varying spatial and logical reasoning complexity. It is able to reject queries that do not point to any object in the given image and to generate multiple output masks for queries that match multiple objects in the given image. The MLLM backbone allows SAM 3 Agent to robustly handle queries of varying length and format. It also enables SAM 3 Agent to accurately segment fine-grained concepts that SAM 3 alone struggle with, such as specific aircraft types and fine-grained food categories.

SAM3 Agent Complete Reasoning Trace Example

```
[  
  {  
    "role": "system",  
    "content": {System prompt including tool definitions, agent instructions, and  
    response format.}  
  },  
  {  
    "role": "user",  
    "content": [  
      {  
        "type": "image",  
        "image": Figure {\ref{fig:agent_mario_1}}  
      },  
      {  
        "type": "text",  
        "text": "the stronger Mario"  
      }  
    ]  
  },  
  {  
    "role": "assistant",  
    "content": [  
      {  
        "type": "text",  
        "text": "<think> There is only one image in the context. The image shows a  
        banner with pixel art characters resembling Mario. The user wants to ground \"the  
        stronger Mario,\" which appears to be the character on the right side of the equation  
        depicted on the banner. I will use the 'segment_phrase' function with the text  
        prompt \"stronger Mario\" to attempt to ground the target object.\nBefore  
        proceeding, I remind myself that each call to 'segment_phrase' will delete all  
        previous segments."  
      }  
    ]  
  }]
```

```

3402
3403     previously generated masks, so I need to ensure my plan is effective in one go. I
3404     will start by attempting to ground \"stronger Mario.\\"\\n\\n</think>\\n<tool> {\"name\":
3405         \"segment_phrase\", \"parameters\": {\"text_prompt\": \"stronger Mario\"}} </tool>\\n
3406         }
3407     ]
3408     {
3409         "role": "user",
3410         "content": [
3411             {
3412                 "type": "text",
3413                 "text": "SAM3 did not generate any masks for the prompt 'stronger Mario',
3414                 please call the segment_phrase tool again with a different, perhaps more general and
3415                 simple noun phrase prompt."
3416             }
3417         ],
3418         {
3419             "role": "assistant",
3420             "content": [
3421                 {
3422                     "type": "text",
3423                     "text": "<think> There is still only one image in the context. The
3424                     previous attempt to ground \"stronger Mario\" did not yield any masks. I need to try
3425                     a simpler and more general noun phrase that still captures the essence of the target.
3426                     Given the pixel art style, I'll try grounding \"Mario\" to see if it captures the
3427                     character on the right side of the equation.\\n\\nI remind myself that each call to 'seg-
3428                     ment_phrase' will delete all previously generated masks, so I need to ensure my
3429                     plan is effective in one go. I will now attempt to ground \"Mario.\\"\\n\\n</think>\\n<
3430                     tool> {\"name\": \"segment_phrase\", \"parameters\": {\"text_prompt\": \"Mario\"}} </
3431                     tool>\\n\\n"
3432                 }
3433             ],
3434             {
3435                 "role": "user",
3436                 "content": [
3437                     {
3438                         "type": "text",
3439                         "text": "SAM3 did not generate any masks for the prompt 'Mario', please
3440                         call the segment_phrase tool again with a different, perhaps more general and simple
3441                         noun phrase prompt."
3442                     }
3443                 ],
3444             },
3445             {
3446                 "role": "assistant",
3447                 "content": [
3448                     {
3449                         "type": "text",
3450                         "text": "The original user query was: 'the stronger Mario', and the SAM3
3451                         prompt you used was: 'pixel character'. SAM3 generated 4 masks for the prompt 'pixel
3452                         character', which are rendered in the latest image, now please analyze this new image
3453                         and decide your next action."
3454                     },
3455                     {
3456                         "type": "image",
3457                         "image": Figure {\\ref{fig:agent_mario_2}}
3458                     }
3459                 ],
3460             },
3461         ],
3462     },
3463 
```



Figure 23: Images from the Complete Reasoning Trace Example of SAM 3 Agent (Qwen2.5-VL 72B), including images (a) and (b) from the agent context and image (c) containing the final agent output mask. The textual user query is “the stronger Mario”. The SAM 3 Intermediate Masks are generated by SAM 3 using the text prompt “pixel character”.

```

3466
3467
3468
3469
3470
3471
3472
3473
3474
3475
3476
3477
3478
3479
3480
3481
3482
3483
3484
3485
3486
3487
3488
3489
3490
3491
3492
3493
3494
3495
3496
3497
3498
3499
3500
3501
3502
3503
3504
3505
3506
3507
3508
3509

{
    "role": "assistant",
    "content": [
        {
            "type": "text",
            "text": "<think> There are two images in the context now: the raw input image and the current output image with four masks rendered. Mask 2 corresponds to the \"stronger Mario\" character, which is the target of the user's query. Masks 1, 3, and 4 are not relevant to the query. \n\nSince I have identified the correct mask, I will proceed to select it as the final answer.\n</think>\n<tool> {\"name\": \"select_masks_and_return\", \"parameters\": {\"final_answer_masks\": [2]}} </tool>"
        }
    ],
    "role": "user",
    "content": [
        {
            "type": "text",
            "text": "Final Answer:"
        },
        {
            "type": "image",
            "image": Figure {\ref{fig:agent_mario_3}}
        }
    ],
]
}

```

H.3 FULL QUANTITATIVE RESULTS

We report the full quantitative experiment results of SAM 3 Agent on ReasonSeg in Tab. 39, OmniLabel in Tab. 40, and RefCOCO-Seg in Tab. 41. SAM 3 Agent achieves the best results on both ReasonSeg and OmniLabel in a zero-shot manner, without training on any referring expression segmentation or reasoning segmentation data. SAM 3 Agent also surpasses previous zero-shot SOTA results on RefCOCO+ and RefCOCOg, and is close to training-based SoTA. We hypothesize that on RefCOCO, where all masks come from the MSCOCO dataset and each query points to exactly one ground-truth object mask, training-based methods learn the specific dataset annotation biases. We show examples of such annotation biases in the RefCOCO-Seg datasets in Fig. 24. SAM 3 Agent, being a zero-shot method, is unable to exploit these (generally undesirable) biases.

3510	3511	Model		Training		Val Set		Test Set		Test (Short)		Test (Long)	
		Name	Version	RES	ReasonSeg	gIoU	cIoU	gIoU	cIoU	gIoU	cIoU	gIoU	cIoU
3513	SEEM	-		x	x	25.5	21.2	24.3	18.7	20.1	11.5	25.6	20.8
3514	Grounded SAM	-		x	x	26.0	14.5	21.3	16.4	17.8	10.8	22.4	18.6
3515	OVSeg	-		x	x	28.5	18.6	26.1	20.8	18.0	15.5	28.7	22.5
3516	GLaMM	Vicuna 7B	✓	x	x	47.4	47.2	-	-	-	-	-	-
3517	SAM4MLLM	Qwen-VL 7B	✓	x	x	46.7	48.1	-	-	-	-	-	-
3518	SAM4MLLM	LLaVA1.6 8B	✓	x	x	58.4	60.4	-	-	-	-	-	-
3519	Seg-Zero	Qwen2.5-VL 3B	✓	x	x	58.2	53.1	56.1	48.6	-	-	-	-
3520	Seg-Zero	Qwen2.5-VL 7B	✓	x	x	62.6	62.0	57.5	52.0	-	-	-	-
3521	X-SAM	Phi3 3.8B	✓	✓	x	56.6	32.9	57.8	41.0	47.7	48.1	56.0	40.8
3522	HyperSeg	Phi2 3B	✓	✓	x	59.2	56.7	-	-	-	-	-	-
3523	Kang et al.	LLaVA1.5 7B	x	x	x	-	52.4	-	48.7	-	48.0	-	49.1
3524	Kang et al.	LLaVA1.5 13B	x	x	x	-	60.5	-	49.9	-	48.7	-	51.0
3525	LISA	LLaVA 7B	✓	x	x	44.4	46.0	36.8	34.1	37.6	34.4	36.6	34.7
3526	LISA	LLaVA 7B	✓	✓	x	52.9	54.0	47.3	34.1	40.6	40.6	49.4	51.0
3527	LISA	LLaVA 13B	✓	x	x	48.9	46.9	44.8	45.8	39.9	43.3	46.4	46.5
3528	LISA	LLaVA 13B	✓	✓	x	56.2	62.9	51.7	51.1	44.3	42.0	54.0	54.3
3529	LISA	Llama2 13B	✓	✓	x	60.0	67.8	51.5	51.3	43.9	45.8	54.0	53.8
3530	LISA	LLaVA1.5 7B	✓	x	x	53.6	52.3	48.8	47.1	48.3	48.8	49.2	48.9
3531	LISA	LLaVA1.5 7B	✓	✓	x	61.3	62.9	55.6	56.9	48.3	46.3	57.9	59.7
3532	LISA	LLaVA1.5 13B	x	x	x	57.7	60.3	53.8	50.8	50.8	50.0	54.7	50.9
3533	LISA	LLaVA1.5 13B	✓	✓	x	65.0	72.9	61.3	62.2	55.4	50.6	63.2	65.3
3534	RSVP	LLaVA1.6 7B	x	x	x	59.2	56.7	56.9	50.7	47.9	42.0	58.4	53.0
3535	RSVP	Qwen2-VL 7B	x	x	x	58.6	48.5	56.1	51.6	48.5	44.3	57.1	53.0
3536	RSVP	Gemini1.5-Flash	x	x	x	56.9	49.2	57.1	59.2	47.3	40.2	60.2	65.6
3537	RSVP	GPT-4o	x	x	x	64.7	63.1	60.3	60.0	55.4	50.4	61.9	62.5
3538	Gemini Seg	Gemini2.5 Flash	?	?	x	28.3	13.3	30.6	9.2	16.5	8.0	35.0	9.5
3539	SAM 3 Agent	Qwen2.5-VL 7B	x	x	x	65.4	50.5	62.6	56.2	59.1	41.8	63.7	57.8
3540	SAM 3 Agent	Qwen2.5-VL 72B	x	x	x	75.0	63.5	71.8	65.2	71.3	55.3	72.0	67.9
3541	SAM 3 Agent	Llama4 Maverick	x	x	x	71.5	<u>67.1</u>	69.3	67.1	70.9	69.6	68.8	66.5
3542	SAM 3 Agent	Gemini2.5 Pro	x	x	x	76.0	63.3	73.8	68.1	74.0	60.4	73.7	70.1

Table 39: SAM 3 Agent experiments on the ReasonSeg dataset (Lai et al., 2024) for Reasoning Segmentation. Training-RES indicates whether the model has been fine-tuned on the Referring Expression Segmentation task. Training-ReasonSeg indicates whether the model has been fine-tuned on the ReasonSeg training set. The overall best performances are shown in **bold** and the best zero-shot performances for models not trained on the ReasonSeg training set are underlined. Notable baselines include: SEEM (Zou et al., 2023), Grounded SAM (Ren et al.), OVSeg (Liang et al., 2023), GLaMM (Rasheed et al., 2024), SAM4MLLM (Chen et al., 2024b), Seg-Zero (Liu et al., 2025), X-SAM (Wang et al., 2025a), HyperSeg (Wei et al., 2024), (Kang et al., 2025), LISA (Lai et al., 2024), RSVP (Lu et al., 2025) and Gemini-seg (Paul Voigtlaender, Valentin Gabeur and Rohan Doshi, 2025)



Figure 24: Examples of annotation bias and ground truth errors from the RefCOCO-Seg datasets (Kazemzadeh et al., 2014; Mao et al., 2016). For each example, see the original dataset ground truth annotation (left image), the textual user query (bottom text), and the SAM 3 Agent (Qwen2.5-VL 72B) final segmentation output (right image). Our error analysis reveals such annotation bias and ground truth errors account for the majority of low-IoU predictions by SAM 3 Agent on the RefCOCO-Seg datasets.

3564

3565

3566

3567

3568

3569

3570

3571

3572

3573

3574

3575

3576

3577

3578

3579

3580	Model		2023 Val Set (AP @ IoU=0.50:0.95)						
	Name	Version	AP	AP-categ	AP-descr	AP-descr-pos	AP-descr-S	AP-descr-M	AP-descr-L
3581	FIBER	FIBER-B	25.7	30.3	22.3	34.8	38.6	19.5	12.4
3582	GLIP	GLIP-T	19.3	23.6	16.4	25.8	29.4	14.8	8.2
3583	GLIP	GLIP-L	25.8	32.9	21.2	33.2	37.7	18.9	10.8
3584	Zhao et al.	GLIP-T	22.2	27.2	18.8	29.0	—	—	—
3585	Zhao et al.	FIBER-B	28.1	32.1	25.1	36.5	—	—	—
3586	DesCo	GLIP-T	23.8	27.4	21.0	30.4	—	—	—
3587	DesCo	FIBER-B	29.3	31.6	27.3	37.7	—	—	—
3588	GLEE	Lite	20.3	37.5	14.0	19.1	23.0	12.7	10.0
3589	GLEE	Lite-Scale	22.7	35.5	16.7	22.3	33.7	14.3	10.2
3590	GLEE	Plus	25.4	46.7	17.5	23.9	28.4	16.3	12.5
3591	GLEE	Plus-Scale	27.0	44.5	19.4	25.9	36.0	17.2	12.4
3592	Real	Swin-B	—	—	36.5	52.1	54.4	33.2	25.5
3593	LED	Qwen2	27.9	33.9	23.7	36.2	—	—	—
3594	LED	InternLM2-2B	27.9	33.4	23.9	36.1	—	—	—
3595	LED	InternLM2-6B	26.3	32.0	22.4	34.3	—	—	—
3596	ROD-MLLM	Vicuna 7B	—	—	25.3	30.9	31.8	24.5	21.0
3597	WSCL	GLIP-T	24.3	23.9	24.7	34.4	39.3	21.6	16.4
3598	WSCL	FIBER-B	30.5	31.6	29.5	40.3	43.7	26.3	21.3
3599	WSCL	Desco-GLIP	26.5	27.1	25.9	35.6	38.1	23.2	18.7
3600	WSCL	Desco-FIBER	32.0	33.1	30.9	40.4	45.2	27.7	22.9
3601	SAM 3 Agent	Qwen2.5-VL 7B	38.6	41.0*	36.5	43.0	52.6	34.3	26.7
3602	SAM 3 Agent	Qwen2.5-VL 72B	42.8	41.0*	44.7	54.1	58.4	42.6	36.1
3603	SAM 3 Agent	Llama4 Maverick	38.5	41.0*	36.2	46.6	47.5	34.9	28.1
3604	SAM 3 Agent	Gemini2.5 Pro	43.7	41.0*	46.7	64.8	54.6	46.2	38.7

Table 40: SAM 3 Agent experiments on the OmniLabel dataset (Schulter et al., 2023) (val 2023) for generalized referring expression comprehension (box prediction). * indicates predictions generated by the base SAM 3 model without MLLM. The overall best performances are shown in **bold**. Notable baselines include: FIBER (Dou et al., 2022), GLIP (Li et al., 2022b), (Zhao et al., 2024), DesCo (Li et al., 2023c), GLEE (Wu et al., 2024a), Real (Chen et al., 2025), LED (Zhou et al., 2025), ROD-MLLM (Yin et al., 2025), and WSCL (Park et al., 2024).

3605

3606

3607

3608

3609

3610

3611

3612

3613

3614

3615

3616

3617

3618
 3619
 3620
 3621
 3622
 3623
 3624
 3625

3626	Model		Training			RefCOCO		RefCOCO+		RefCOCOg		
	Name	Version	RES	val	testA	testB	val	testA	testB	val (U)	test (U)	val (G)
3627	LISA	LLaVA 7B	✓	74.9	79.1	72.3	65.1	70.8	58.1	67.9	70.6	—
3628	GSVA	13B	✓	79.2	81.7	77.1	70.3	73.8	63.6	75.7	77.0	—
3629	GLaMM	Vicuna 7B	✓	79.5	83.2	76.9	72.6	78.7	64.6	74.2	74.9	—
3630	SAM4MLLM	LLaVA1.6 7B	✓	79.6	82.8	76.1	73.5	77.8	65.8	74.5	75.6	—
3631	SAM4MLLM	LLaVA1.6 8B	✓	79.8	82.7	74.7	74.6	80.0	67.2	75.5	76.4	—
3632	GLEE	Plus	✓	79.5	—	—	68.3	—	—	70.6	—	—
3633	GLEE	Pro	✓	80.0	—	—	69.6	—	—	72.9	—	—
3634	DETTRIS	DETTRIS-L	✓	81.0	81.9	79.0	75.2	78.6	70.2	74.6	75.3	—
3635	UniLSeg	UniLSeg-20	✓	80.5	81.8	78.4	72.7	77.0	67.0	78.4	79.5	—
3636	UniLSeg	UniLSeg-100	✓	81.7	83.2	79.9	73.2	78.3	68.2	79.3	80.5	—
3637	PSALM	Phi1.5 1.3B	✓	83.6	84.7	81.6	72.9	75.5	70.1	73.8	74.4	—
3638	EVF-SAM	RC	✓	82.1	83.7	80.0	75.2	78.3	70.1	76.8	77.4	—
3639	EVF-SAM	Extra Data	✓	82.4	84.2	80.2	76.5	80.0	71.9	78.2	78.3	—
3640	RICE	Qwen2.5-7B	✓	83.5	85.3	81.7	79.4	82.8	75.4	79.8	80.4	—
3641	MLCD-seg	Qwen2.5-7B	✓	83.6	85.3	81.5	79.4	82.9	75.6	79.7	80.5	—
3642	HyperSeg	Phi2 2.7B	✓	84.8	85.7	83.4	79.0	83.5	75.2	79.4	78.9	—
3643	X-SAM	Phi3 3.8B	✓	85.1	87.1	83.4	78.0	81.0	74.4	83.8	83.9	—
3644	GL-CLIP	ResNet-50	✗	32.7	35.3	30.1	37.7	40.7	34.9	41.6	42.9	44.0
3645	GL-CLIP	ViT-B/32	✗	32.9	34.9	30.1	38.4	42.1	32.7	42.0	42.0	42.7
3646	CaR	ViT-B/16	✗	33.6	35.4	30.5	34.2	36.0	31.0	36.7	36.6	36.6
3647	Ref-Diff	VAE	✗	37.2	38.4	37.2	37.3	40.5	33.0	44.0	44.5	44.3
3648	TAS	ResNet-50	✗	39.9	42.9	35.9	44.0	50.6	36.4	47.7	47.4	48.7
3649	TAS	ViT-B/32	✗	39.8	41.1	36.2	43.6	49.1	36.5	46.6	46.8	48.1
3650	IteRPrimeE	—	✗	40.2	46.5	33.9	44.2	51.6	35.3	46.0	45.1	45.8
3651	Pseudo-RIS	CRIS	✗	39.8	44.8	33.0	42.2	46.3	34.5	43.7	43.4	43.8
3652	Pseudo-RIS	ETRIS	✗	41.1	48.2	33.5	44.3	51.4	35.1	46.0	46.7	46.8
3653	LGD+DINO	ViT-B/32	✗	49.5	54.7	41.0	49.6	58.4	38.6	50.3	51.1	52.5
3654	VLM-VG	ResNet-50	✗	47.7	51.8	44.7	41.2	45.9	34.7	46.6	47.1	—
3655	VLM-VG	ResNet-101	✗	49.9	53.1	46.7	42.7	47.3	36.2	48.0	48.5	—
3656	HybridGL	ViT-B/32	✗	49.5	53.4	45.2	43.4	49.1	37.2	51.3	51.6	—
3657	Kang et al.	LLaVA-1.5 7B	✗	74.2	76.5	70.4	62.5	65.2	56.0	64.2	68.1	—
3658	Kang et al.	LLaVA-1.5 13B	✗	<u>76.1</u>	<u>78.9</u>	<u>72.8</u>	64.1	67.1	57.3	67.7	69.0	—
3659	SAM 3 Agent	Qwen2.5-VL 7B	✗	53.4	58.4	48.0	46.3	52.2	40.8	54.5	55.1	54.6
3660	SAM 3 Agent	Qwen2.5-VL 72B	✗	70.5	74.3	64.6	64.9	70.7	56.6	69.2	69.5	69.8
3661	SAM 3 Agent	Llama4 Maverick	✗	72.8	77.7	67.1	65.8	<u>72.2</u>	59.2	69.3	70.1	69.9
3662	SAM 3 Agent	Gemini2.5 Pro	✗	74.9	77.8	69.9	<u>66.9</u>	71.1	<u>62.4</u>	<u>73.3</u>	<u>73.6</u>	74.2

Table 41: SAM 3 Agent experiments on the RefCOCO / RefCOCO+ / RefCOCOg datasets (Kazemzadeh et al., 2014; Mao et al., 2016) for Referring Expression Segmentation (RES). Training-RES indicates whether the model has been fine-tuned on the RefCOCO/RefCOCO+/RefCOCOg segmentation datasets (notice that nearly all MLLMs were trained on RefCOCO/RefCOCO+/RefCOCOg bbox datasets). The overall best performances are shown in **bold** and the best zero-shot performances for models *not* trained on the RES task are underlined. SAM 3 Agent achieves SoTA performance on RefCOCO+/RefCOCOg in the zero-shot setting, and is close to the fine-tuning SoTA. Notable baselines include: LISA (Lai et al., 2024), GSVA (Xia et al., 2024), GLaMM (Rasheed et al., 2024), SAM4MLLM (Chen et al., 2024b), GLEE (Wu et al., 2024a), DETRIS (Huang et al., 2025), UniLSeg (Liu et al., 2024b), PSALM (Zhang et al., 2024c), EVF-SAM (Zhang et al., 2024b), RICE (Xie et al., 2025), MLCD-seg (An et al., 2024), HyperSeg (Wei et al., 2024), X-SAM (Wang et al., 2025a), GL-CLIP (Yu et al., 2023b), CaR (Sun et al., 2024), Ref-Diff (Ni et al., 2023), TAS (Suo et al., 2023), IteRPrimeE (Wang et al., 2025c), Pseudo-RIS (Yu et al., 2024), LGD+DINO (Li et al., 2025), VLM-VG (Wang et al., 2025b), HybridGL (Liu & Li, 2025), and (Kang et al., 2025).

3665
 3666
 3667
 3668
 3669
 3670
 3671